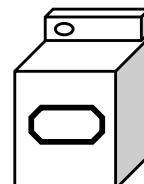


---

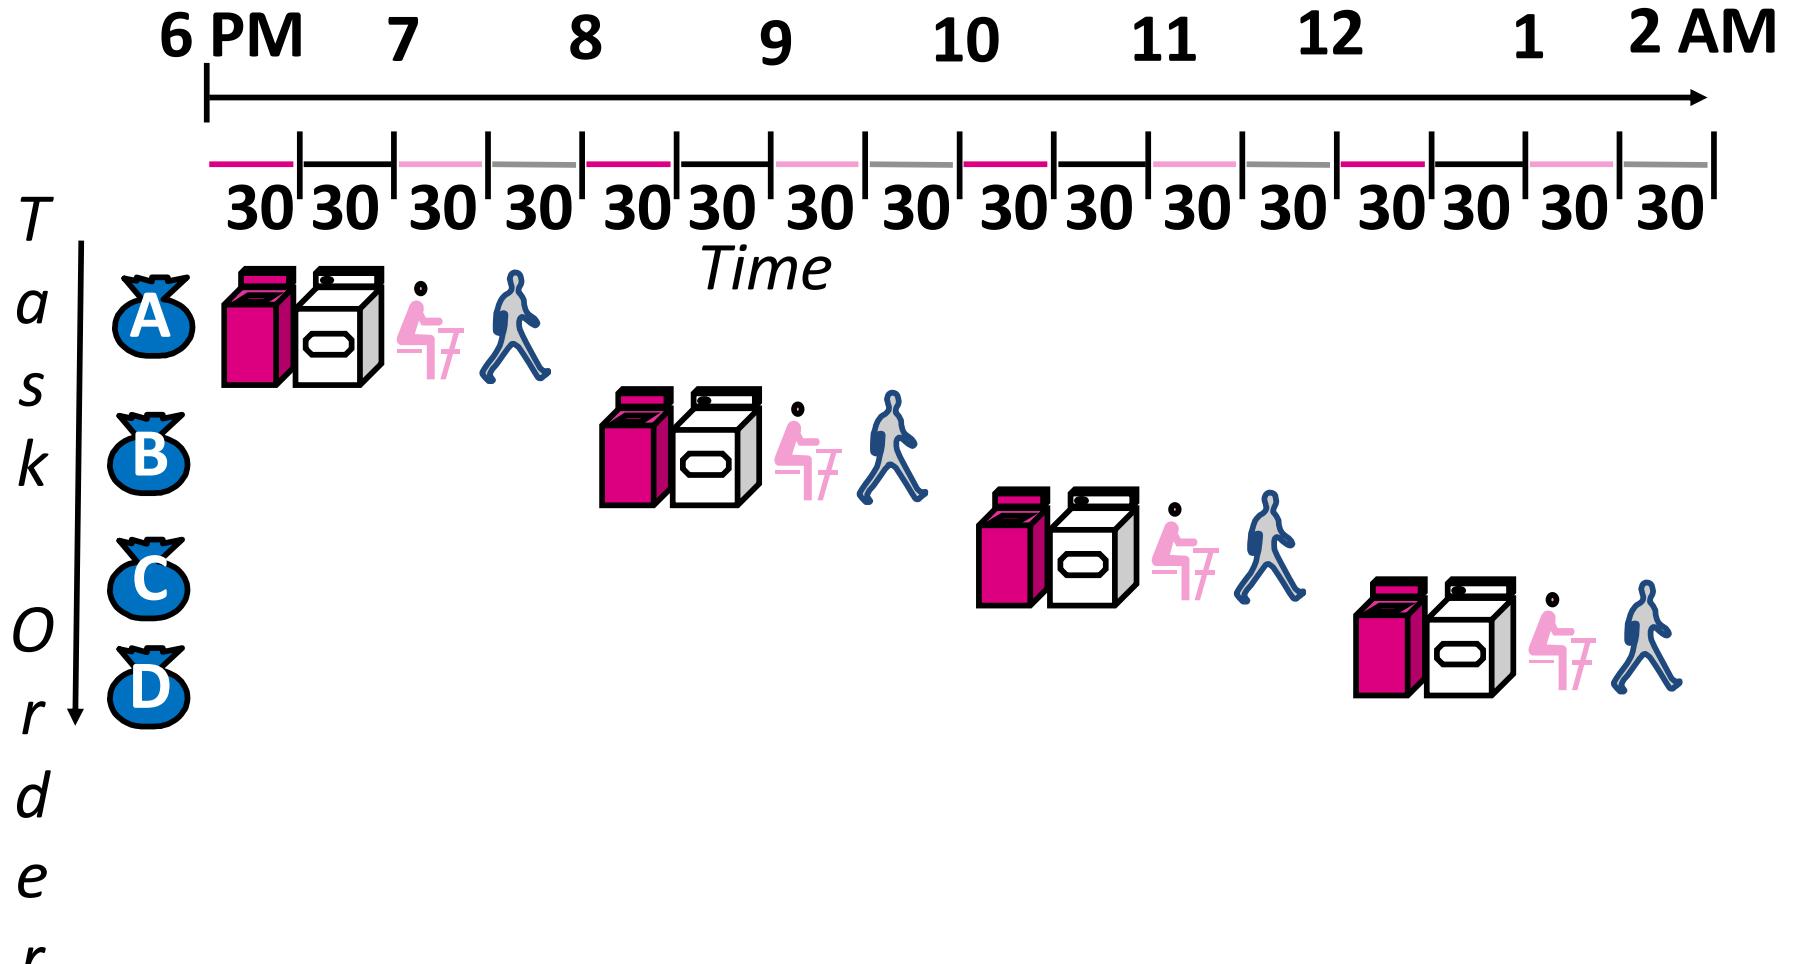
# Pipelining and Hazard

# Pipeline Analogy: Doing Laundry

- Ann, Brian, Cathy, and Dave each have one load of clothes to wash, dry, fold, and put away
  - Washer takes 30 minutes
  - Dryer takes 30 minutes
  - “Folder” takes 30 minutes
  - “Stasher” takes 30 minutes to put clothes into drawers

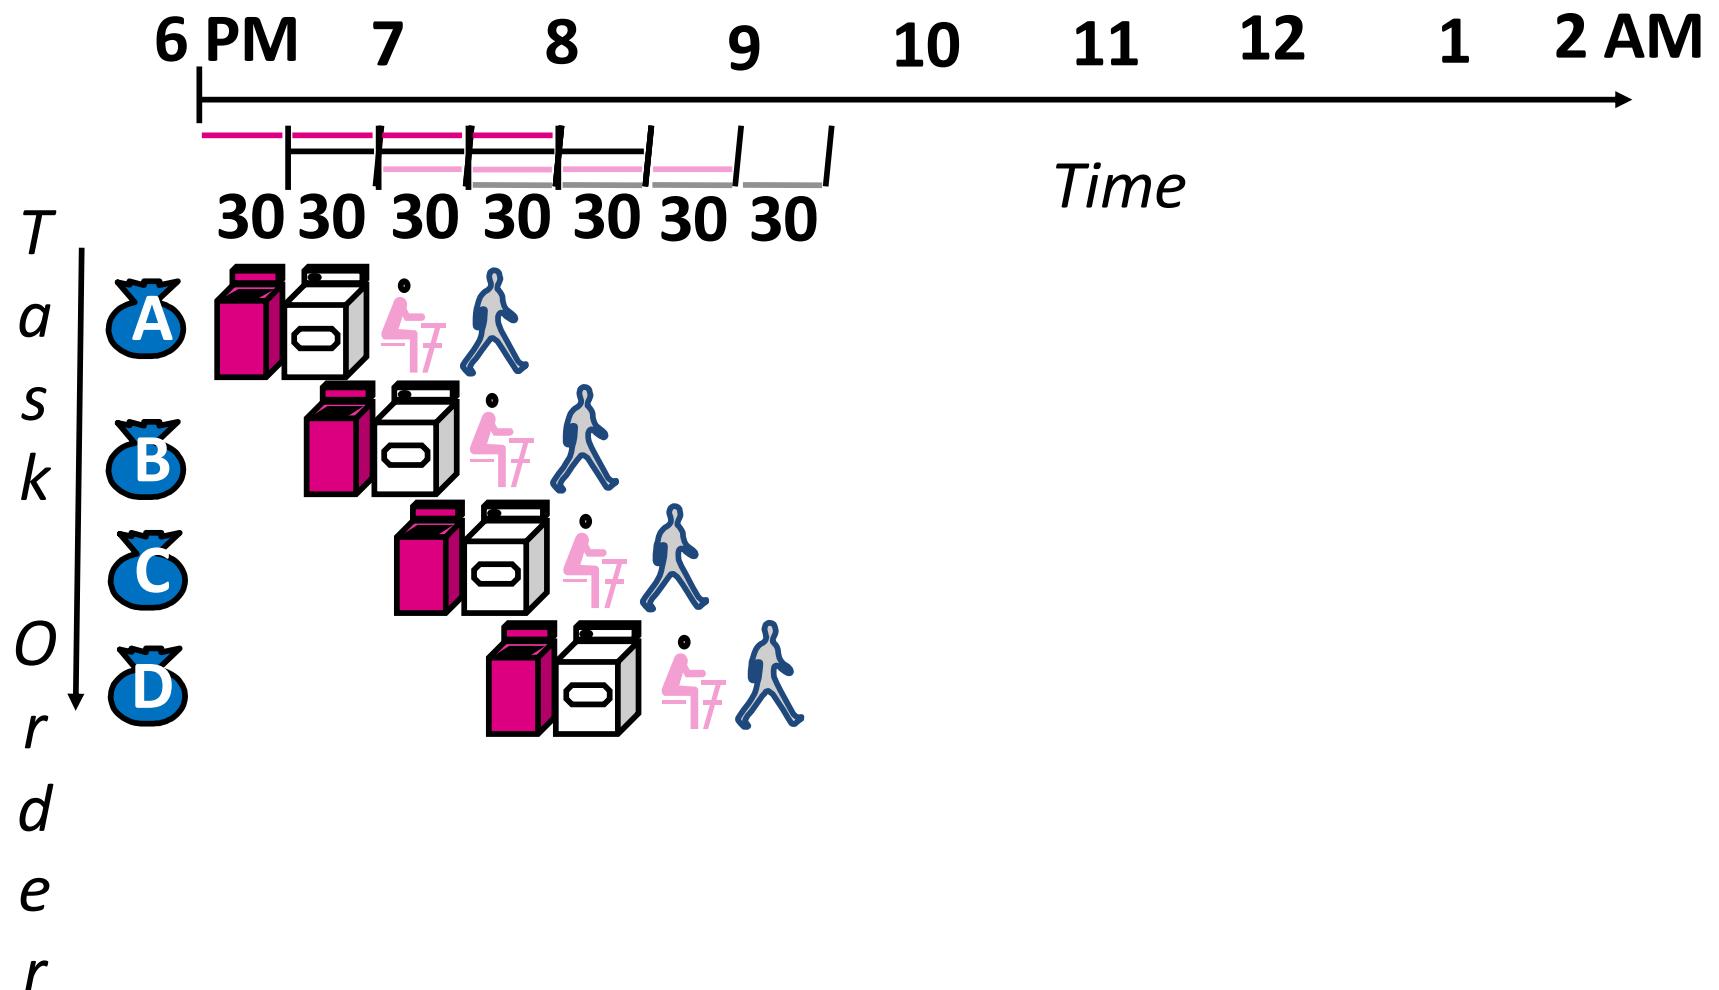


# Sequential Laundry



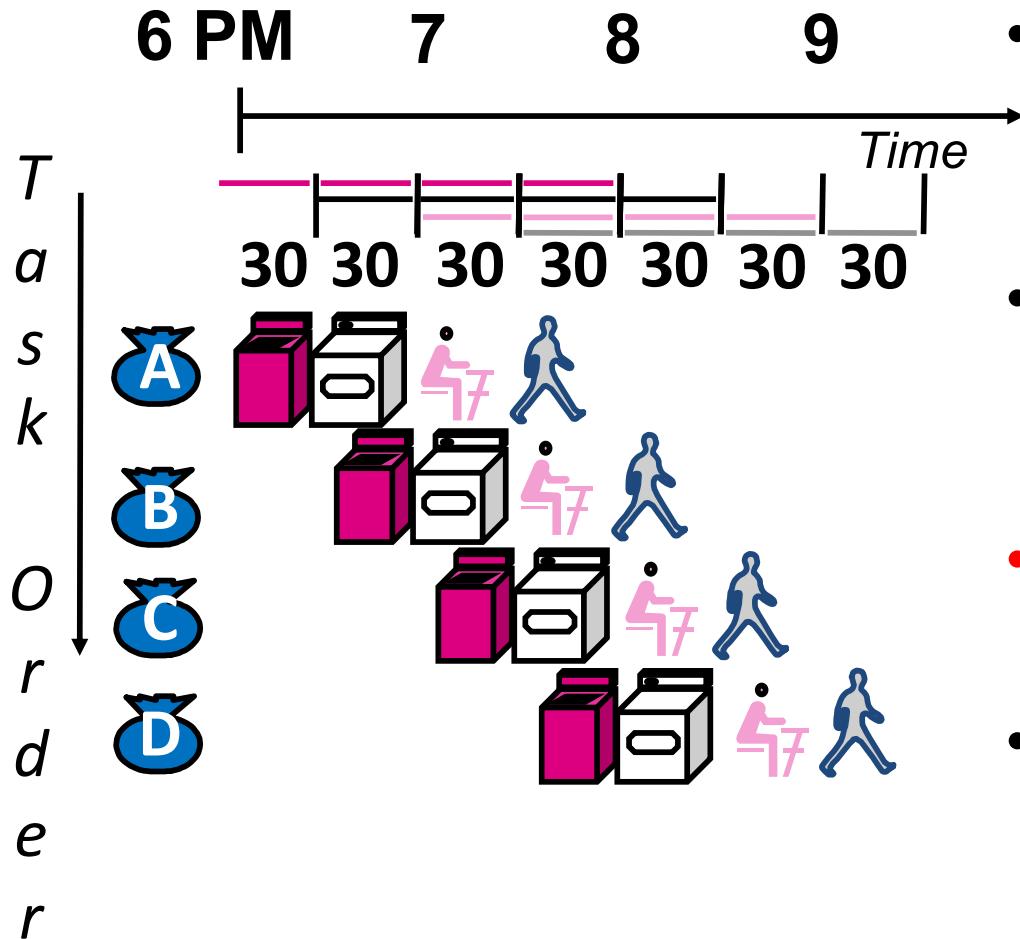
- Sequential laundry takes 8 hours for 4 loads

# Pipelined Laundry



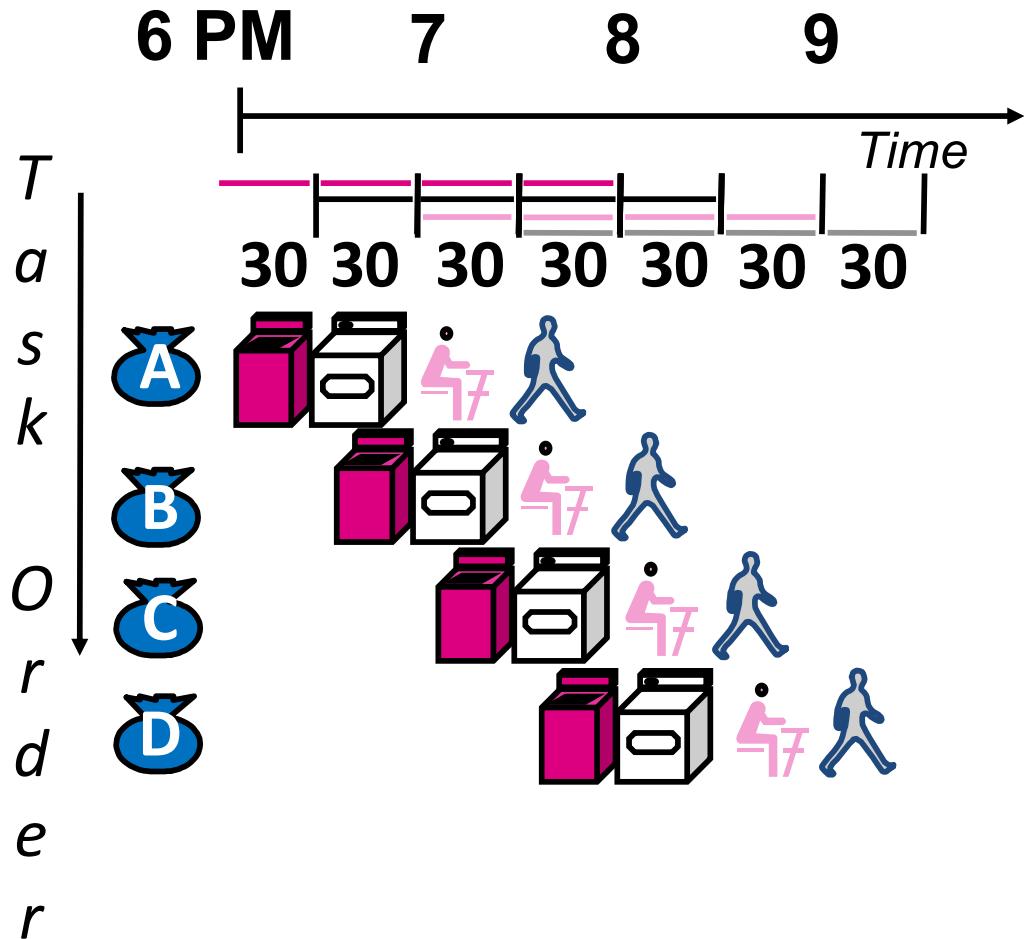
- Pipelined laundry takes 3.5 hours for 4 loads!

# Pipelining Lessons (1/2)



- Pipelining doesn't help *latency* of single task, just *throughput* of entire workload
- *Multiple* tasks operating simultaneously using different resources
- **Potential speedup = number of pipeline stages**
- Speedup reduced by time to *fill* and *drain* the pipeline: 8 hours/3.5 hours or 2.3X v. potential 4X in this example

# Pipelining Lessons (2/2)

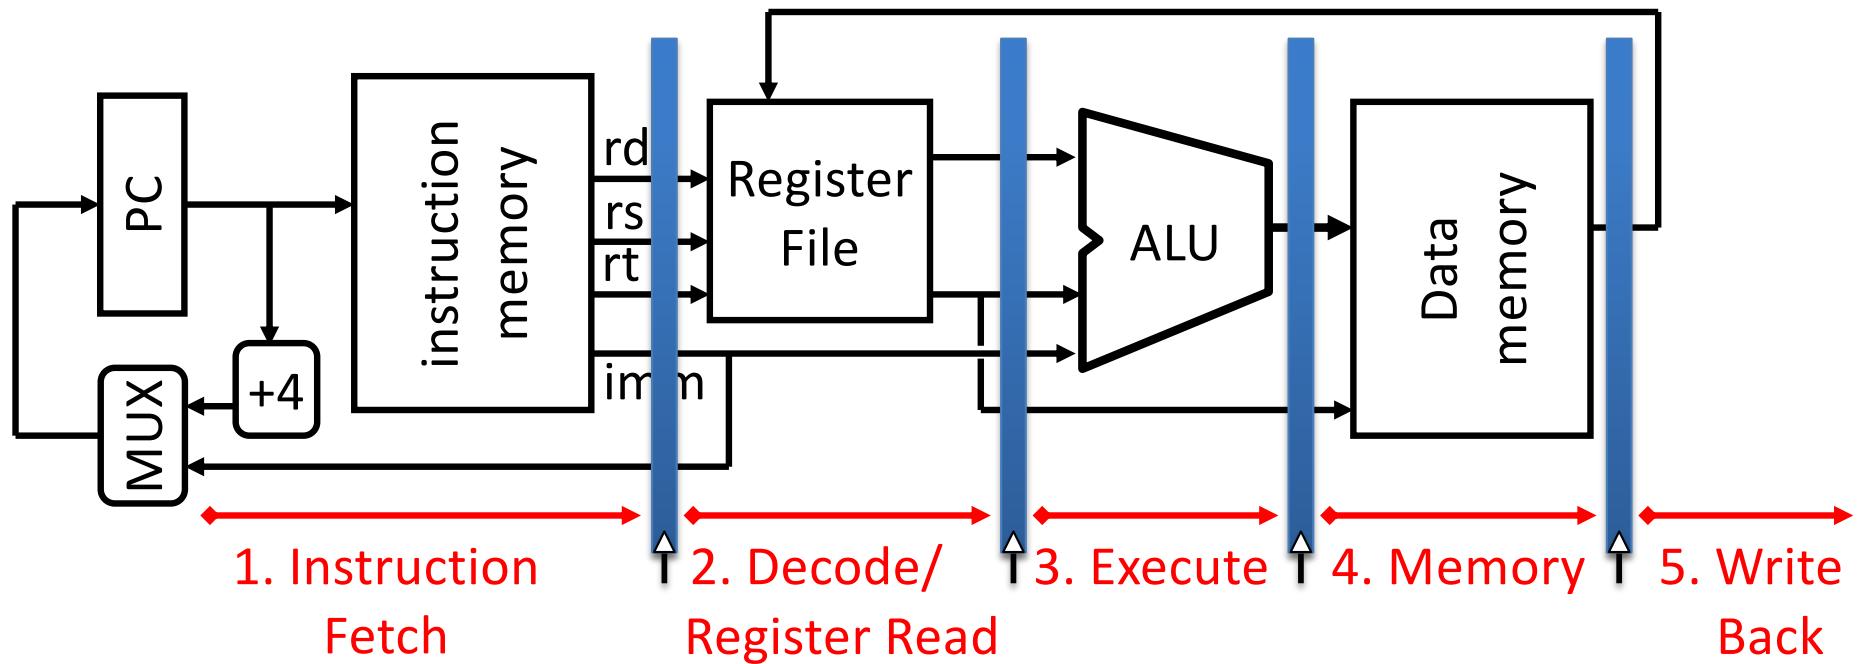


- Suppose new Washer takes 20 minutes, new Stasher takes 20 minutes. How much faster is pipeline?
  - Pipeline rate limited by *slowest* pipeline stage
  - Unbalanced lengths of pipeline stages reduces speedup

# Recall: 5 Stages of MIPS Datapath

- 1) IF: Instruction Fetch, Increment PC
- 2) ID: Instruction Decode, Read Registers
- 3) EX: Execution (ALU)
  - Load/Store: Calculate Address
  - Others: Perform Operation
- 4) MEM:
  - Load: Read Data from Memory
  - Store: Write Data to Memory
- 5) WB: Write Data Back to Register

# Pipelined Datapath



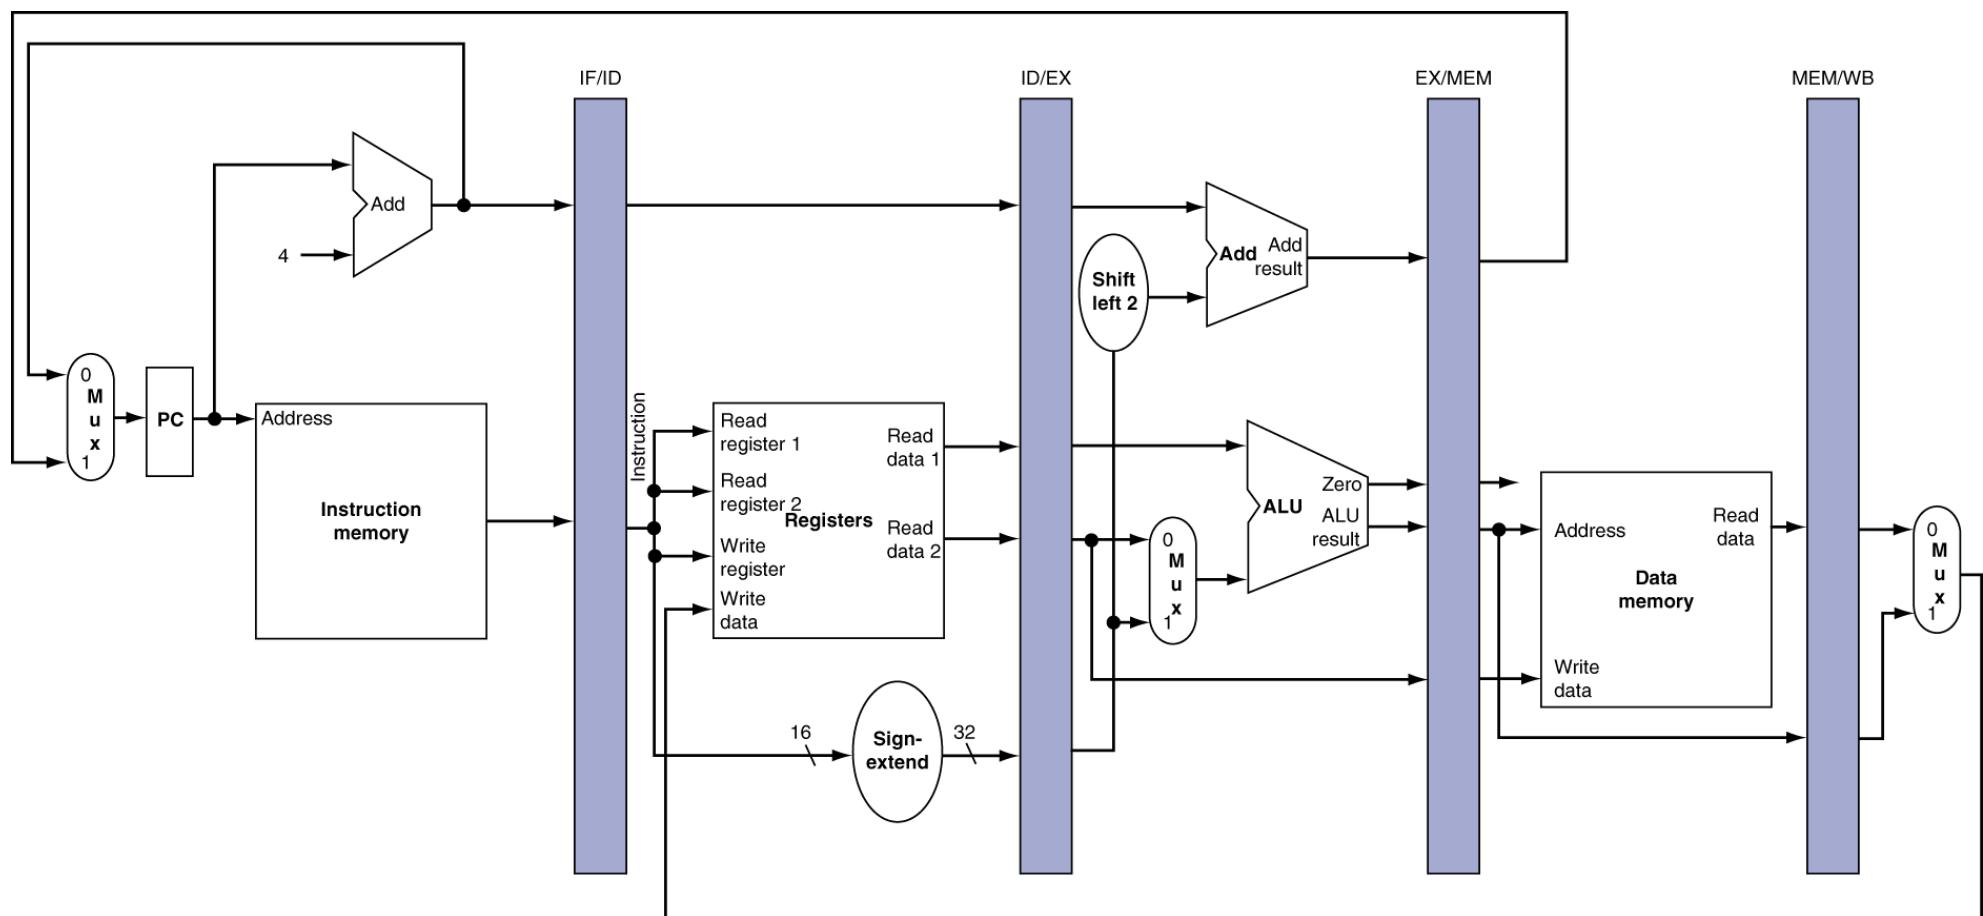
- Add registers between stages
  - Hold information produced in previous cycle
- 5 stage pipeline
  - Clock rate *potentially* 5x faster

# Pipelining Changes

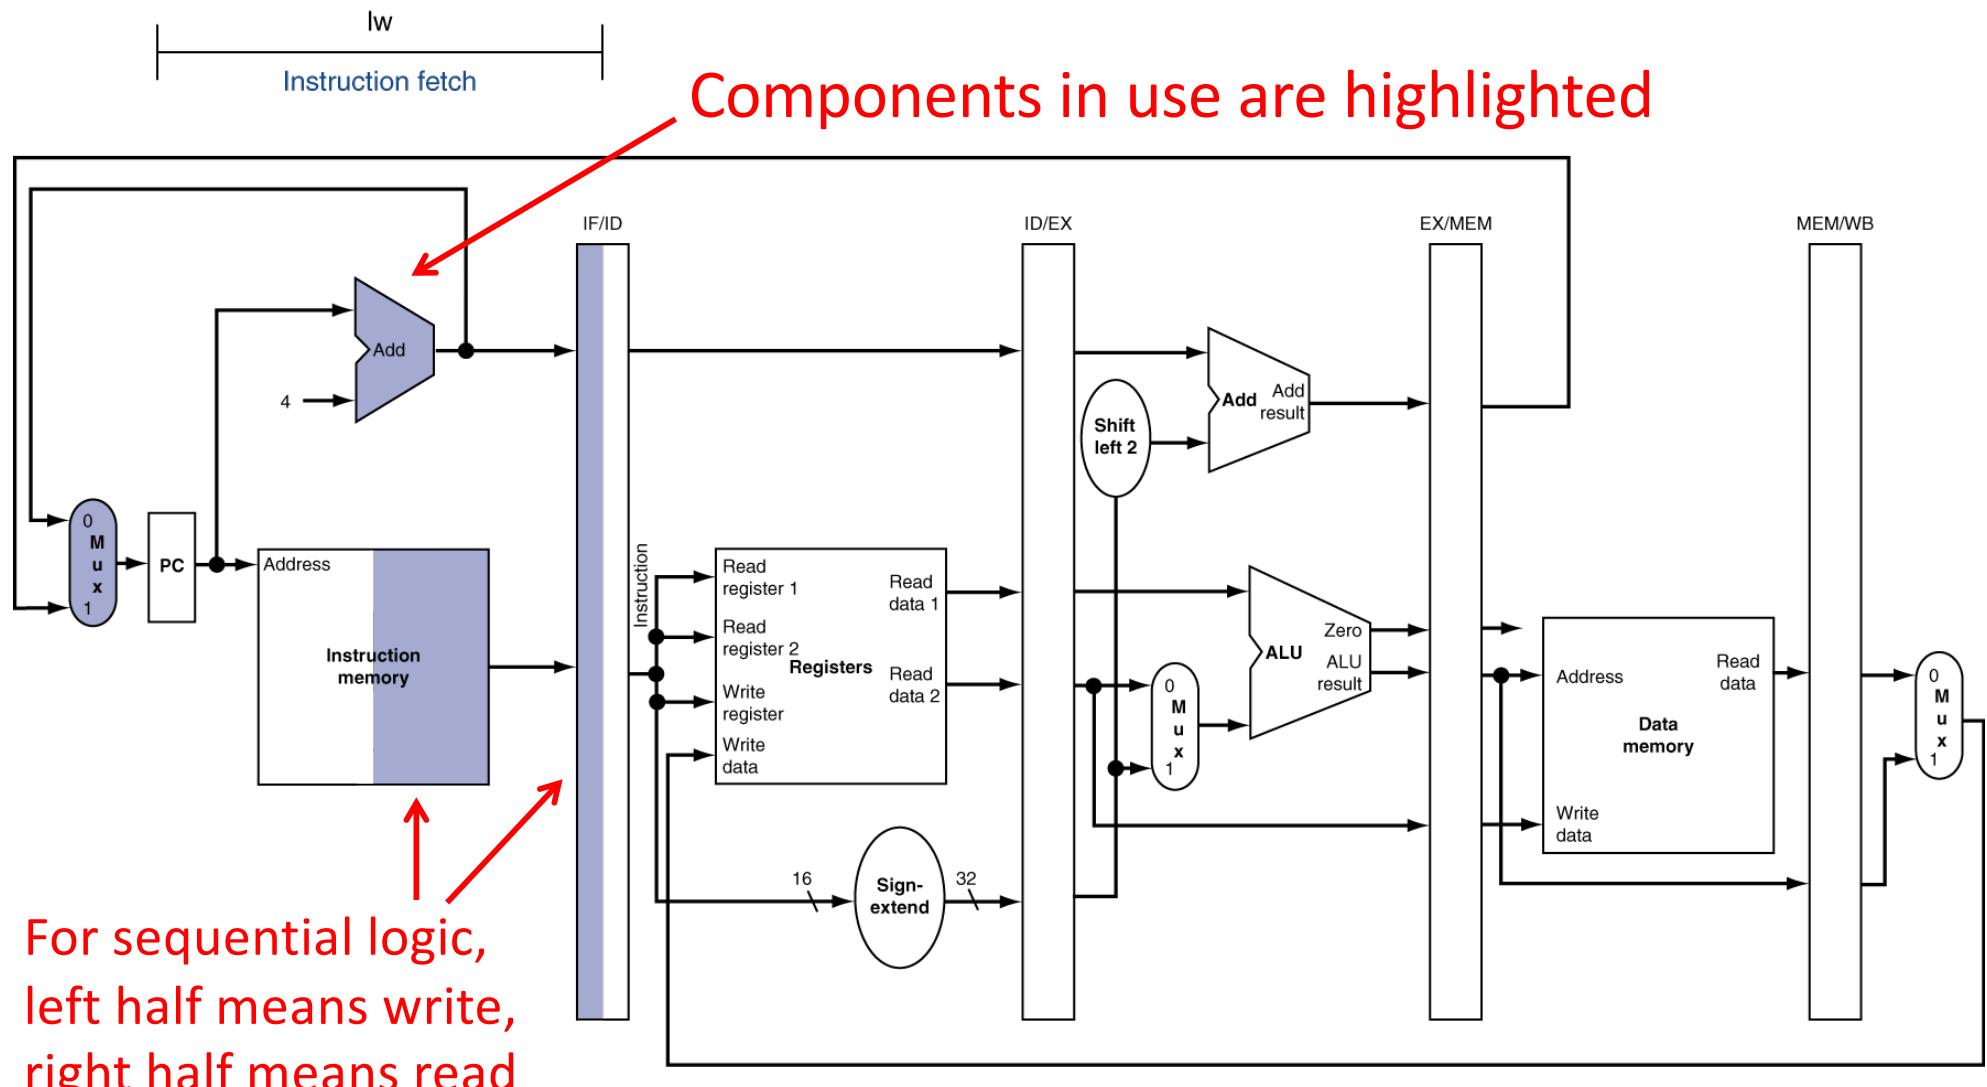
- Registers affect flow of information
  - Name registers for adjacent stages (e.g. IF/ID)
  - Registers *separate* the information between stages
  - At any instance of time, each stage working on a *different* instruction!
- Will need to re-examine placement of wires and hardware in datapath

# More Detailed Pipeline

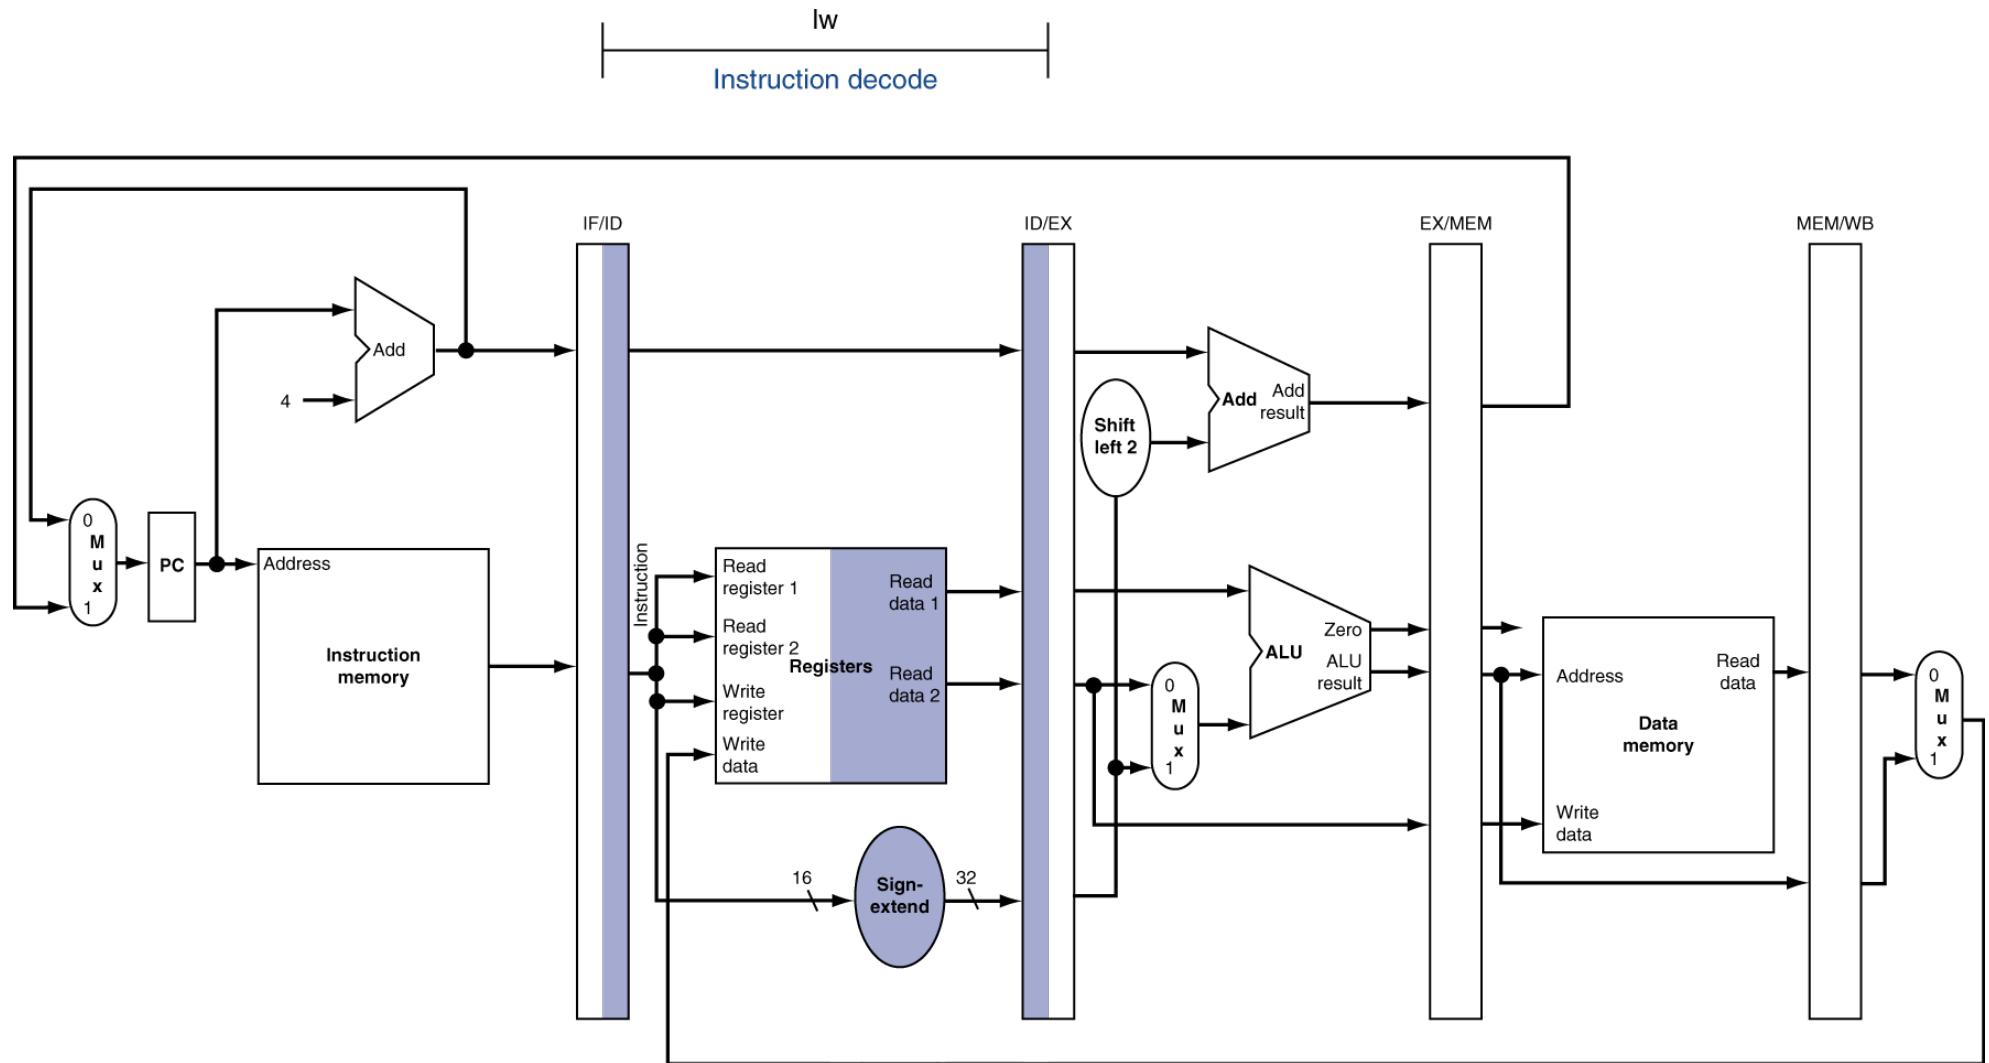
- Examine flow through pipeline for  $l_w$



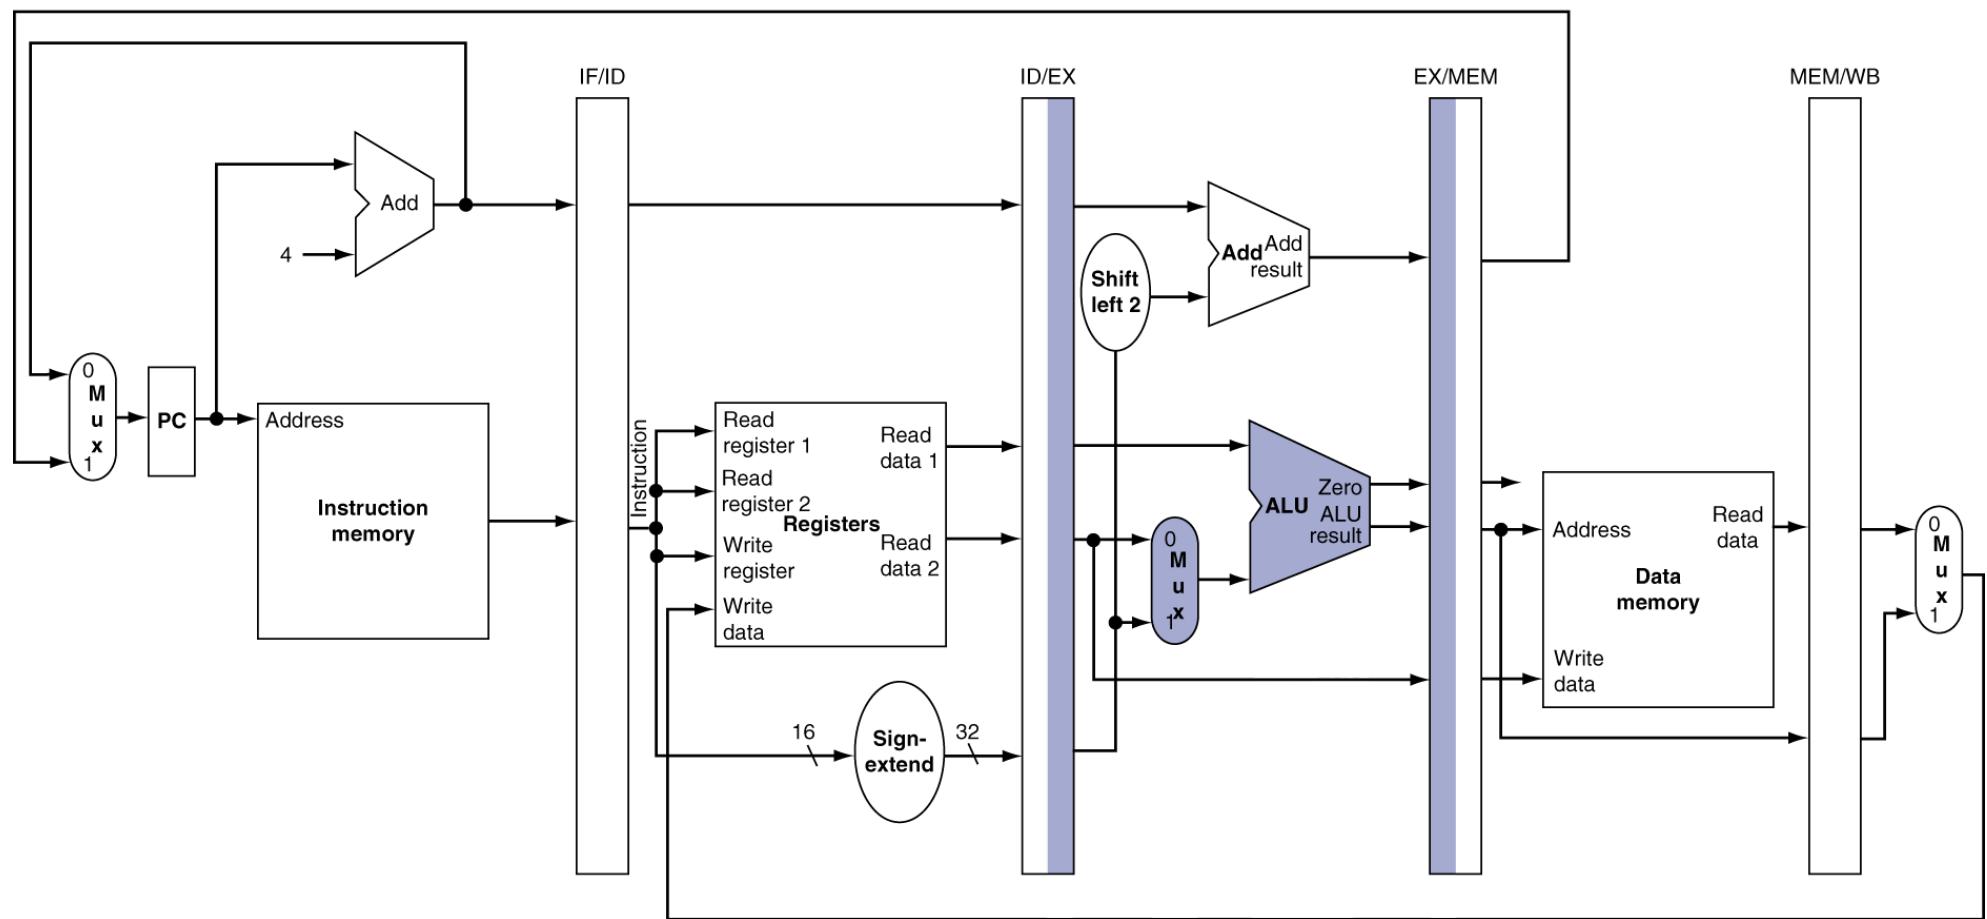
# Instruction Fetch (IF) for Load



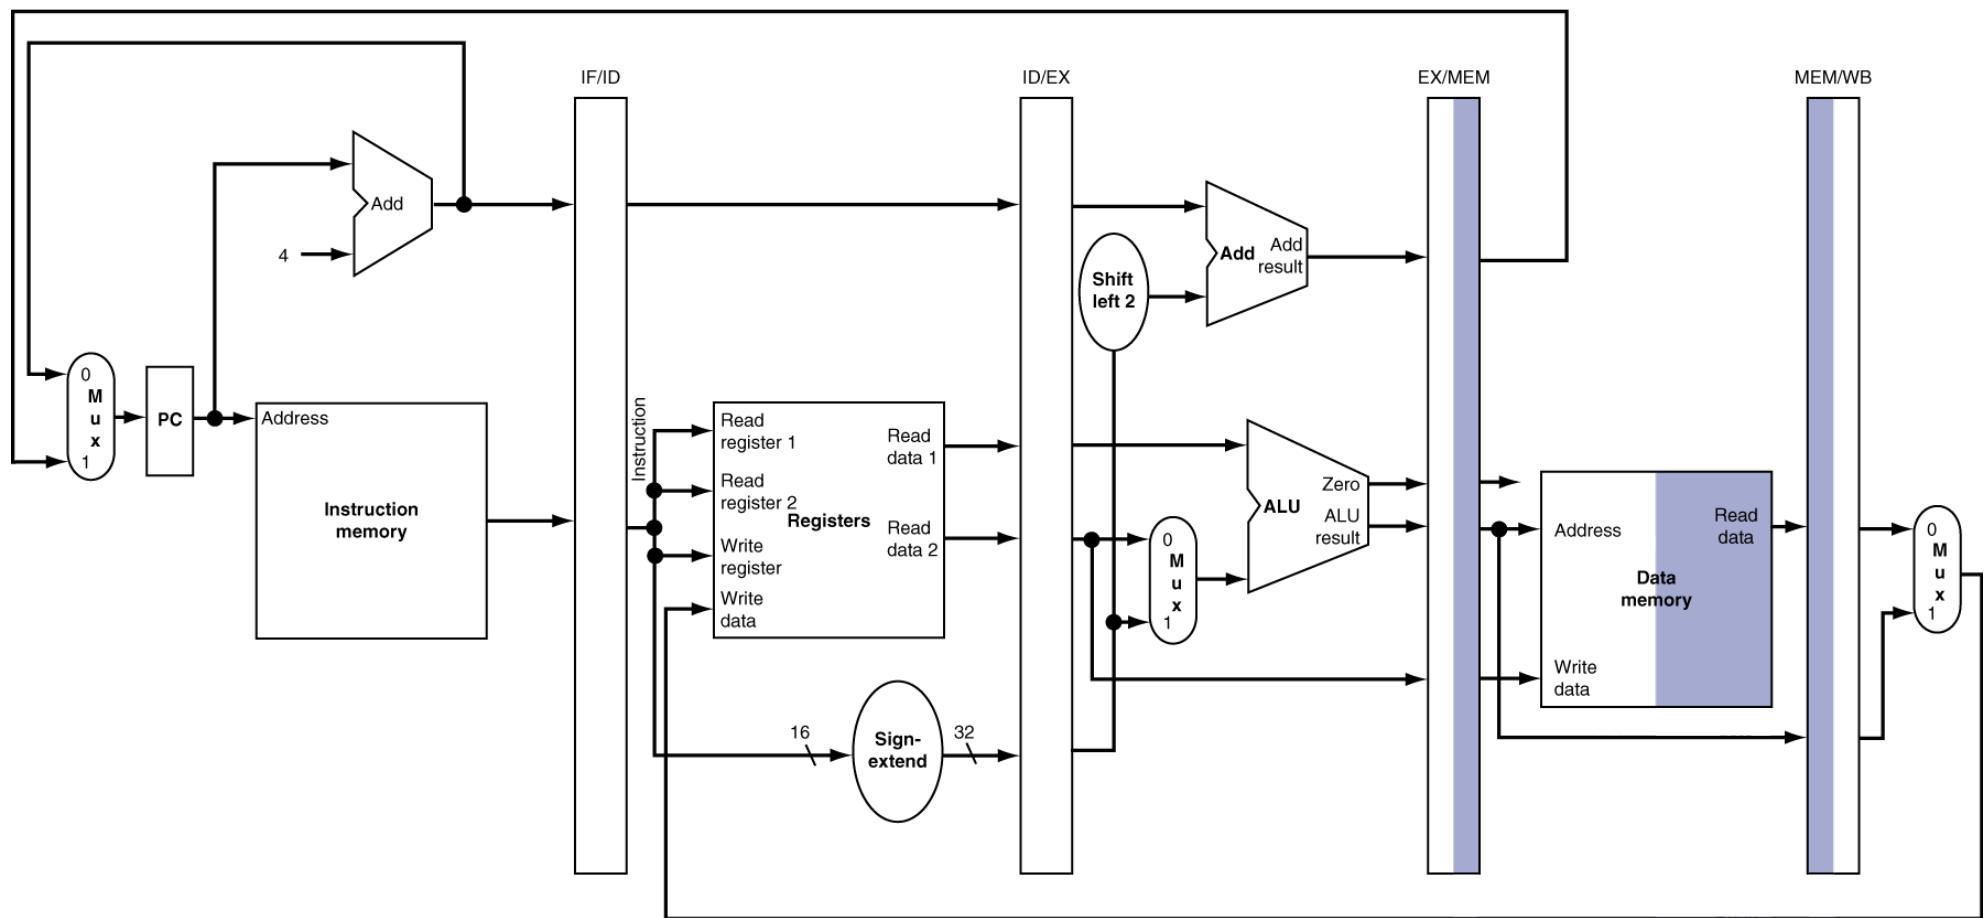
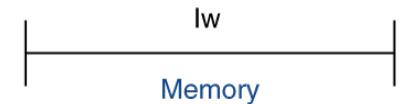
# Instruction Decode (ID) for Load



# Execute (EX) for Load



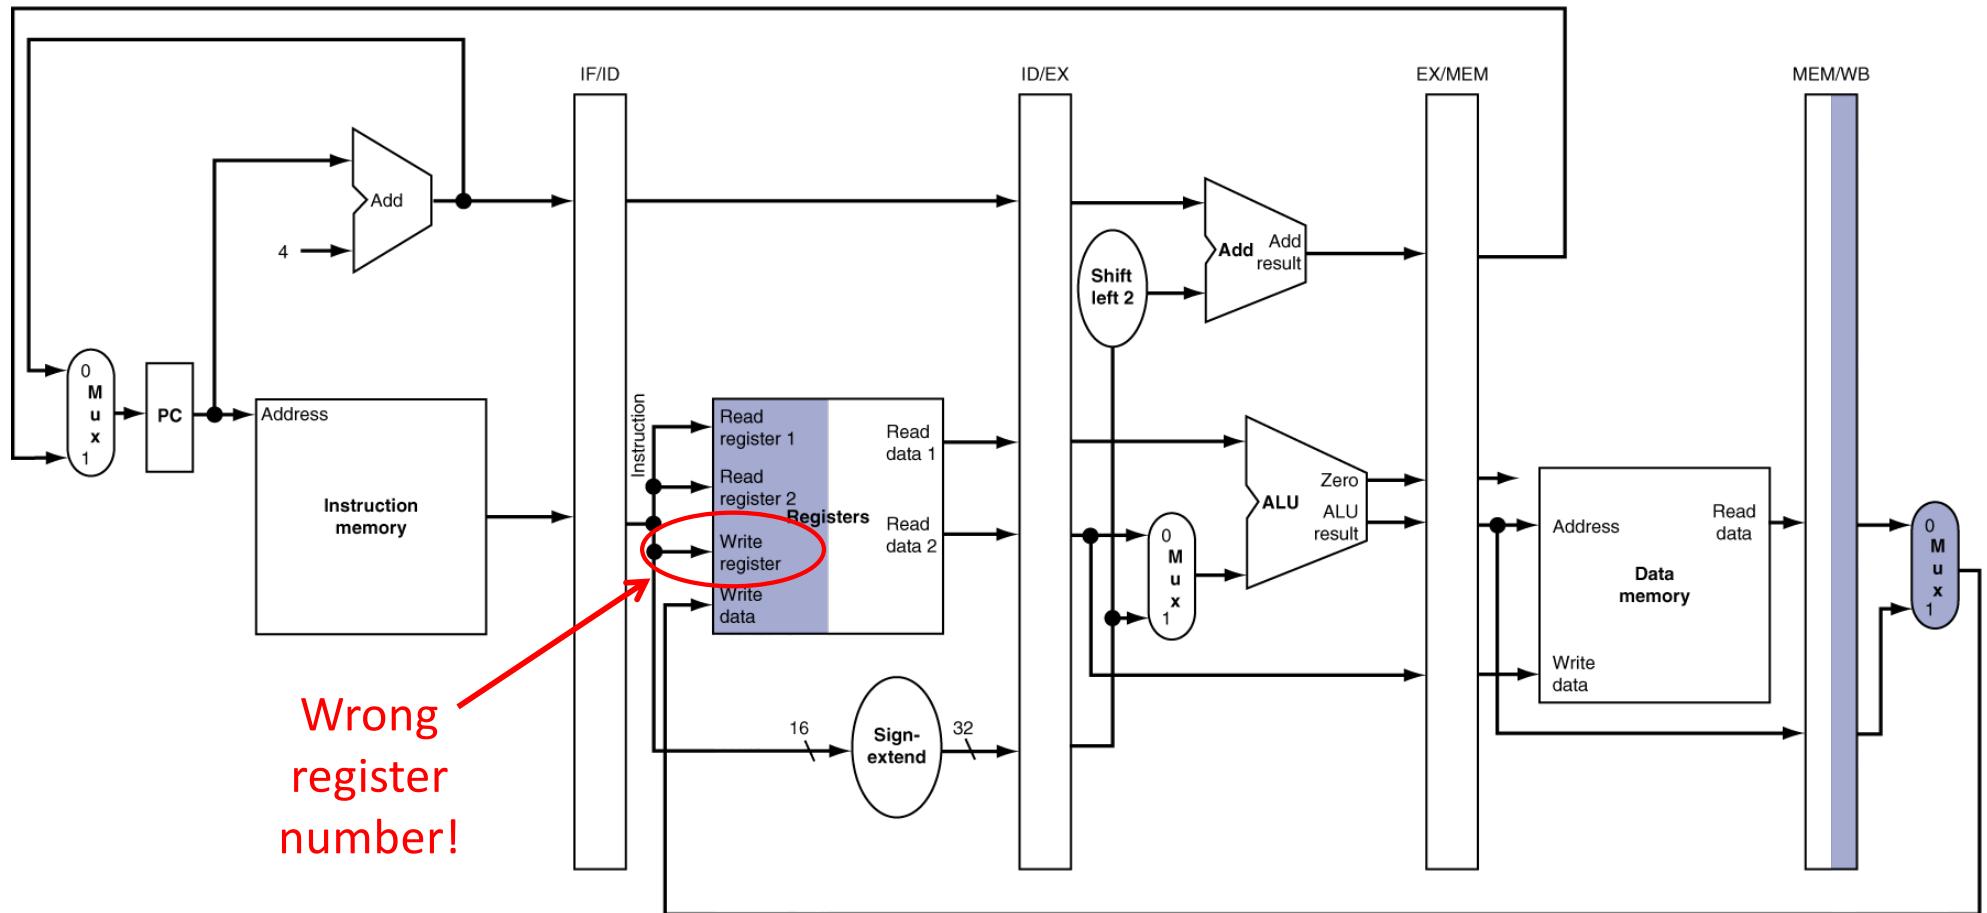
# Memory (MEM) for Load



# Write Back (WB) for Load

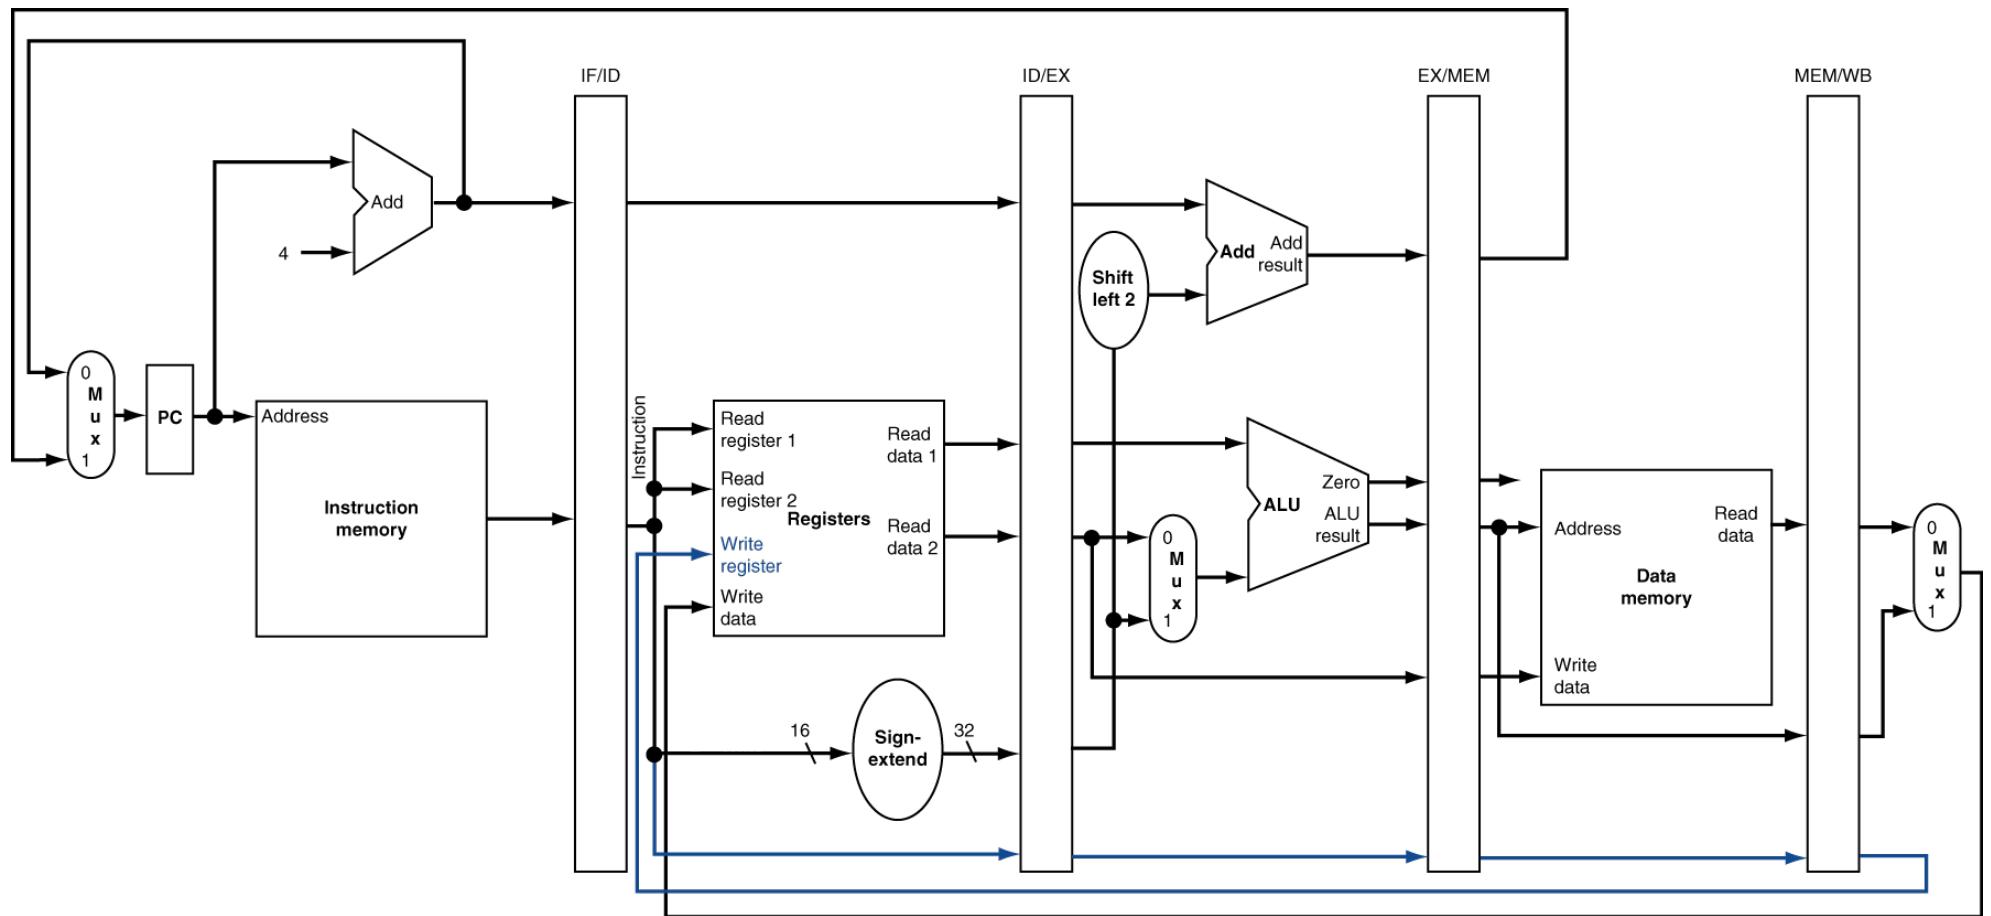
There's something wrong here! (Can you spot it?)

lw  
Write back

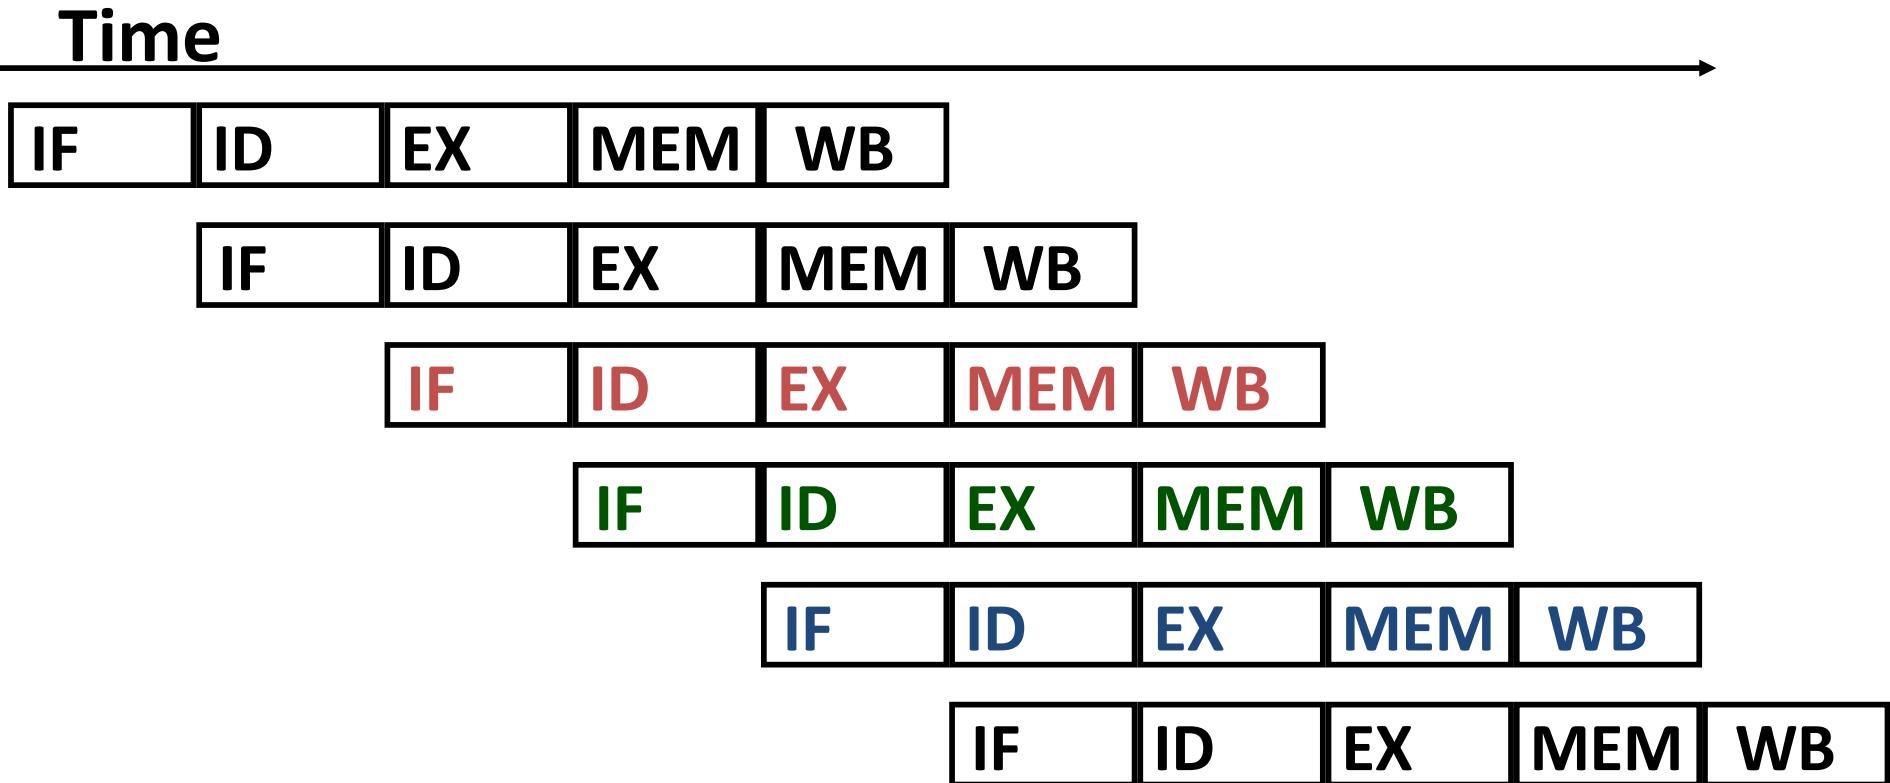


# Corrected Datapath

- Now any instruction that writes to a register will work properly

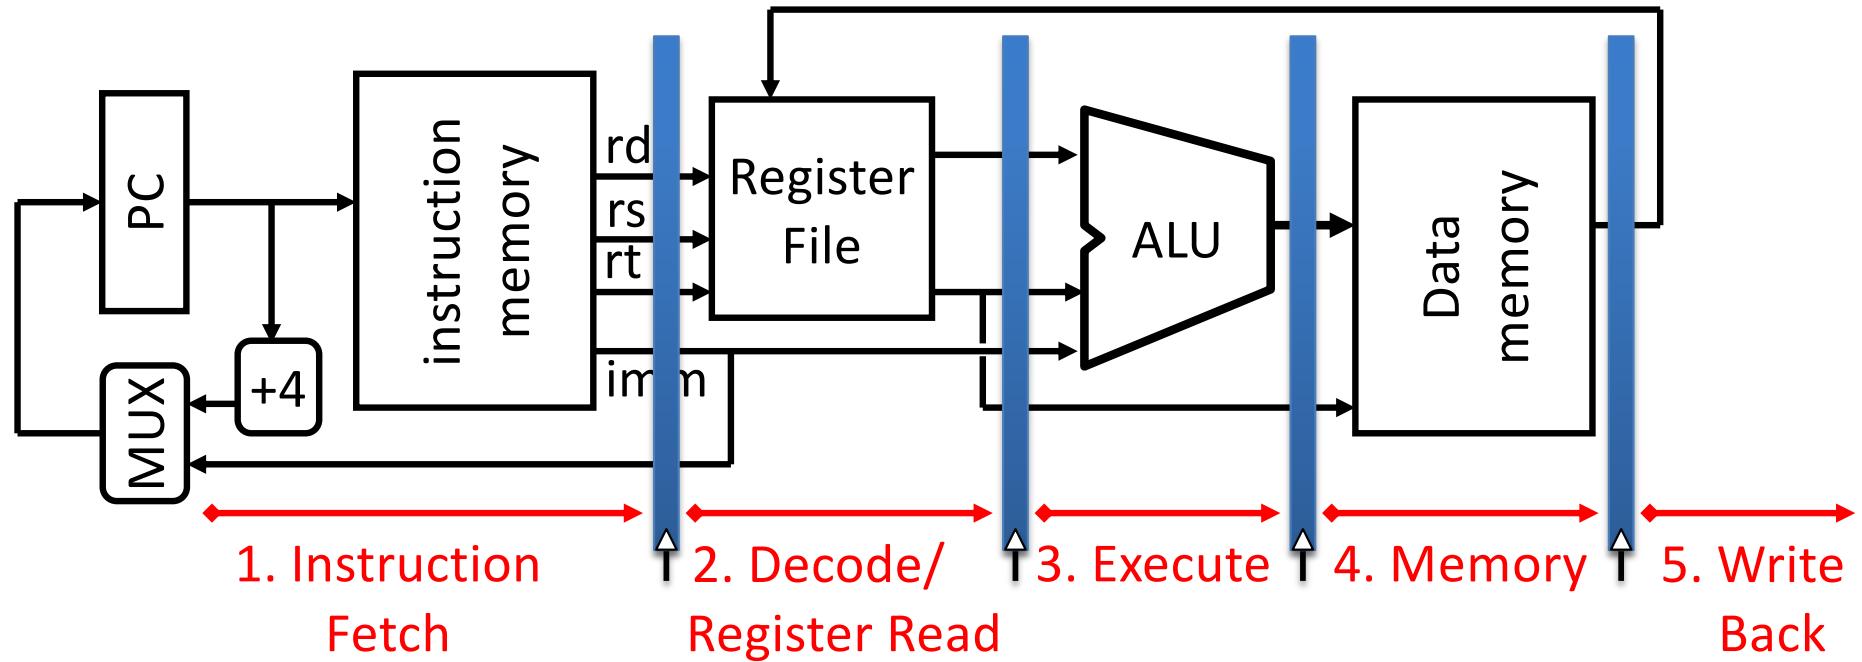


# Pipelined Execution Representation

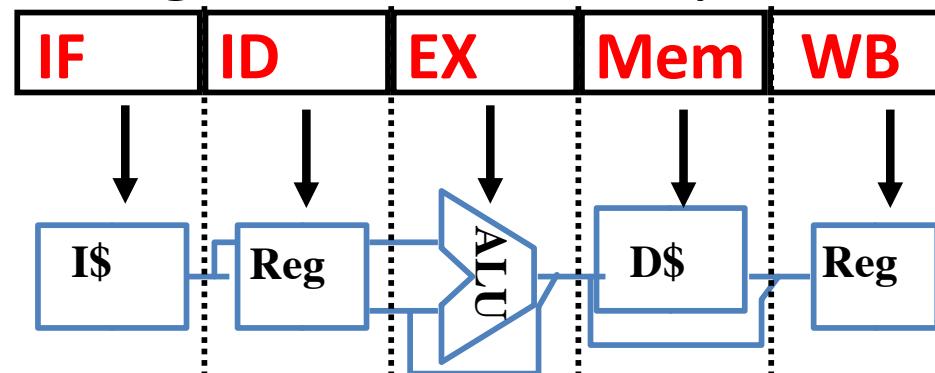


- Every instruction must take same number of steps, so some will idle
  - e.g. MEM stage for any arithmetic instruction

# Graphical Pipeline Diagrams

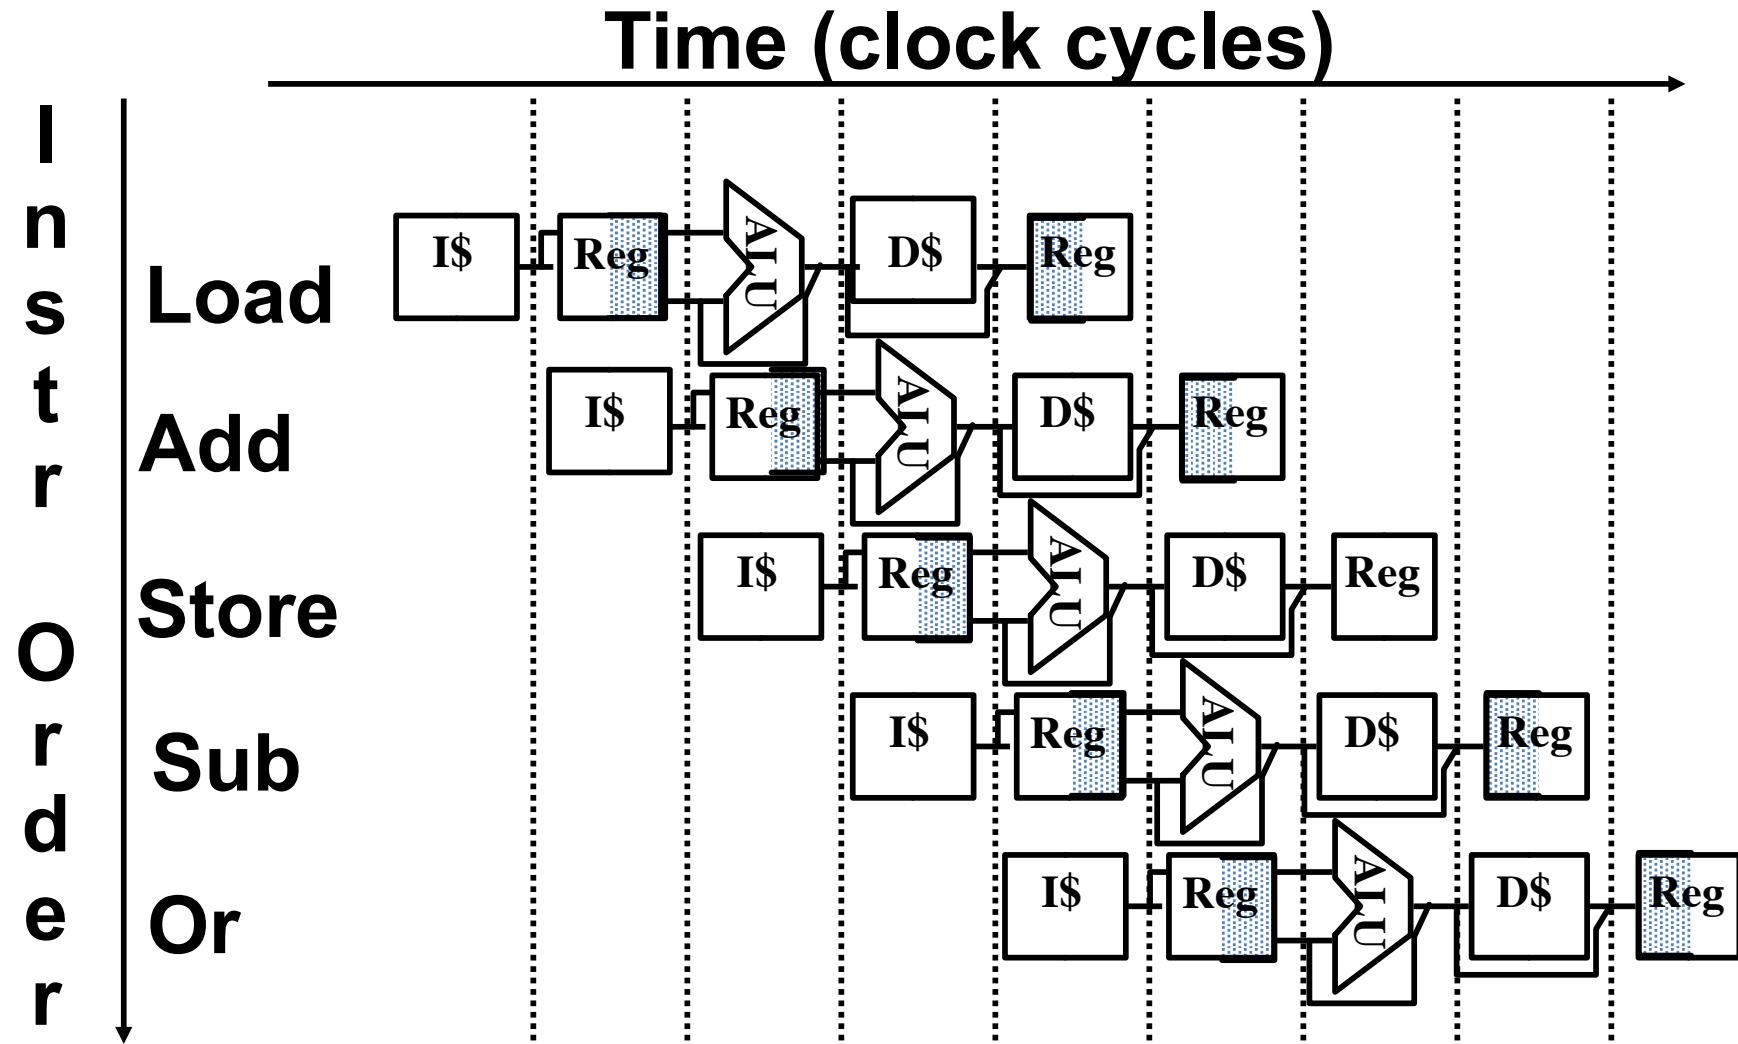


- Use datapath figure below to represent pipeline:



# Graphical Pipeline Representation

- RegFile: right half is read, left half is write



# Instruction Level Parallelism (ILP)

- Pipelining allows us to execute parts of multiple instructions at the same time using the same hardware!
  - This is known as *instruction level parallelism*
- **Recall:** Types of parallelism
  - DLP: same operation on lots of data (SIMD)
  - TLP: executing multiple threads “simultaneously” (OpenMP)

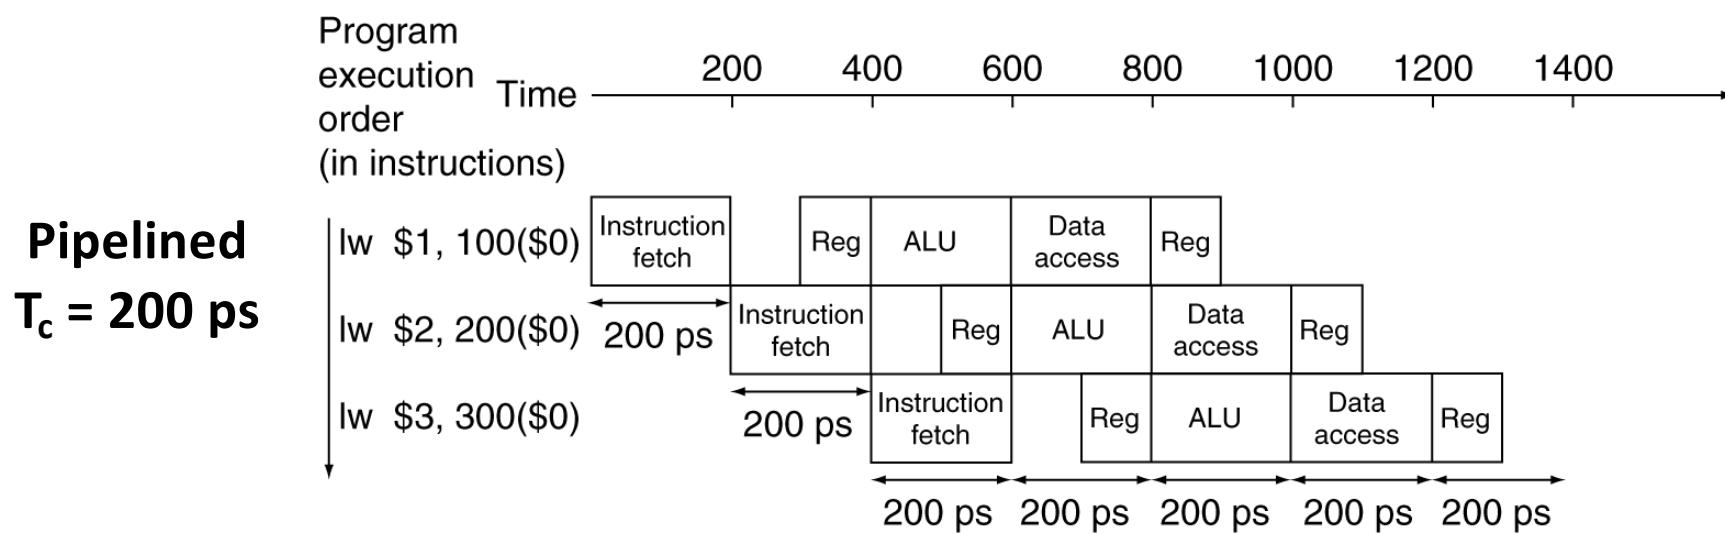
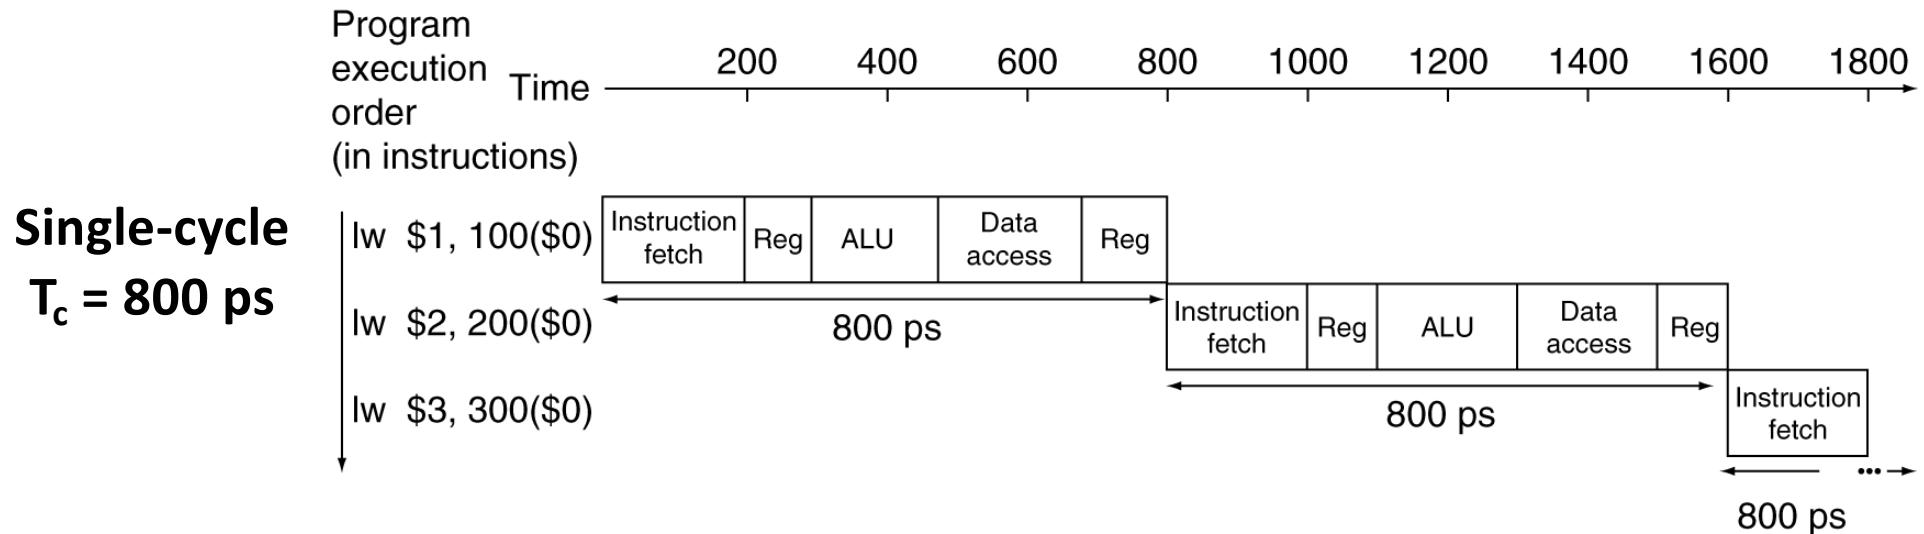
# Pipeline Performance (1/2)

- Assume time for stages is
  - 100ps for register read or write
  - 200ps for other stages

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

- What is pipelined clock rate?
  - Compare pipelined datapath with single-cycle datapath

# Pipeline Performance (2/2)



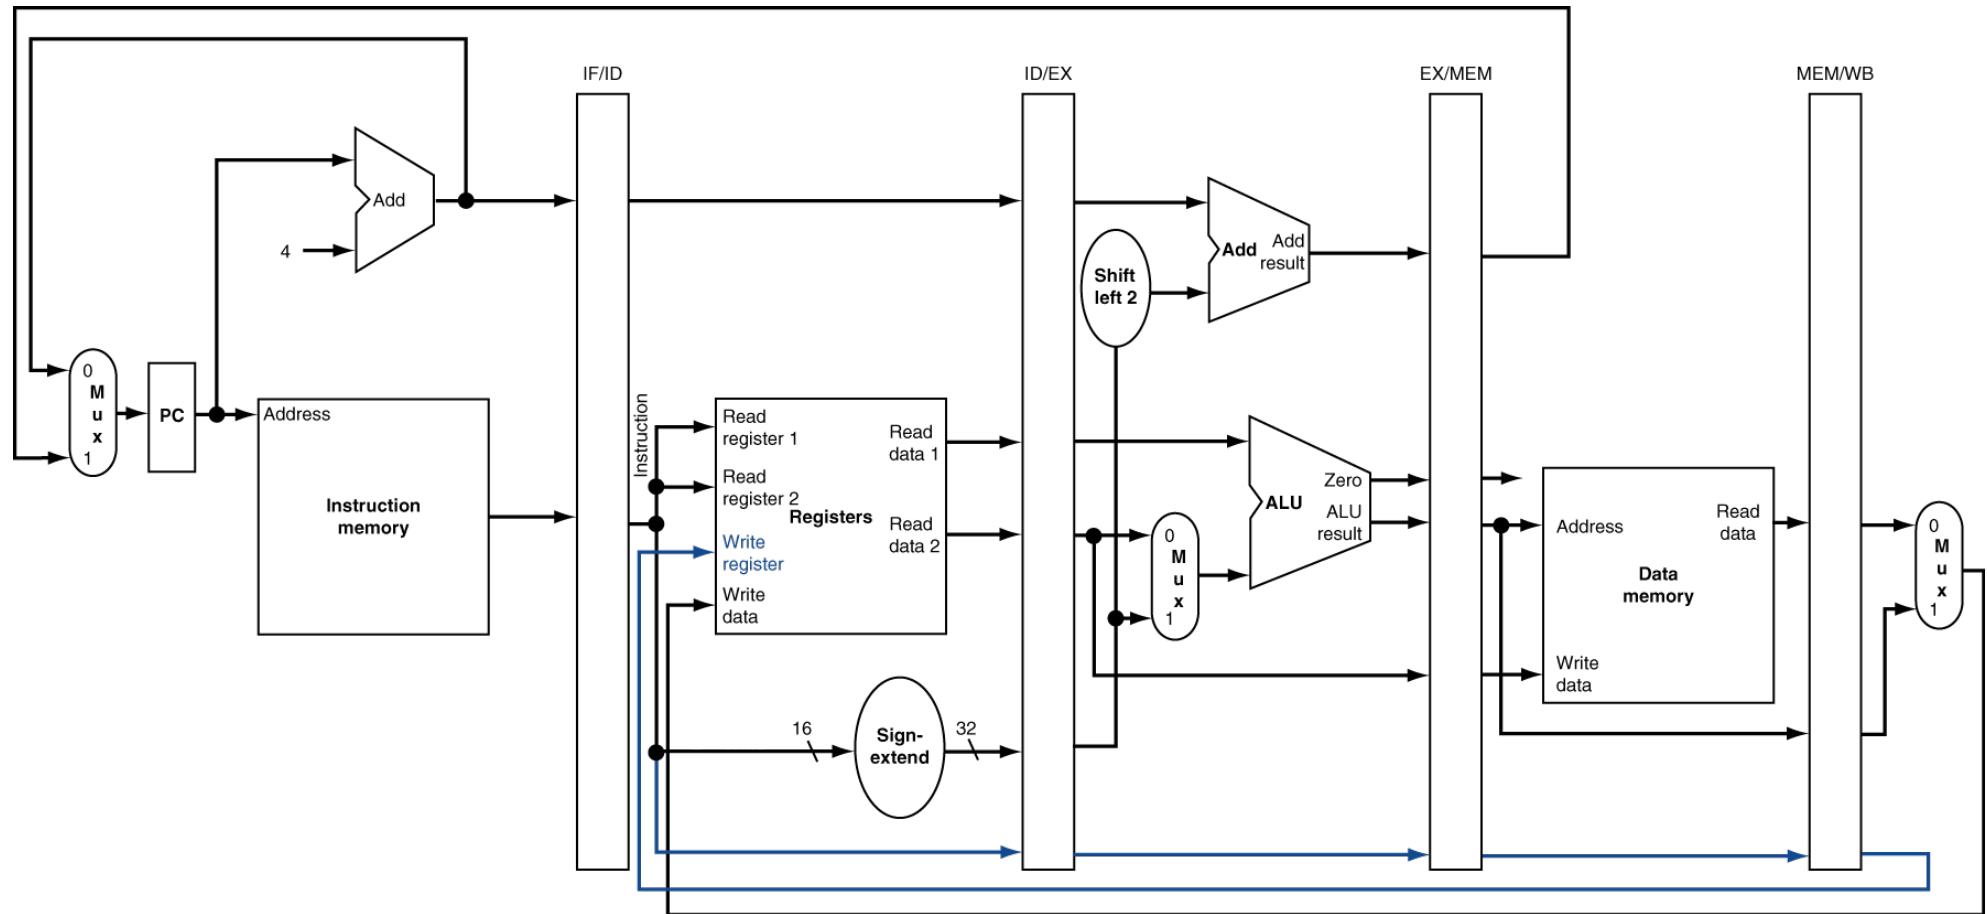
# Pipeline Speedup

- Use  $T_c$  (“time between completion of instructions”) to measure speedup
  - $T_{c,\text{pipelined}} \geq \frac{T_{c,\text{single-cycle}}}{\text{Number of stages}}$
  - Equality only achieved if stages are *balanced* (i.e. take the same amount of time)
- If not balanced, speedup is reduced
- Speedup due to increased throughput
  - Latency for each instruction does not decrease

# Pipelining and ISA Design

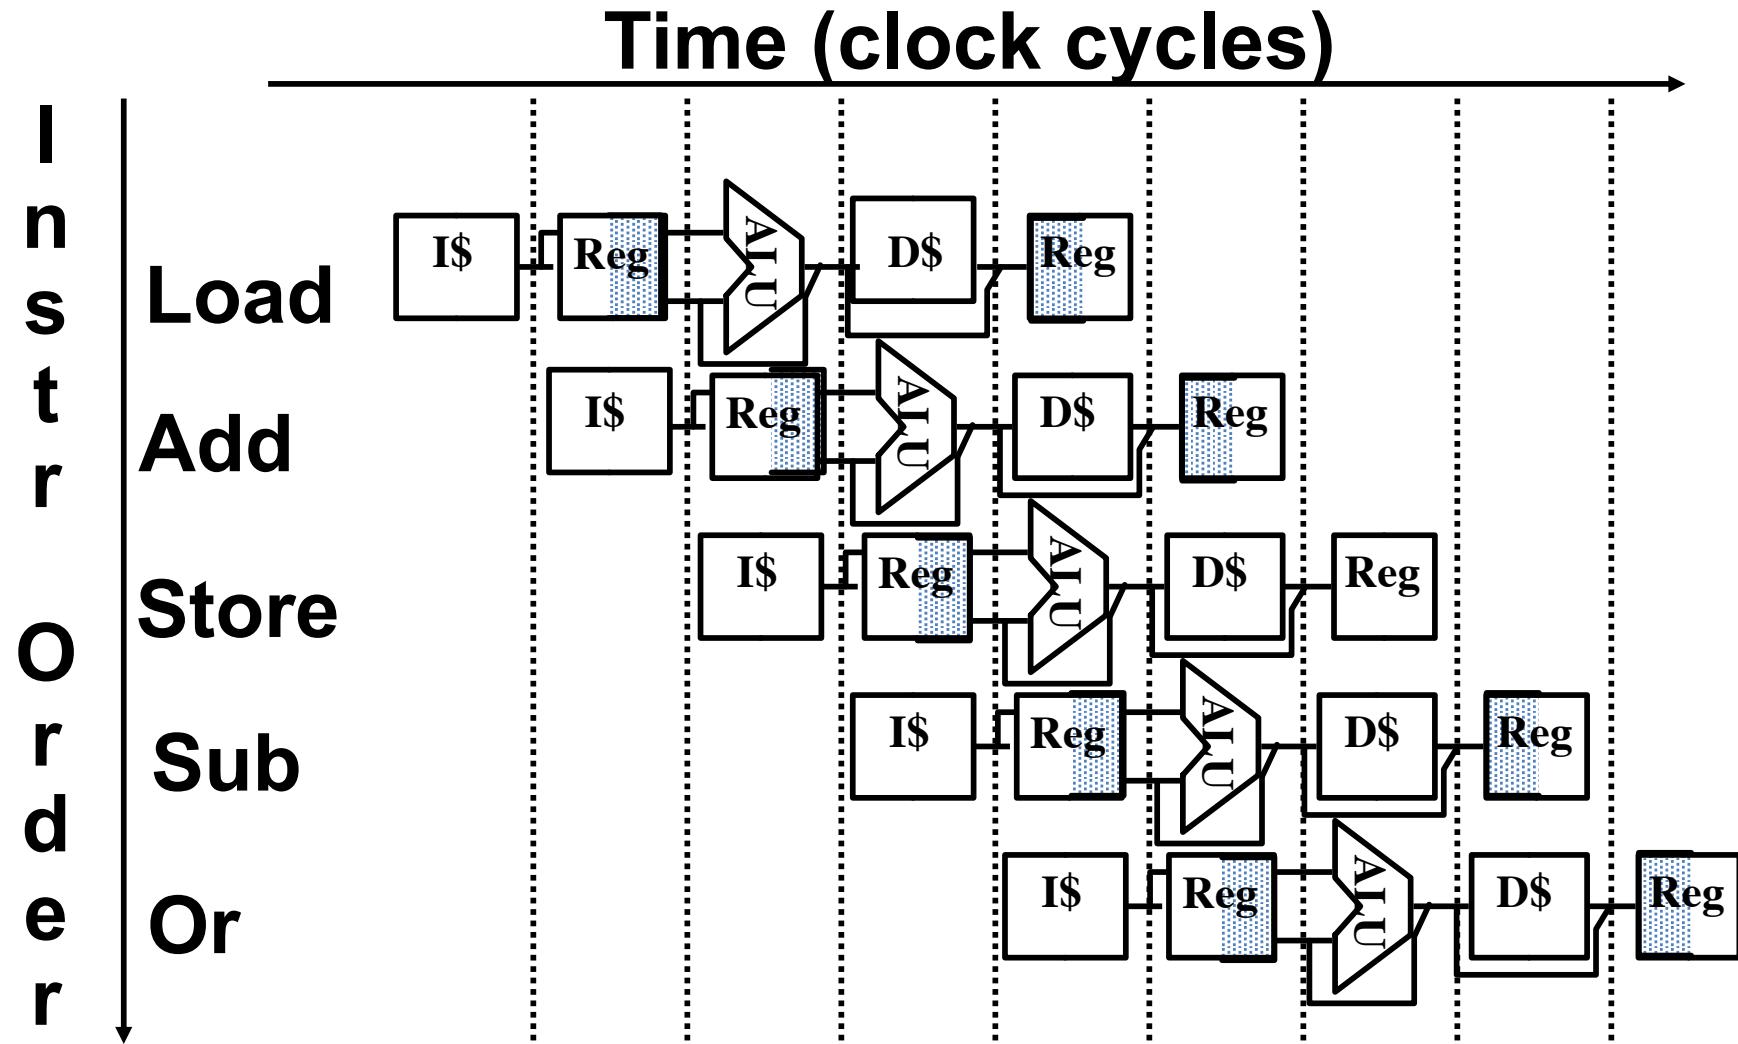
- MIPS Instruction Set designed for pipelining!
- All instructions are 32-bits
  - Easier to fetch and decode in one cycle
- Few and regular instruction formats, 2 source register fields always in same place
  - Can decode and read registers in one step
- Memory operands only in Loads and Stores
  - Can calculate address 3<sup>rd</sup> stage, access memory 4<sup>th</sup> stage
- Alignment of memory operands
  - Memory access takes only one cycle

# Review: Pipelined Datapath



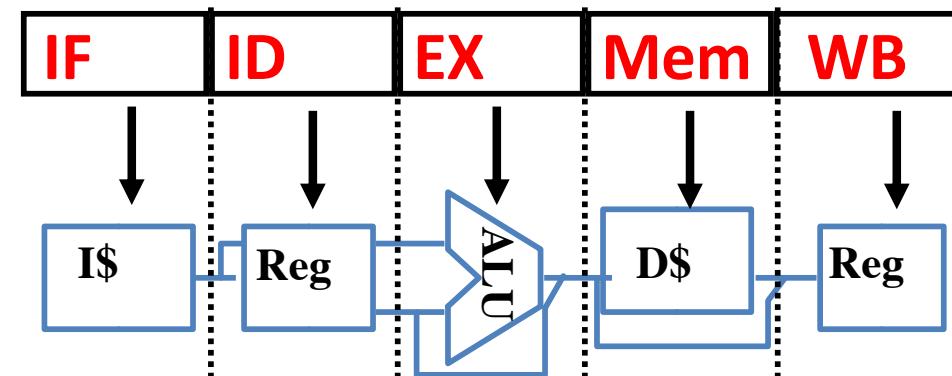
# Graphical Pipeline Representation

- RegFile: right half is read, left half is write



**Question:** Which of the following signals (buses or control signals) for MIPS-lite does NOT need to be passed into the EX pipeline stage?

- PC + 4**
- MemWr**
- RegWr**
- imm16**



# Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

## 1) *Structural hazard*

- A required resource is busy  
(e.g. needed in multiple stages)

## 2) *Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

## 3) *Control hazard*

- Flow of execution depends on previous instruction

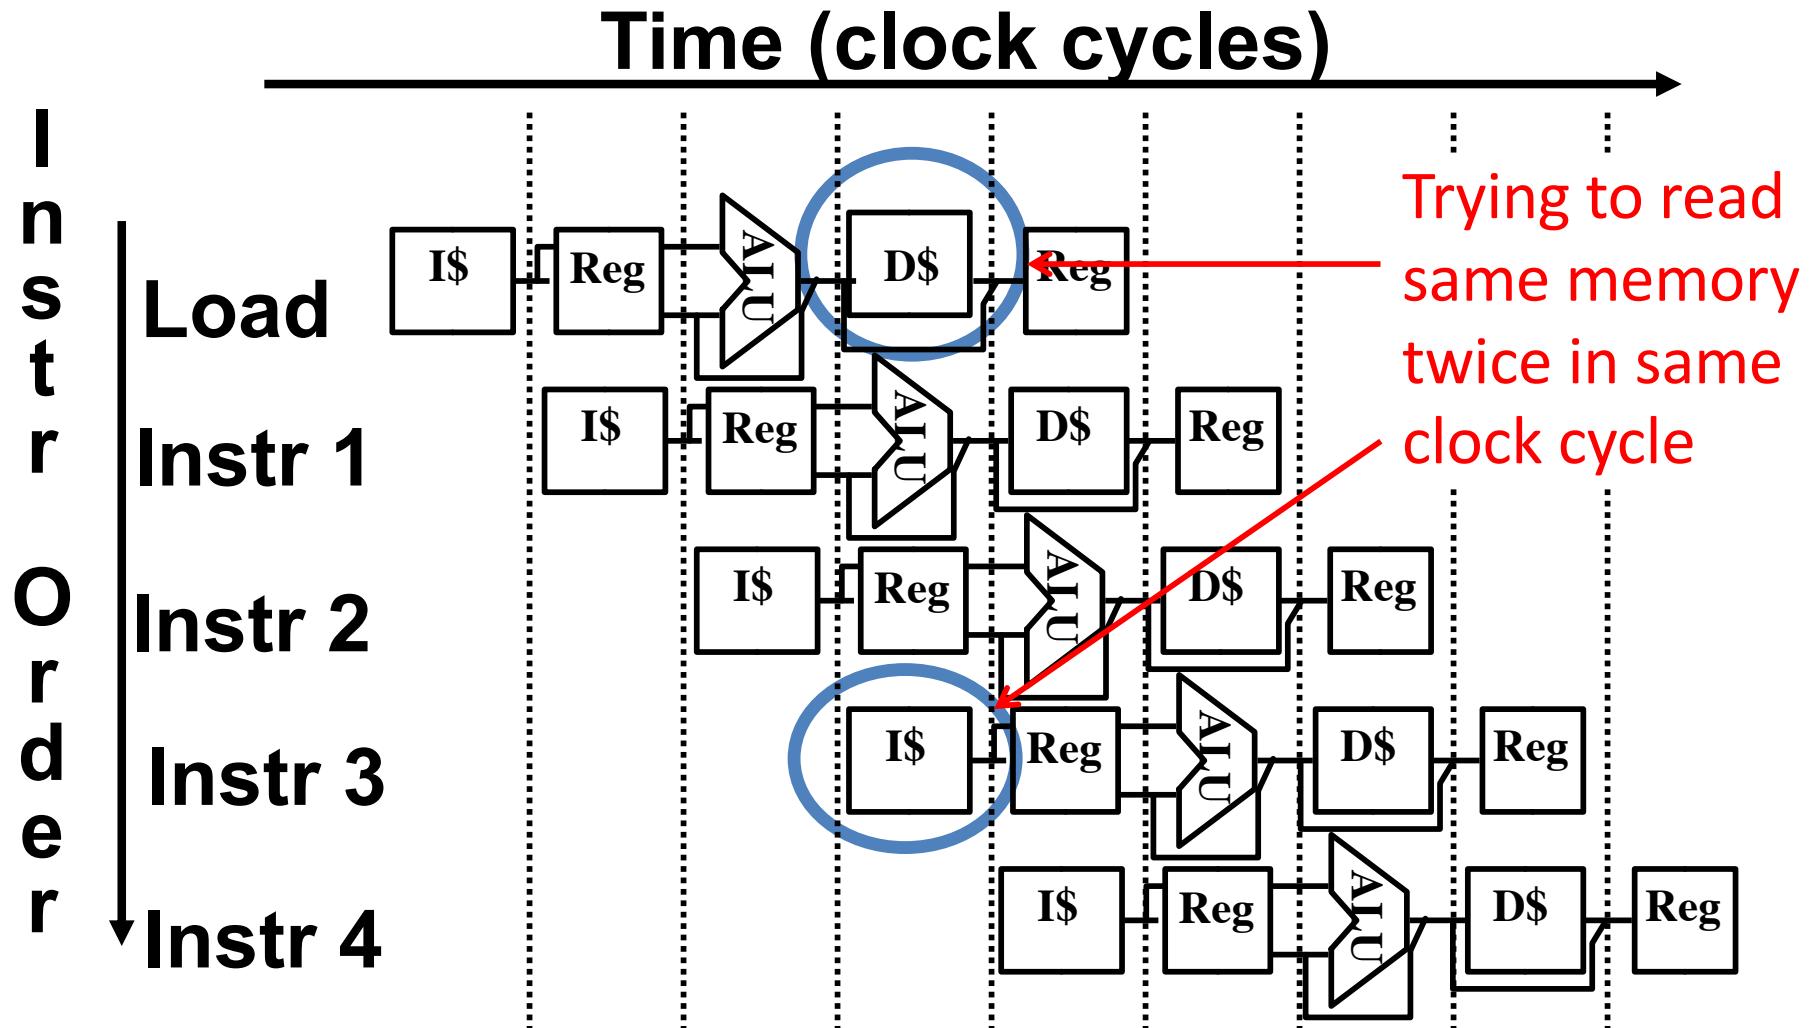
# Agenda

- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- Control Hazards
  - Branch and Jump Delay Slots
  - Branch Prediction

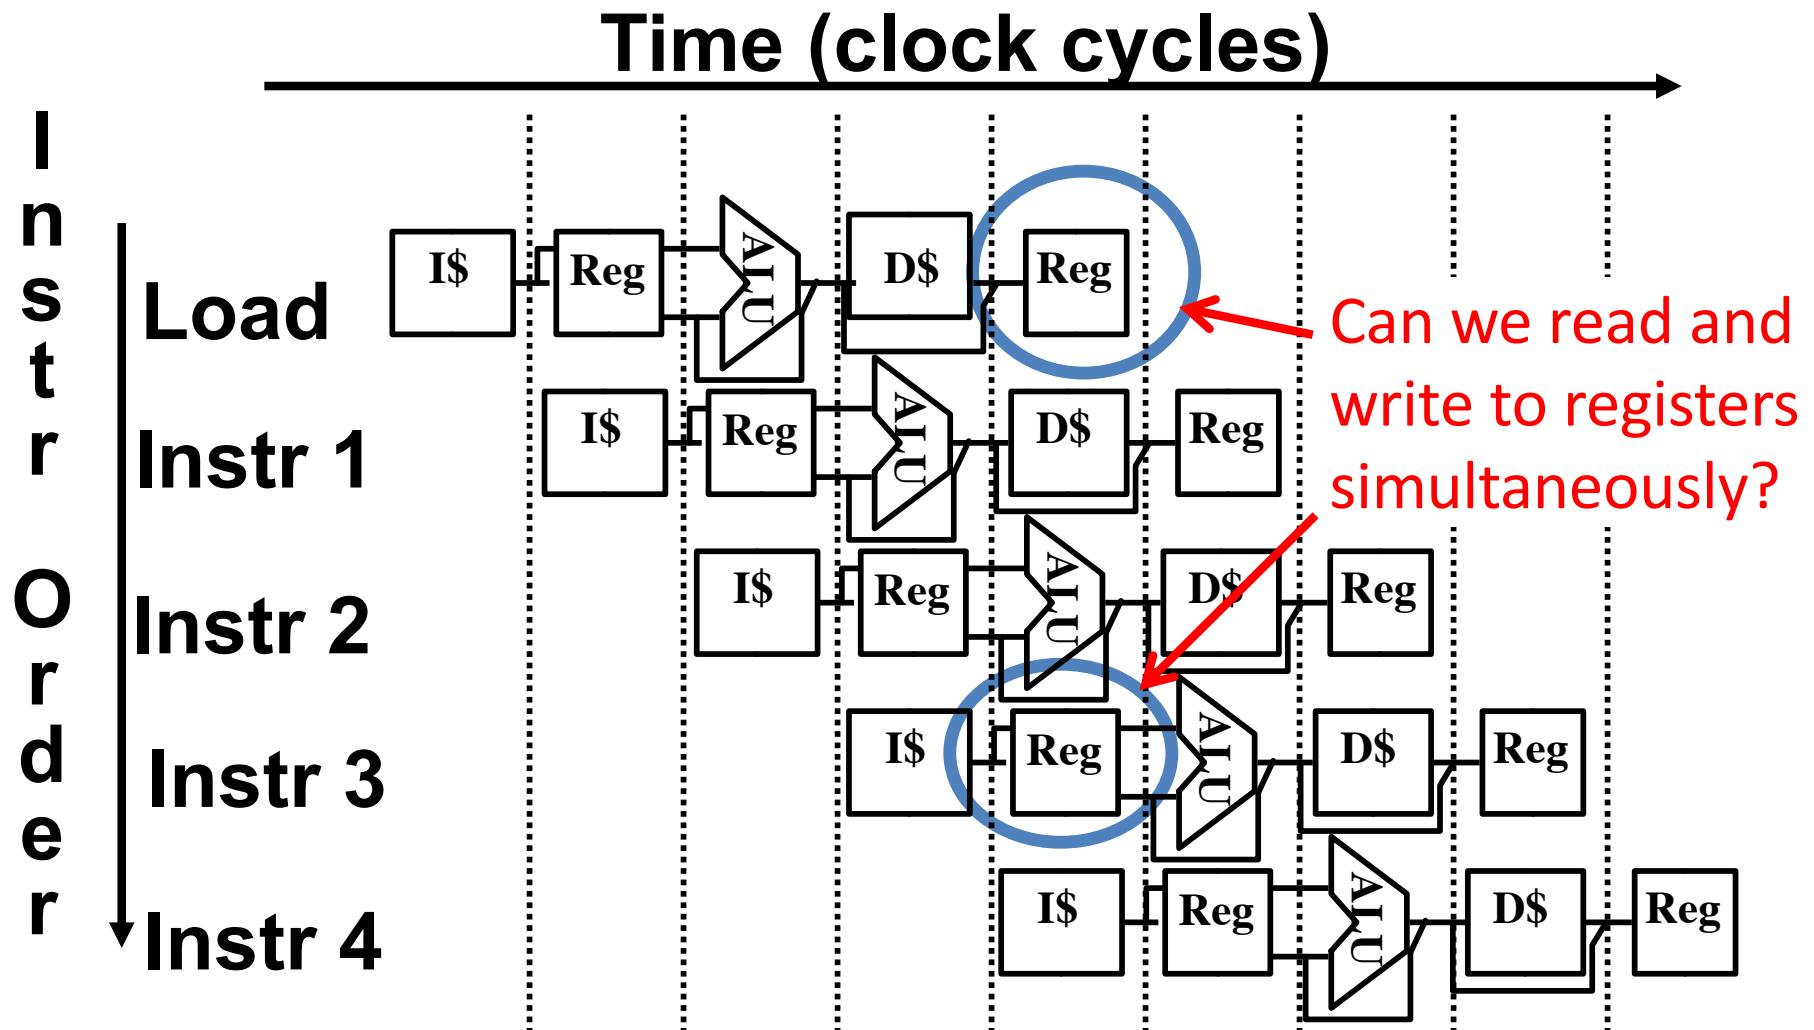
# 1. Structural Hazards

- Conflict for use of a resource
- MIPS pipeline with a single memory?
  - Load/Store requires memory access for data
  - Instruction fetch would have to *stall* for that cycle
    - Causes a pipeline “*bubble*”
- Hence, pipelined datapaths require separate instruction/data memories
  - Separate L1 I\$ and L1 D\$ take care of this

# Structural Hazard #1: Single Memory



# Structural Hazard #2: Registers (1/2)



## Structural Hazard #2: Registers (2/2)

- Two different solutions have been used:
  - 1) Split RegFile access in two: Write during 1<sup>st</sup> half and Read during 2<sup>nd</sup> half of each clock cycle
    - Possible because RegFile access is *VERY* fast (takes less than half the time of ALU stage)
  - 2) Build RegFile with independent read and write ports
- **Conclusion:** Read and Write to registers during same clock cycle is okay

# Agenda

- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- Control Hazards
  - Branch and Jump Delay Slots
  - Branch Prediction

## 2. Data Hazards (1/2)

- Consider the following sequence of instructions:

add \$t0, \$t1, \$t2

sub \$t4, \$t0, \$t3

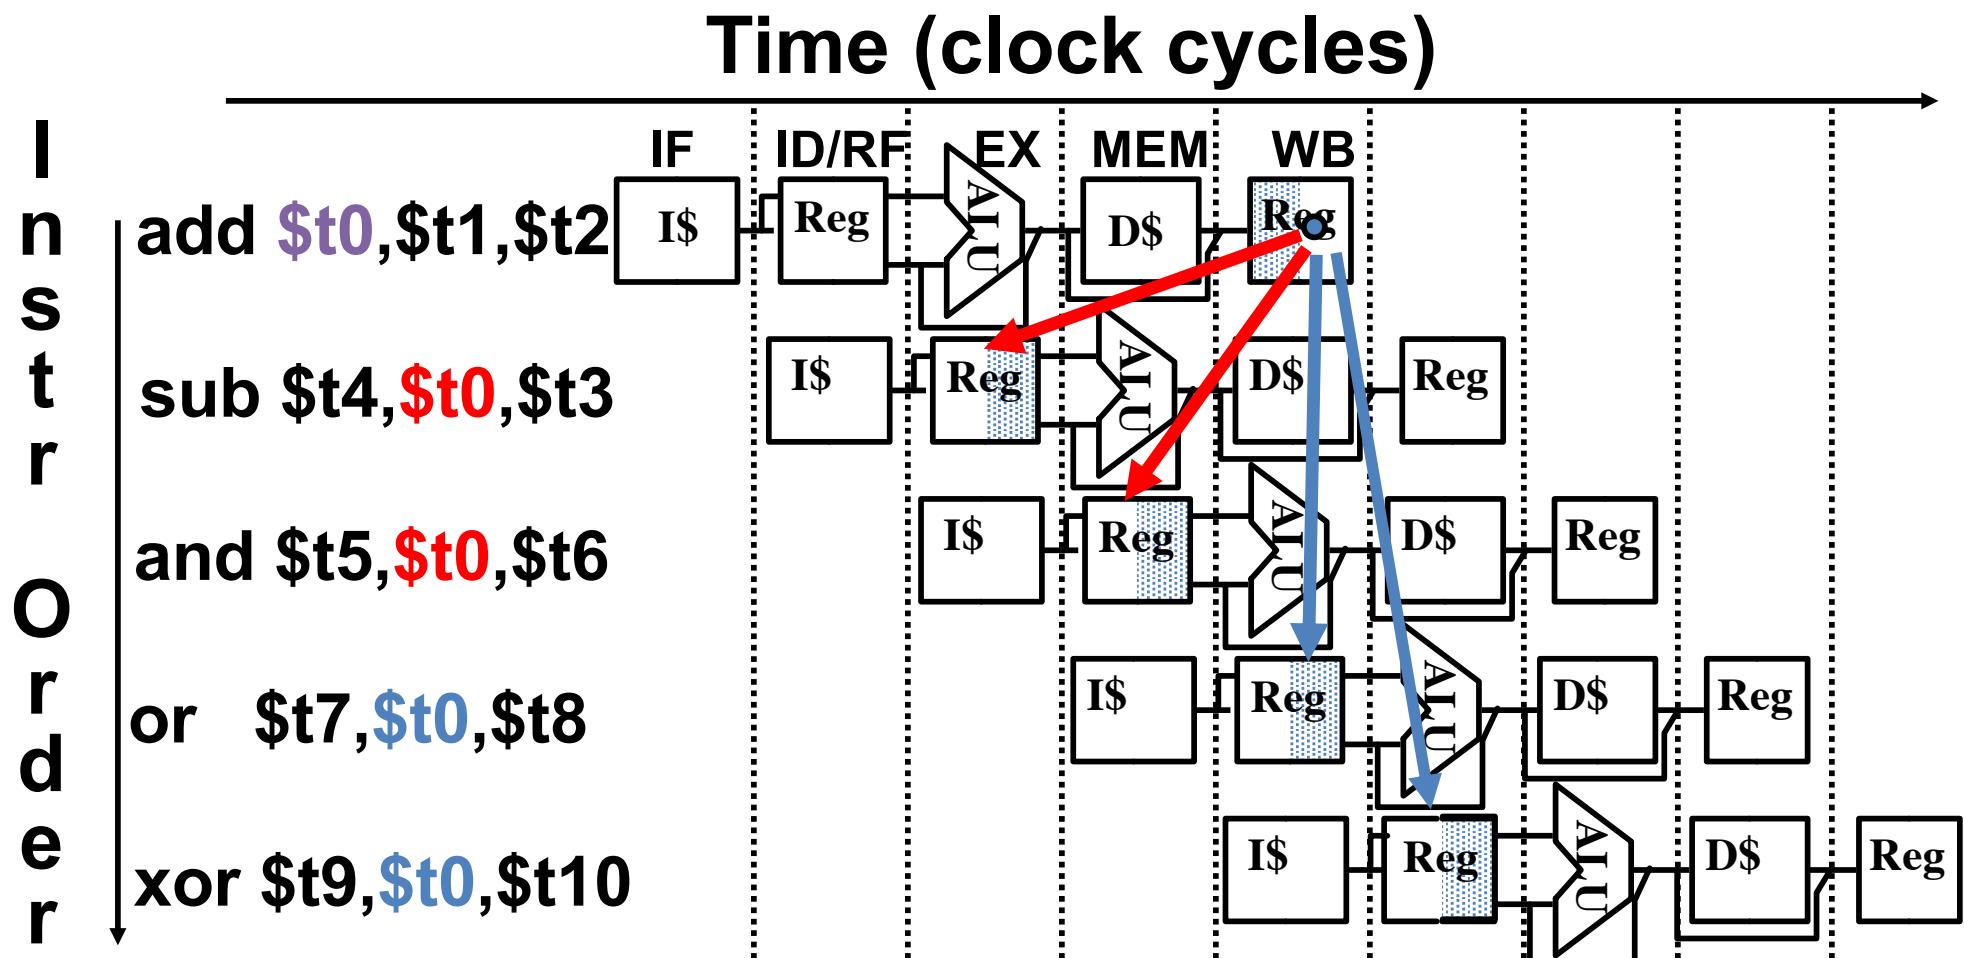
and \$t5, \$t0, \$t6

or \$t7, \$t0, \$t8

xor \$t9, \$t0, \$t10

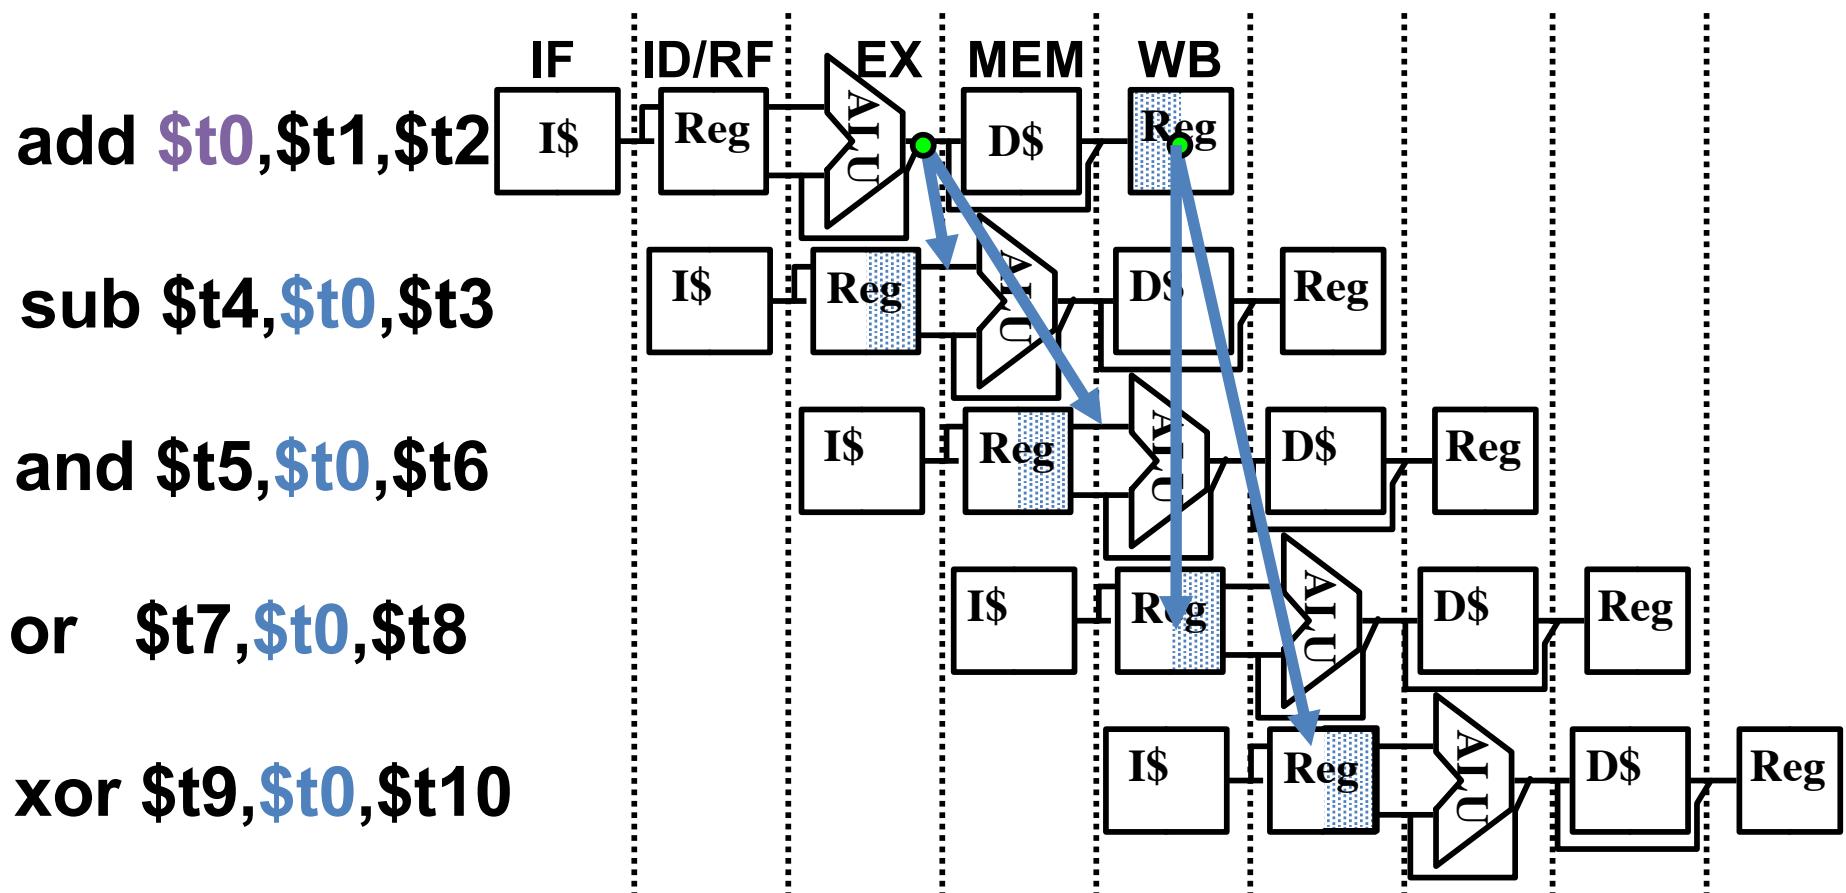
## 2. Data Hazards (2/2)

- Data-flow *backward* in time are hazards



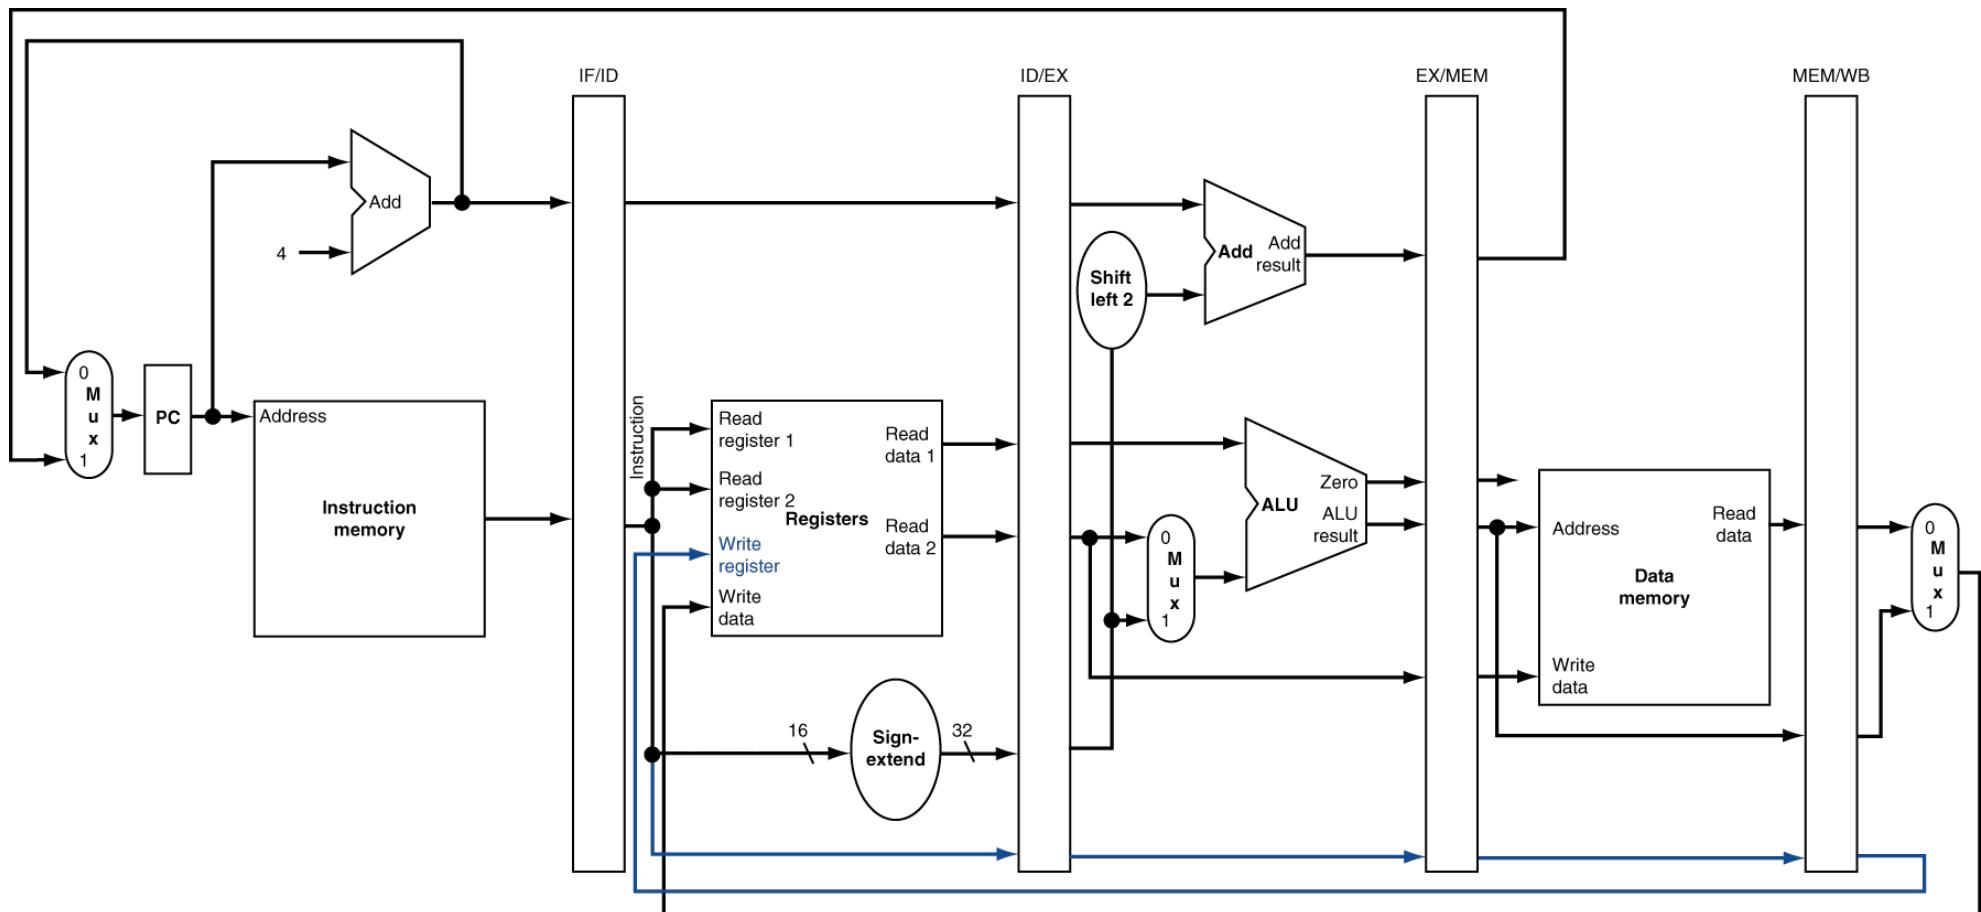
# Data Hazard Solution: Forwarding

- Forward result as soon as it is available
  - OK that it's not stored in RegFile yet



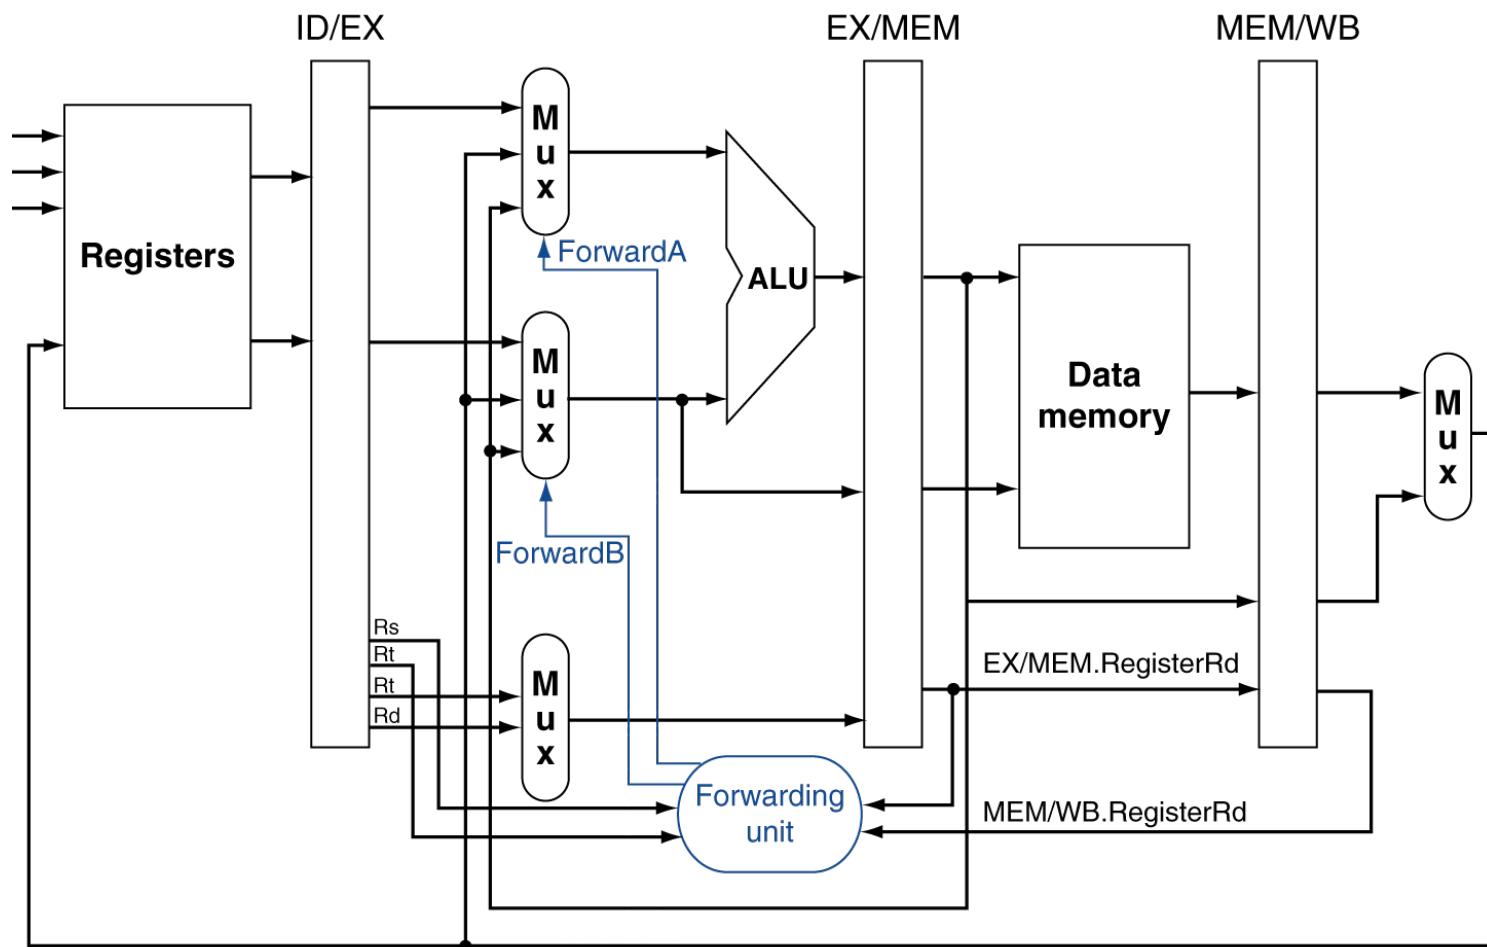
# Datapath for Forwarding (1/2)

- What changes need to be made here?



# Datapath for Forwarding (2/2)

- Handled by *forwarding unit*

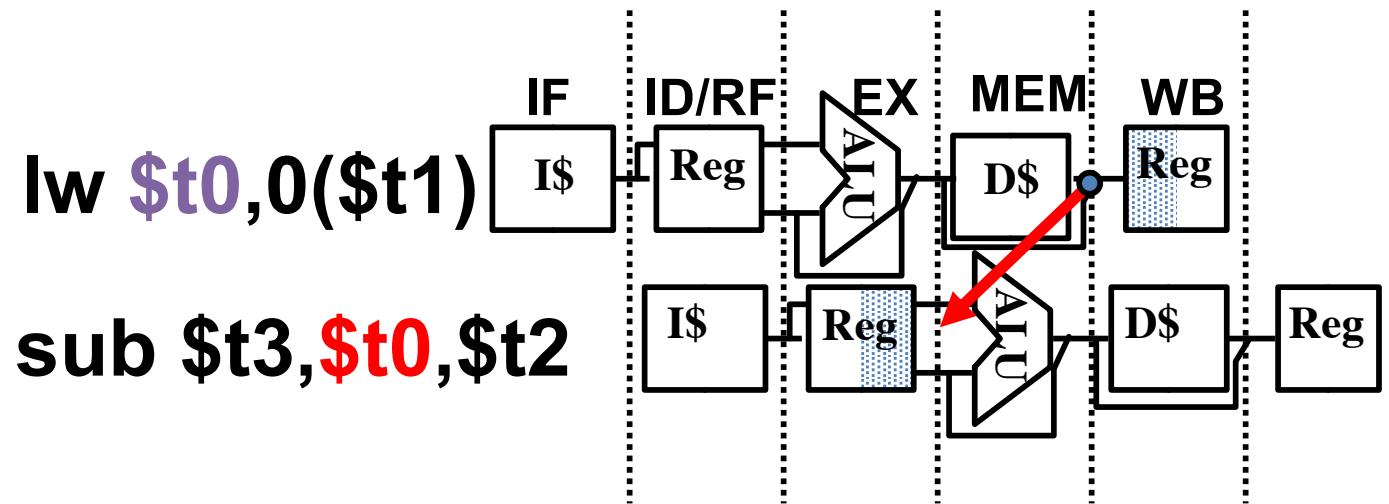


# Agenda

- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- Control Hazards
  - Branch and Jump Delay Slots
  - Branch Prediction

# Data Hazard: Loads (1/4)

- **Recall:** Dataflow backwards in time are hazards



- Can't solve all cases with forwarding
  - Must *stall* instruction dependent on load, then forward (more hardware)

# Data Hazard: Loads (2/4)

- *Hardware* stalls pipeline
  - Called “hardware interlock”

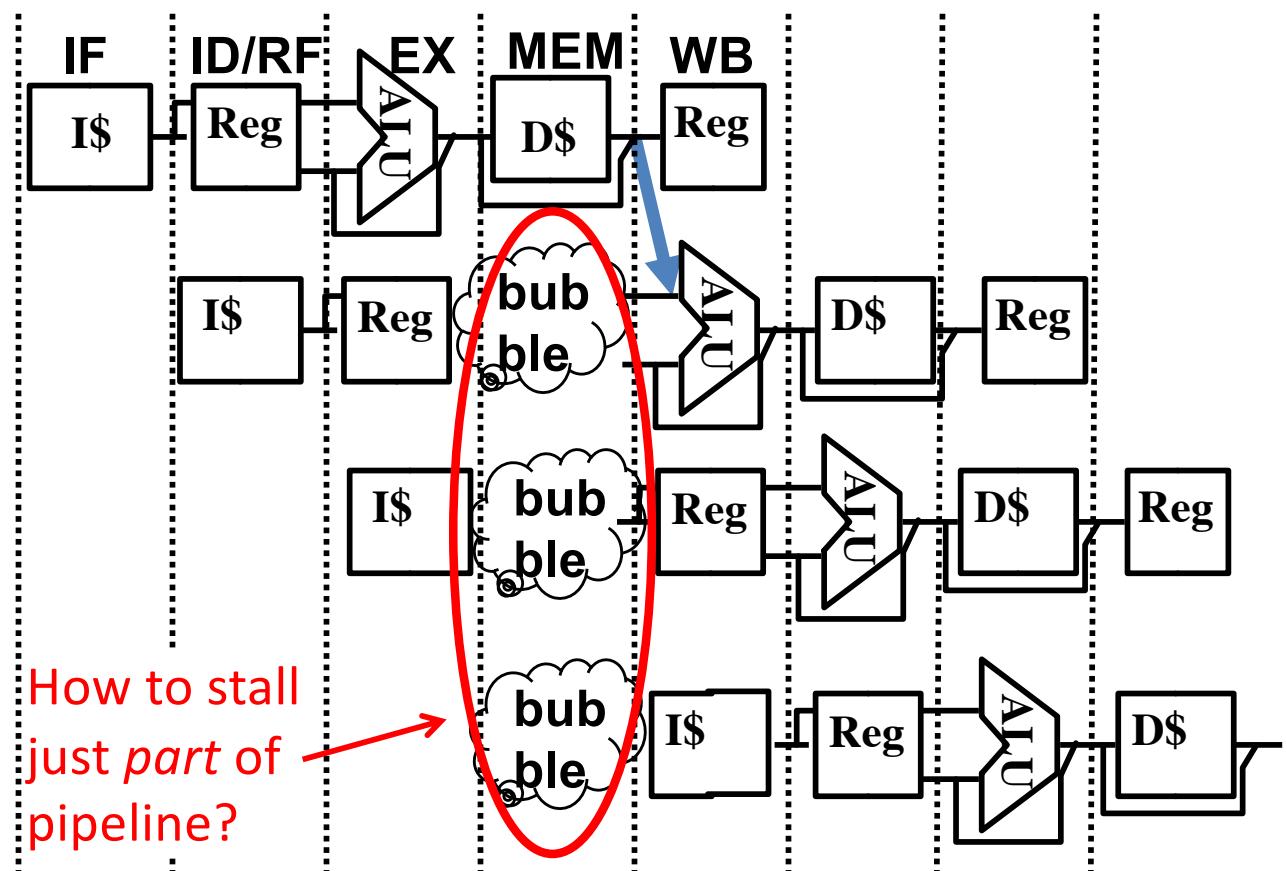
Schematically, this is what we want, but in reality stalls done “horizontally”

**lw \$t0, 0(\$t1)**

**sub \$t3,\$t0,\$t2**

**and \$t5,\$t0,\$t4**

**or \$t7,\$t0,\$t6**



# Data Hazard: Loads (3/4)

- Stall is equivalent to nop

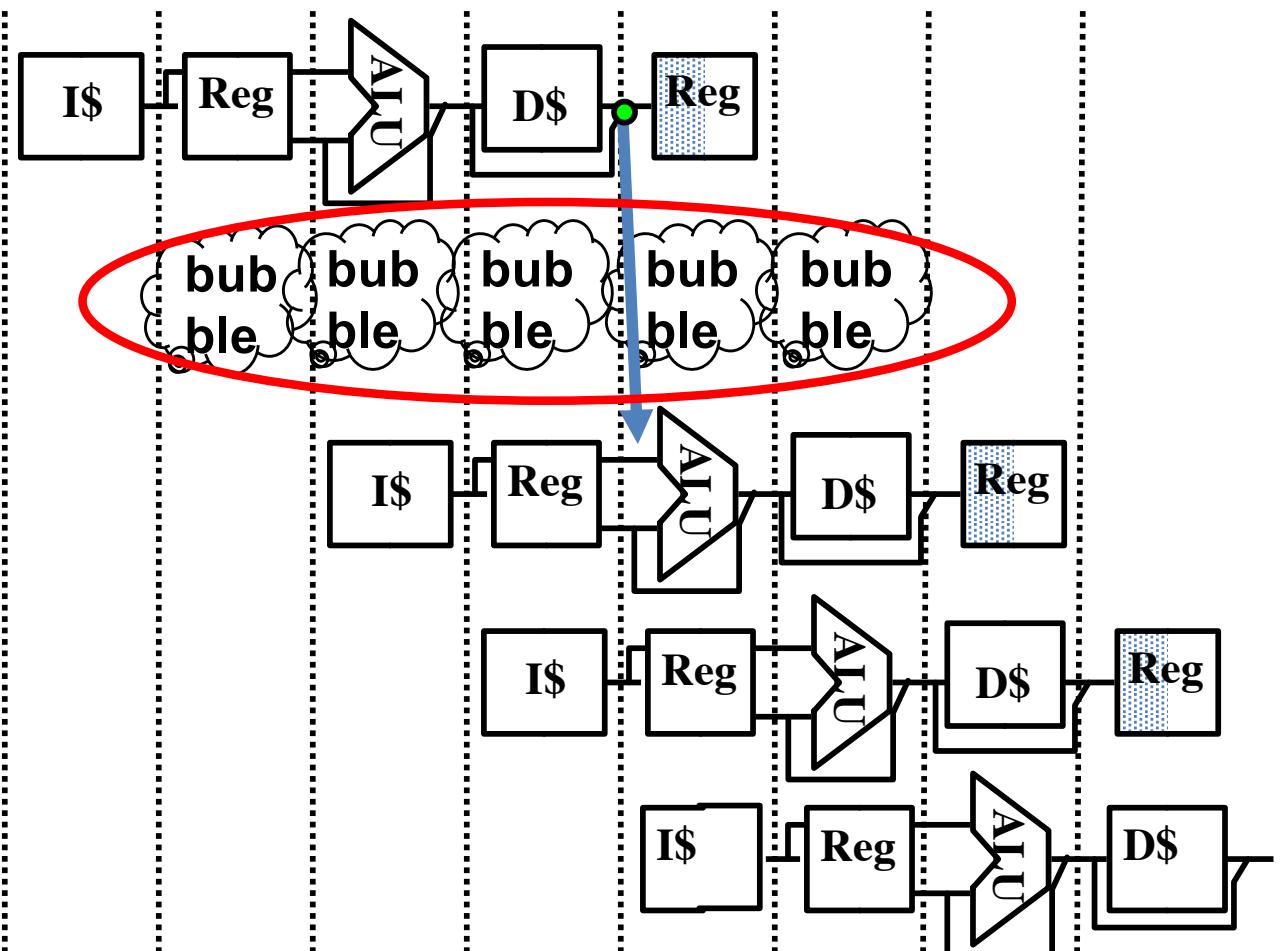
**lw \$t0, 0(\$t1)**

**nop**

**sub \$t3,\$t0,\$t2**

**and \$t5,\$t0,\$t4**

**or \$t7,\$t0,\$t6**

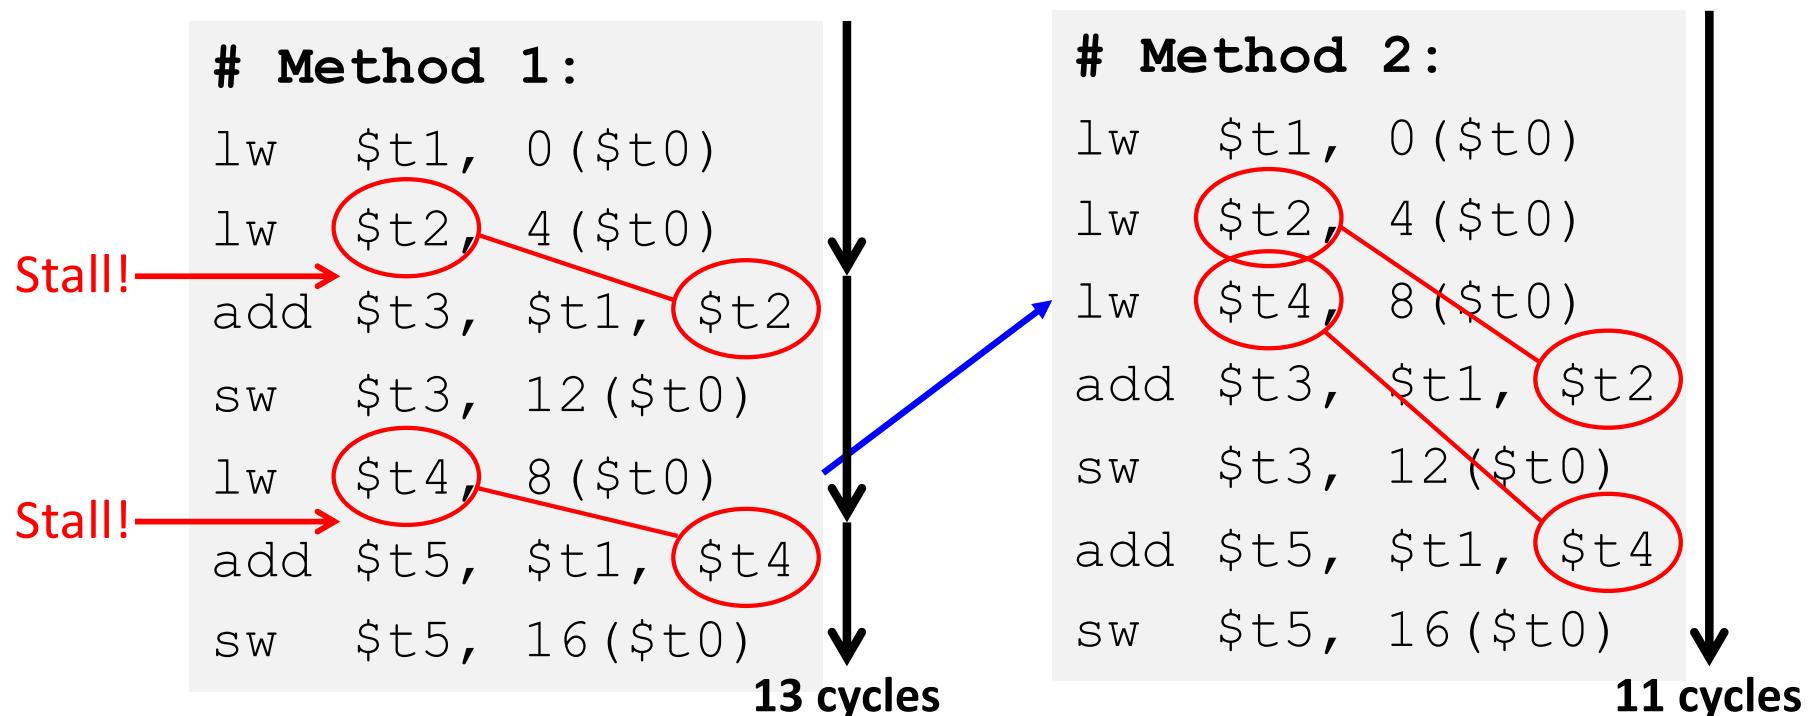


# Data Hazard: Loads (4/4)

- Slot after a load is called a *load delay slot*
  - If that instruction uses the result of the load, then the hardware interlock will stall it for one cycle
  - Letting the hardware stall the instruction in the delay slot is equivalent to putting a `nop` in the slot (except the latter uses more code space)
- **Idea:** Let the compiler put an unrelated instruction in that slot → no stall!

# Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!
- MIPS code for  $A=B+E; C=B+F;$



# Agenda

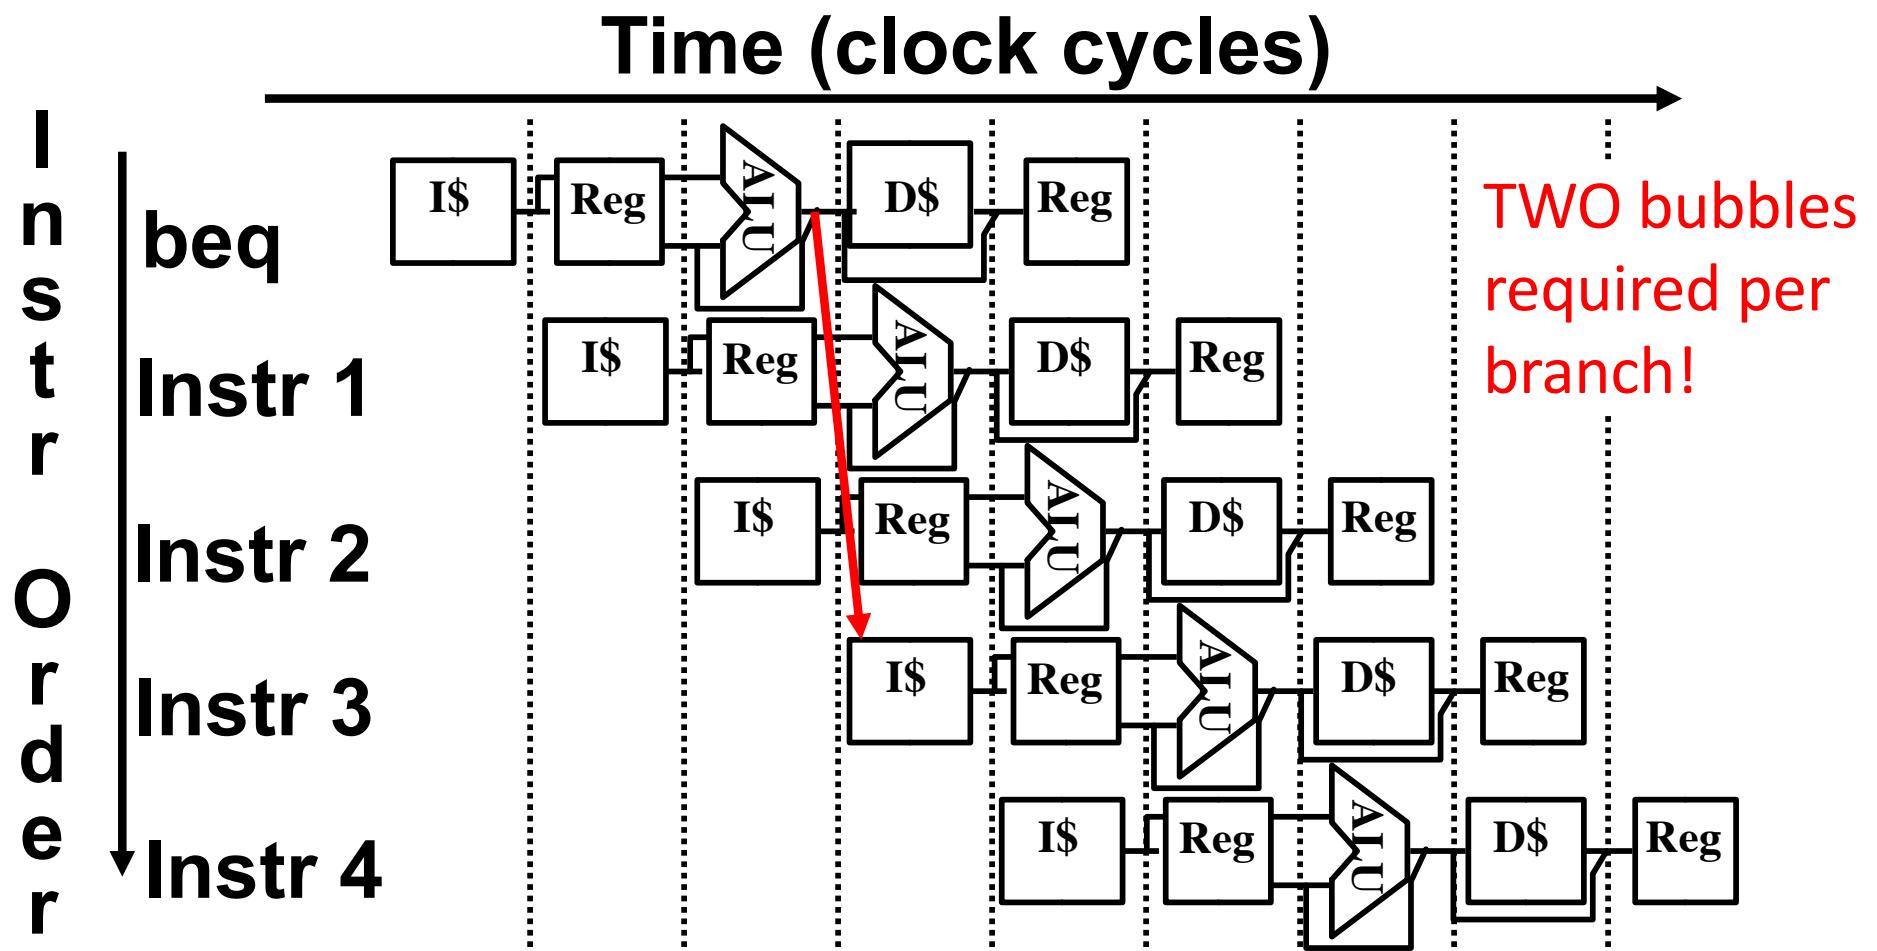
- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- Control Hazards
  - Branch and Jump Delay Slots
  - Branch Prediction

## 3. Control Hazards

- Branch (beq, bne) determines flow of control
  - Fetching next instruction depends on branch outcome
  - Pipeline can't always fetch correct instruction
    - Still working on ID stage of branch
- **Simple Solution:** Stall on *every* branch until we have the new PC value
  - How long must we stall?

# Branch Stall

- When is comparison result available?

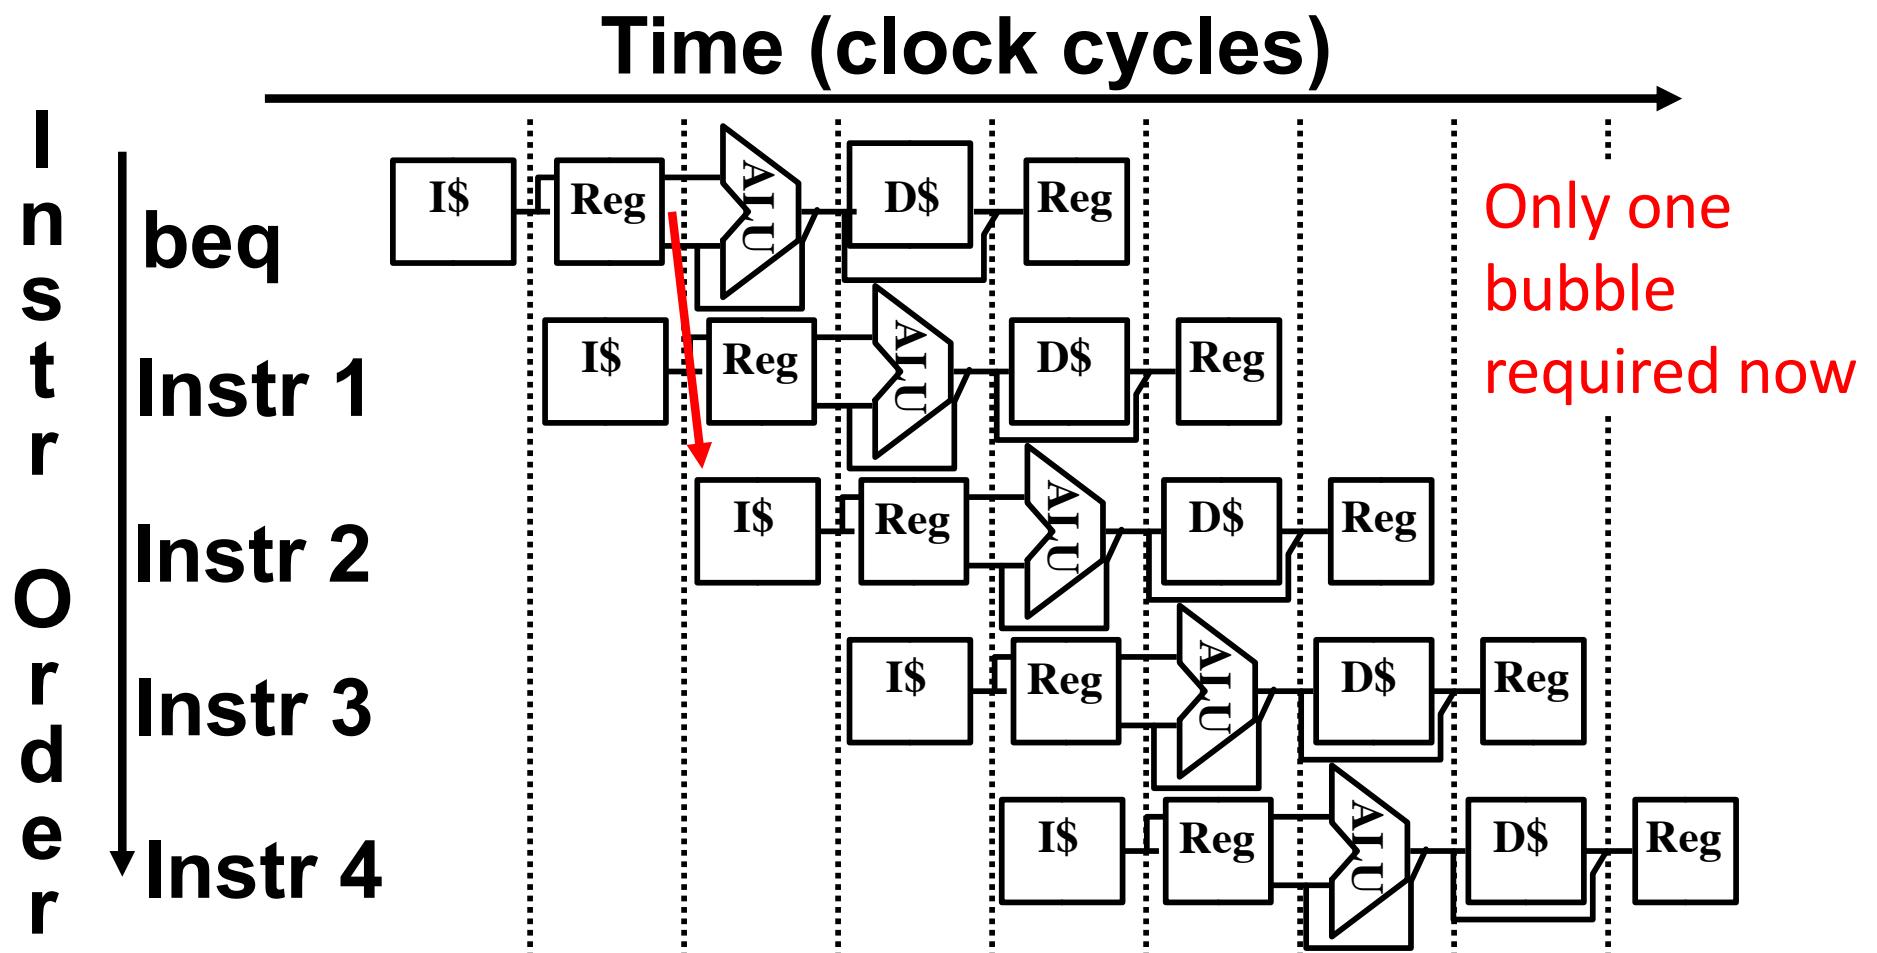


### 3. Control Hazard: Branching

- **Option #1:** Insert **special branch comparator** in ID stage
  - As soon as instruction is decoded, immediately make a decision and set the new value of the PC
  - **Benefit:** Branch decision made in 2<sup>nd</sup> stage, so only one `nop` is needed instead of two
  - **Side Note:** This means that branches are idle in EX, MEM, and WB

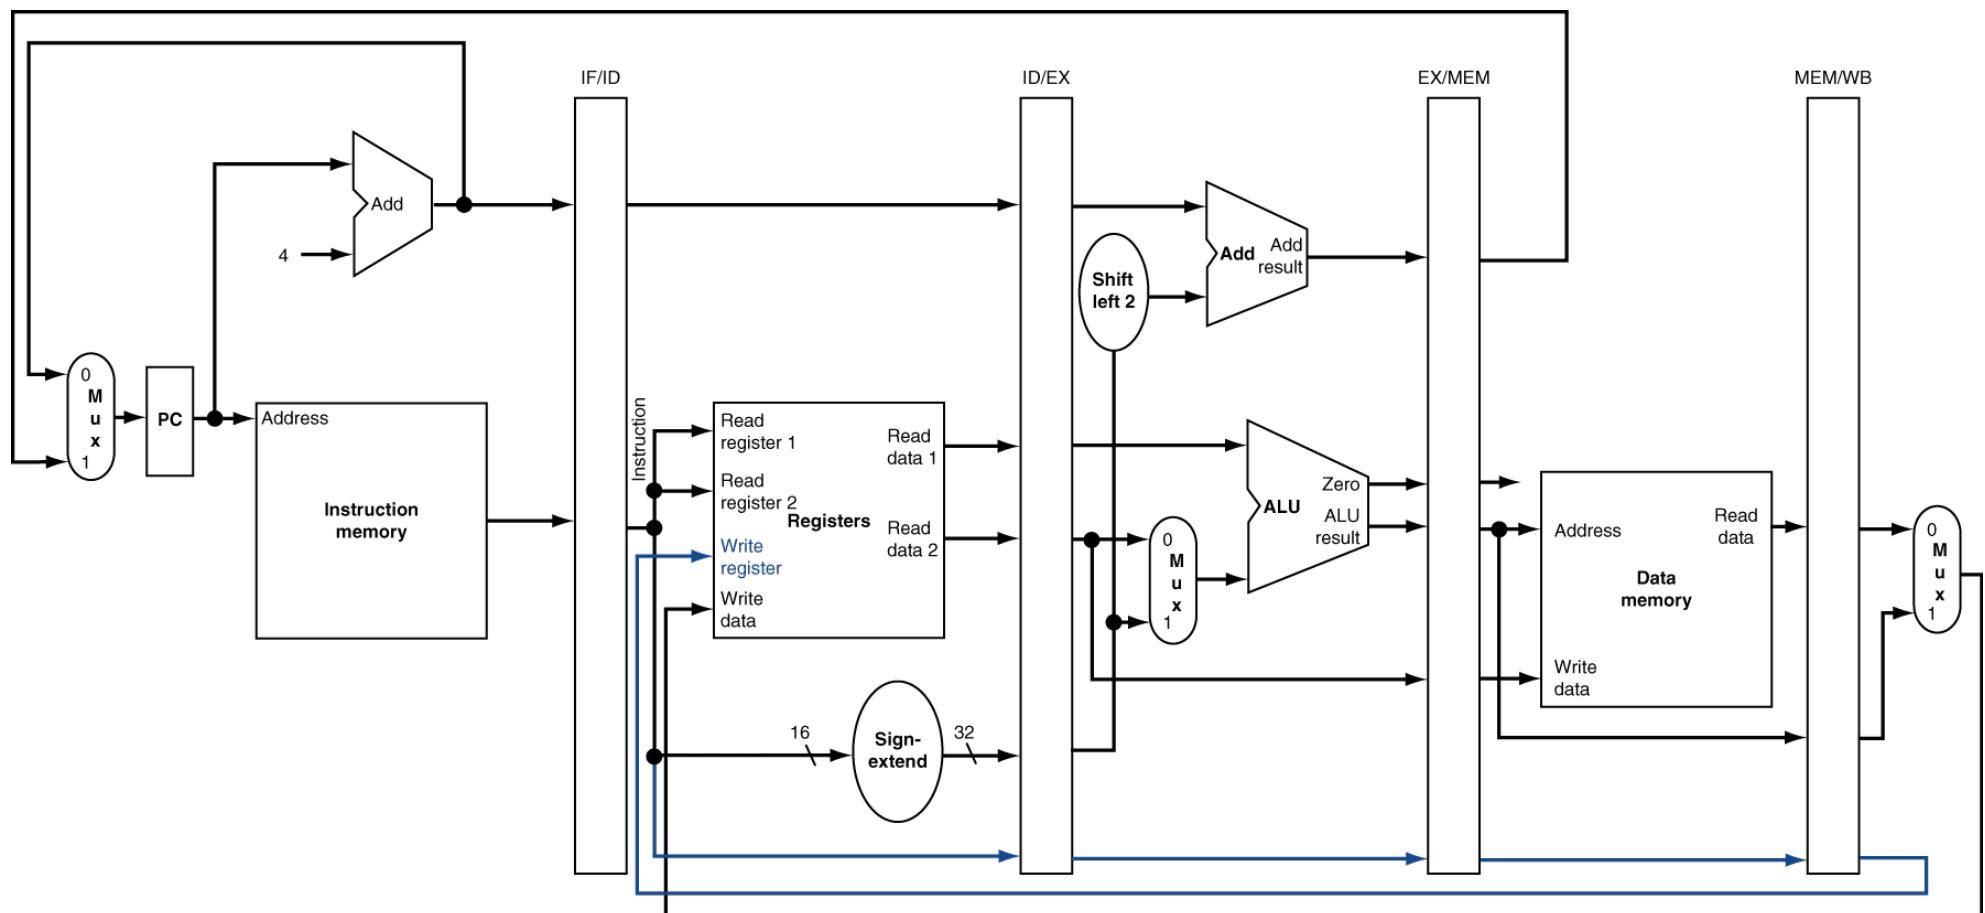
# Improved Branch Stall

- When is comparison result available?



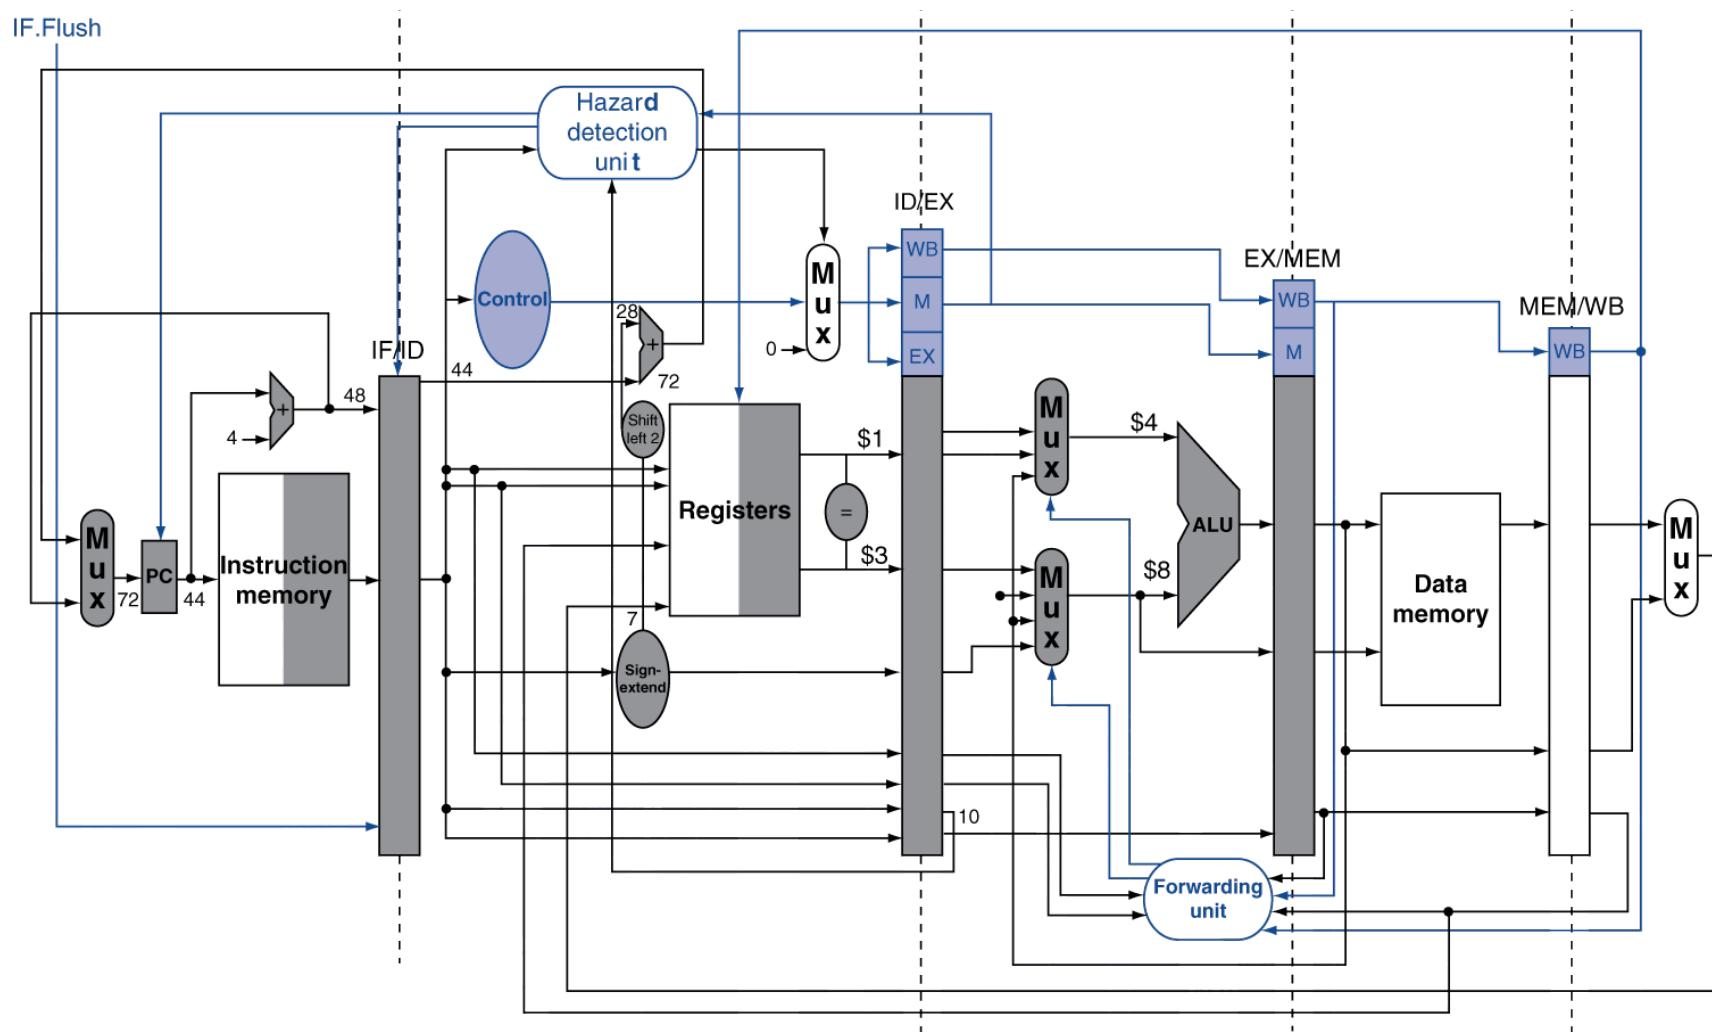
# Datapath for ID Branch Comparator

- What changes need to be made here?



# Datapath for ID Branch Comparator

- Handled by *hazard detection unit*



### 3. Control Hazard: Branching

- **Option #2:** *Branch Prediction* – guess outcome of a branch, fix afterwards if necessary
  - Must cancel (*flush*) all instructions in pipeline that depended on guess that was wrong
  - How many instructions do we end up flushing?
- Achieve simplest hardware if we predict that all branches are NOT taken

### 3. Control Hazard: Branching

- **Option #3:** *Branch delay slot*
  - Whether or not we take the branch, *always* execute the instruction immediately following the branch
  - Worst-Case: Put a `nop` in the branch-delay slot
  - Better Case: Move an instruction from before the branch into the branch-delay slot
    - Must not affect the logic of program

### 3. Control Hazard: Branching

- MIPS uses this *delayed branch* concept
  - Re-ordering instructions is a common way to speed up programs
  - Compiler finds an instruction to put in the branch delay slot  $\approx 50\%$  of the time
- Jumps also have a delay slot
  - Why is one needed?

# Delayed Branch Example

## Nondelayed Branch

```
or $8, $9, $10
```

```
add $1, $2, $3
```

```
sub $4, $5, $6
```

```
beq $1, $4, Exit
```

```
xor $10, $1, $11
```

Exit:

## Delayed Branch

```
add $1, $2, $3
```

```
sub $4, $5, $6
```

```
beq $1, $4, Exit
```

or \$8, \$9, \$10

```
xor $10, $1, $11
```

Why not any of the  
other instructions?

# Delayed Jump in MIPS

- MIPS Green Sheet for `jal`:

$R[31] = PC + 8;$     $PC = \text{JumpAddr}$

- $PC + 8$  because of *jump delay slot!*
- Instruction at  $PC + 4$  always gets executed before `jal` jumps to label, so return to  $PC + 8$

# Agenda

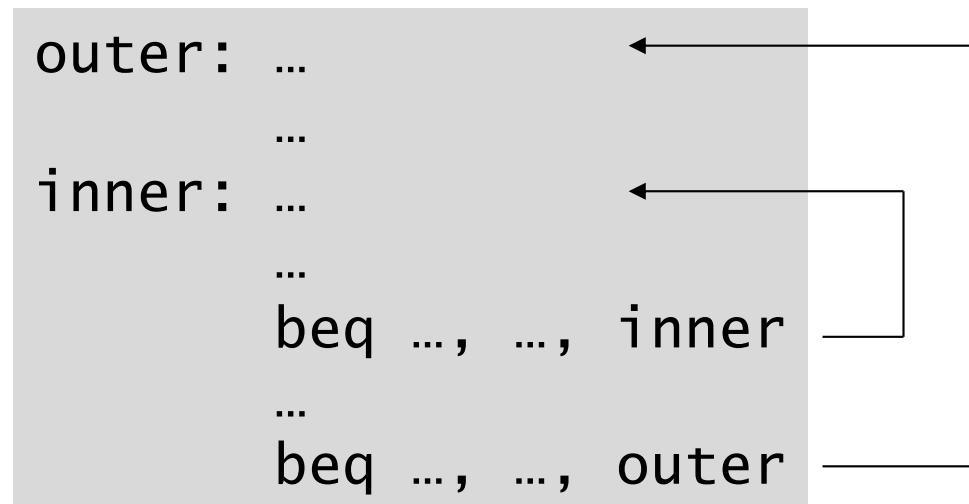
- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- **Control Hazards**
  - Branch and Jump Delay Slots
  - **Branch Prediction**

# Dynamic Branch Prediction

- Branch penalty is more significant in deeper pipelines
  - Also superscalar pipelines
- Use *dynamic branch prediction*
  - Have branch prediction buffer (a.k.a. branch history table) that stores outcomes (taken/not taken) indexed by recent branch instruction addresses
  - To execute a branch
    - Check table and predict the same outcome for next fetch
    - If wrong, flush pipeline and flip prediction

# 1-Bit Predictor: Shortcoming

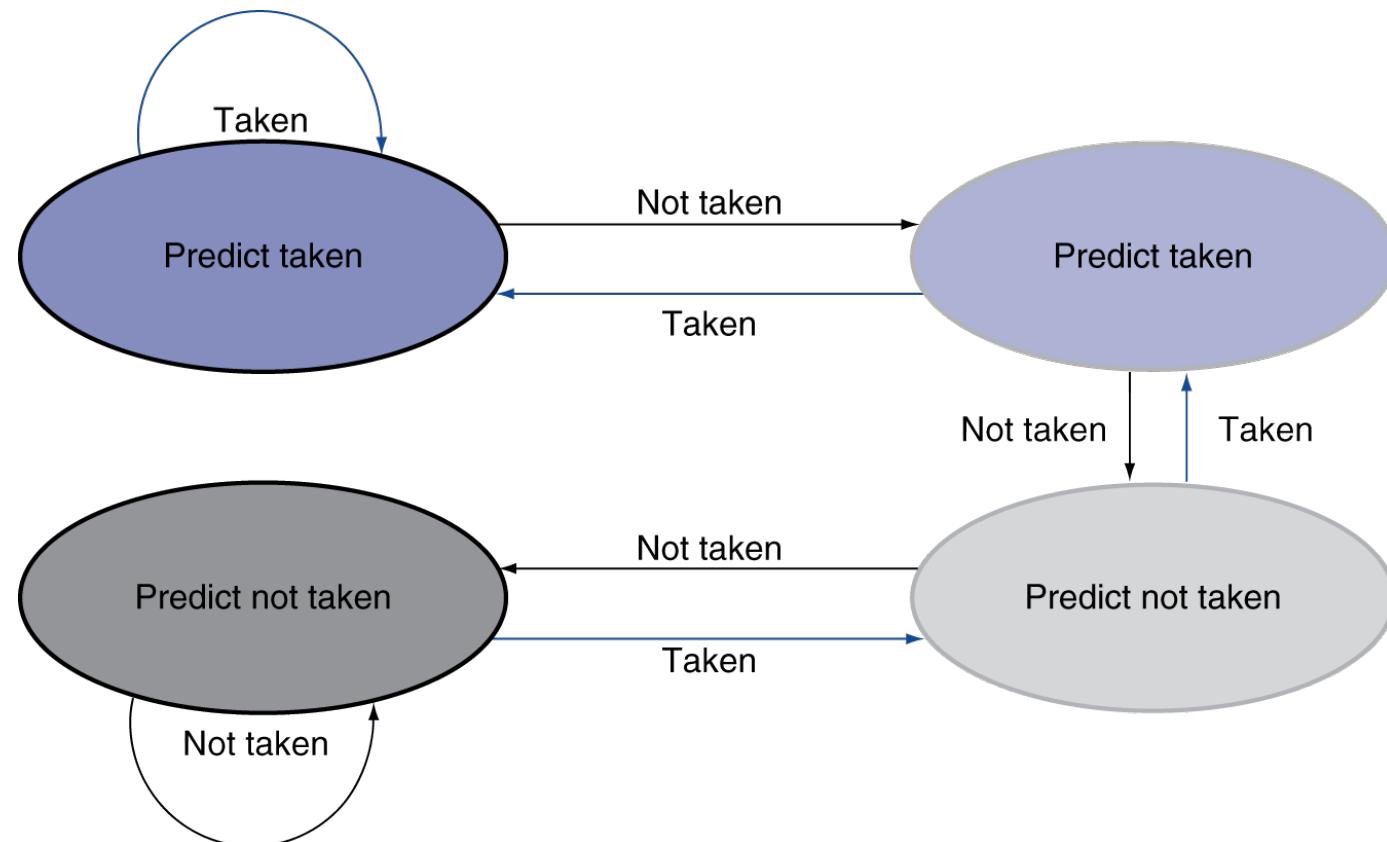
- Examine the code below, assuming both loops will be executed multiple times:



- Inner loop branches are predicted wrong twice!
  - Predict as taken on last iteration of inner loop
  - Then predict as not taken on first iteration of inner loop next time around

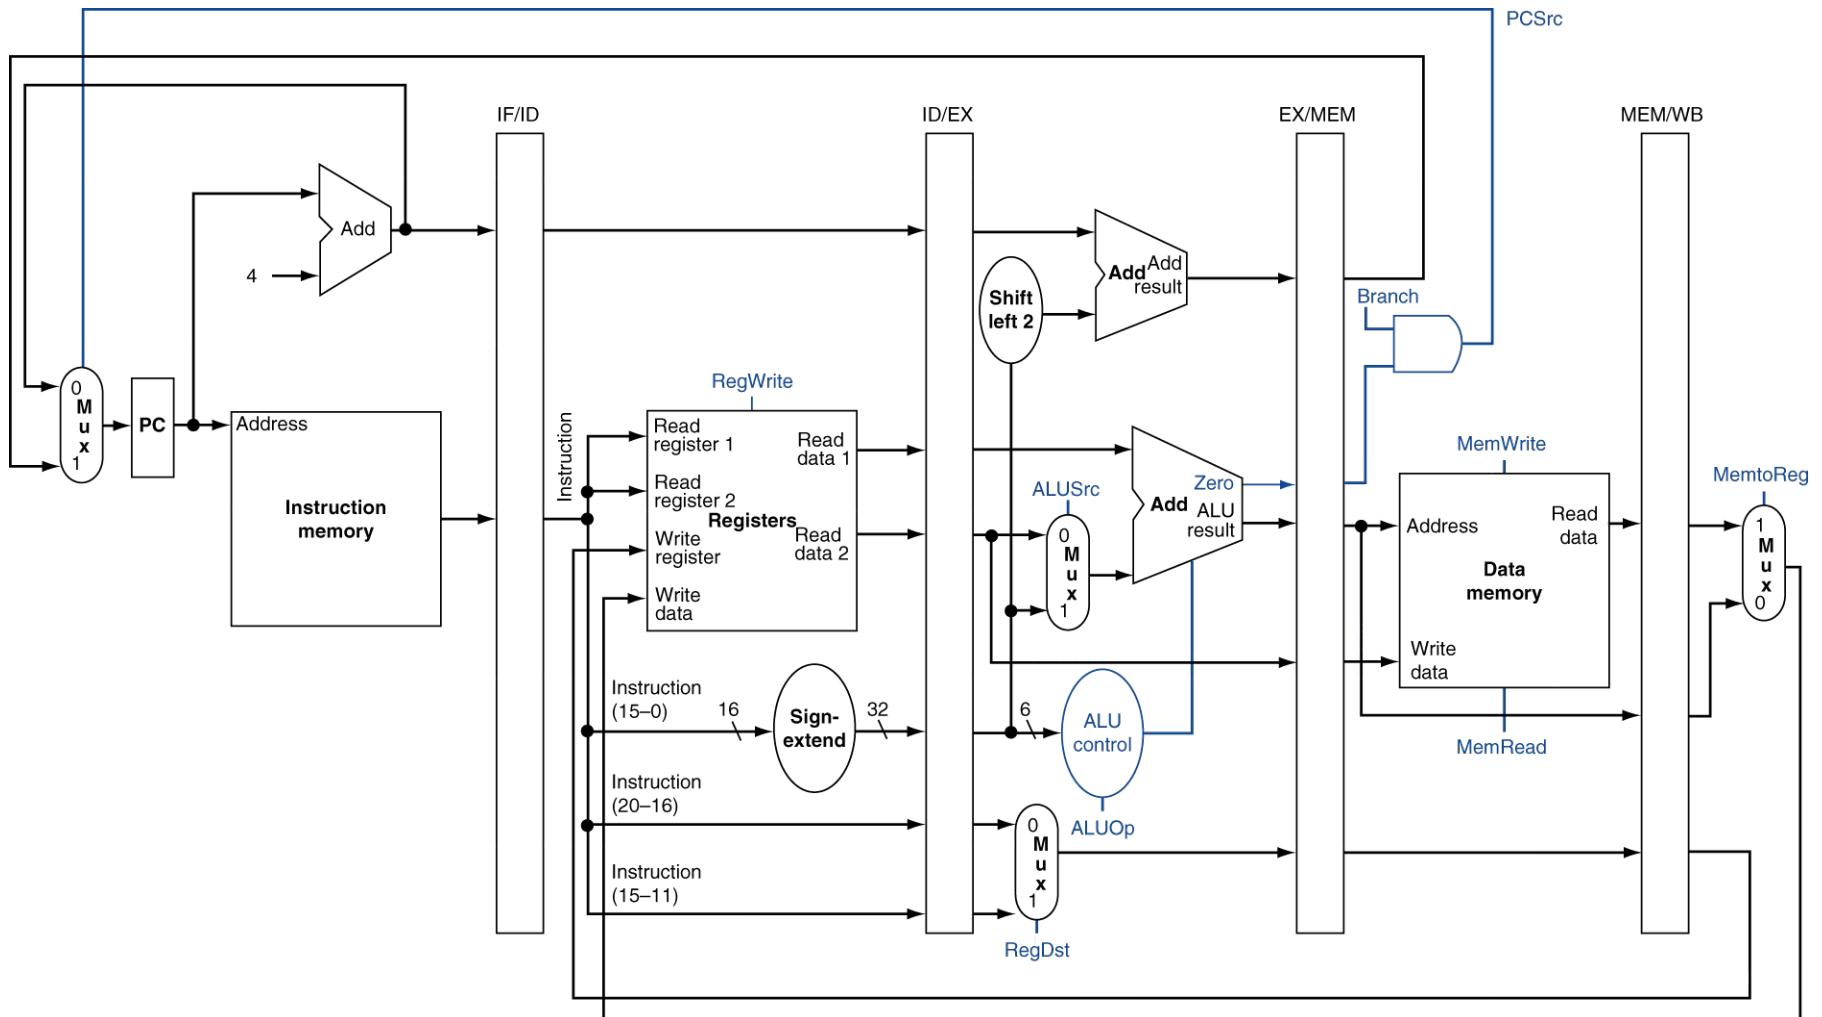
# 2-Bit Predictor

- Only change prediction after two successive incorrect predictions



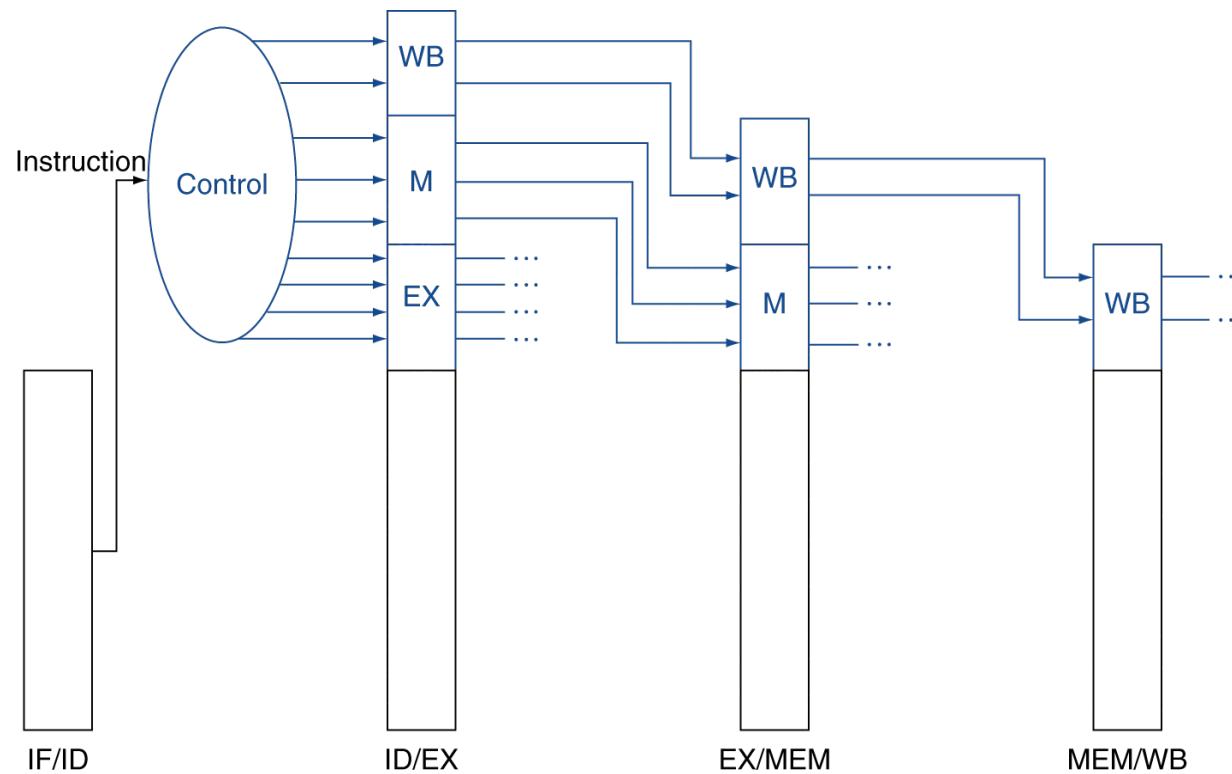
# Summary

## Pipelined Control (Simplified)



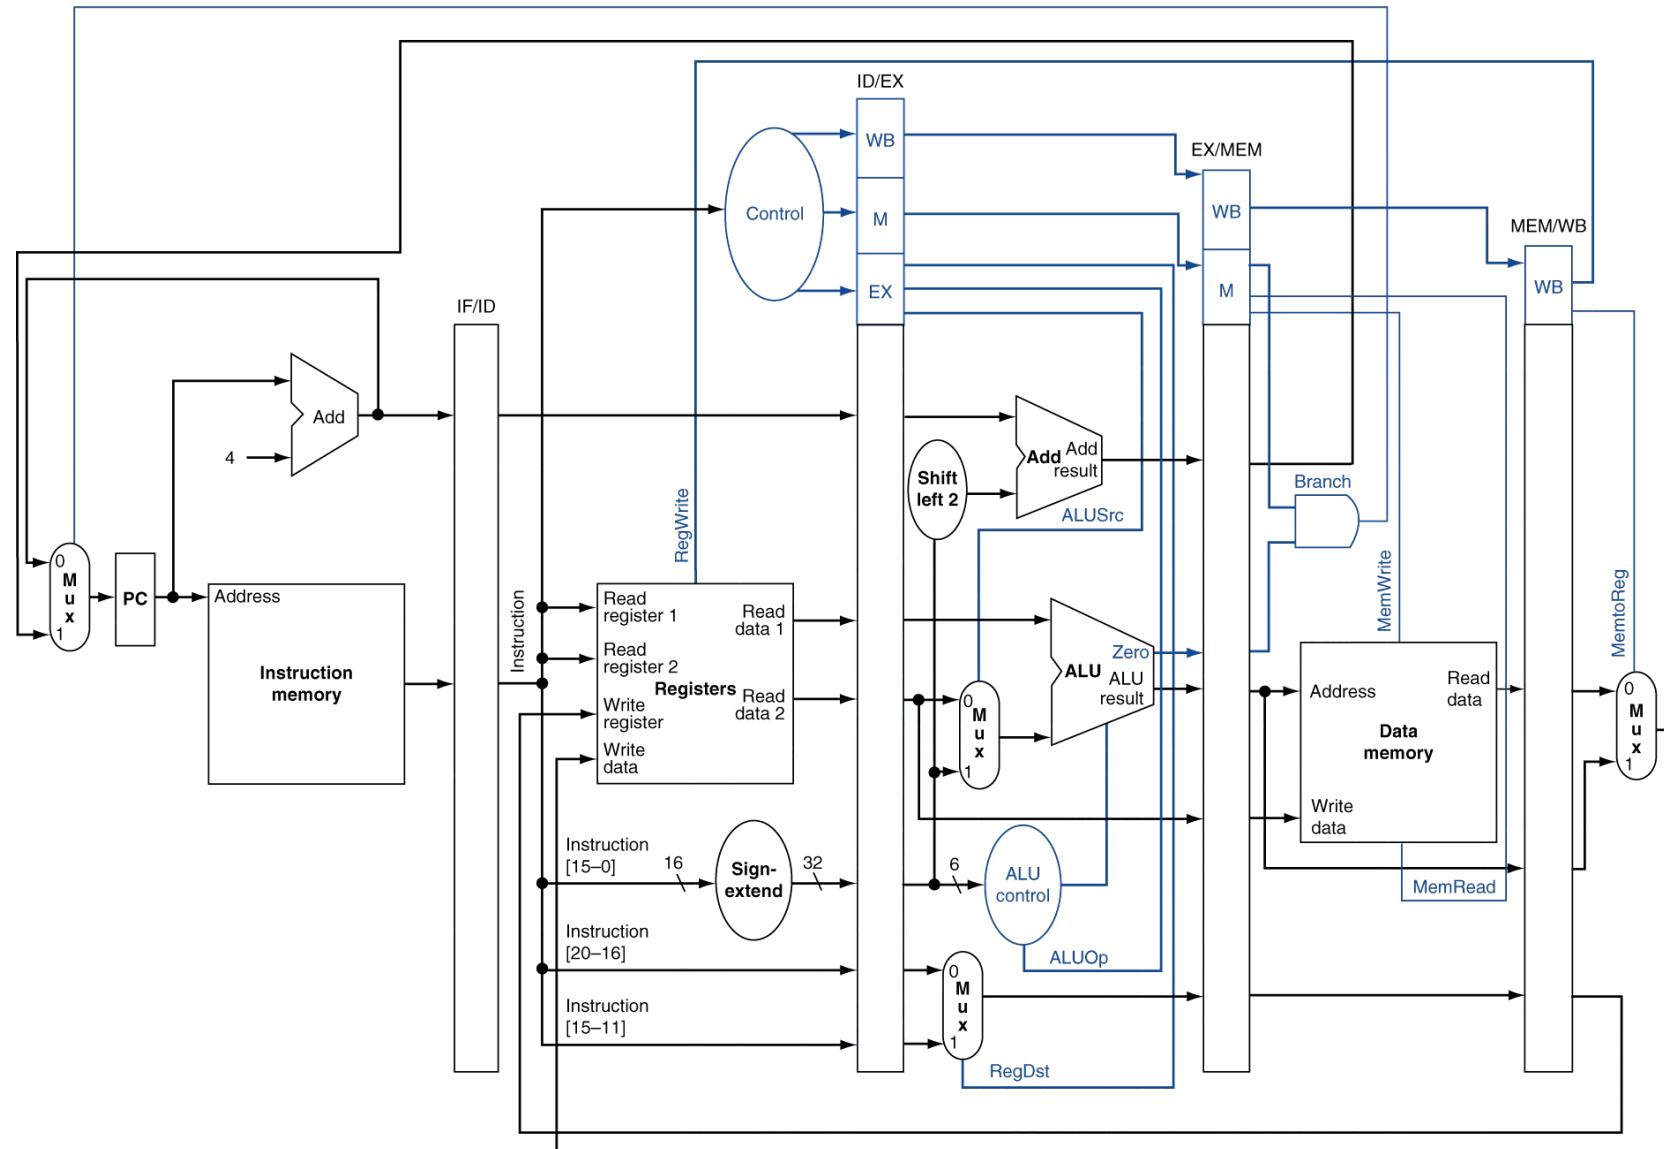
# Pipelined Control

- Control signals derived from instruction
  - As in single-cycle implementation



# Pipelined Control

PCSrc



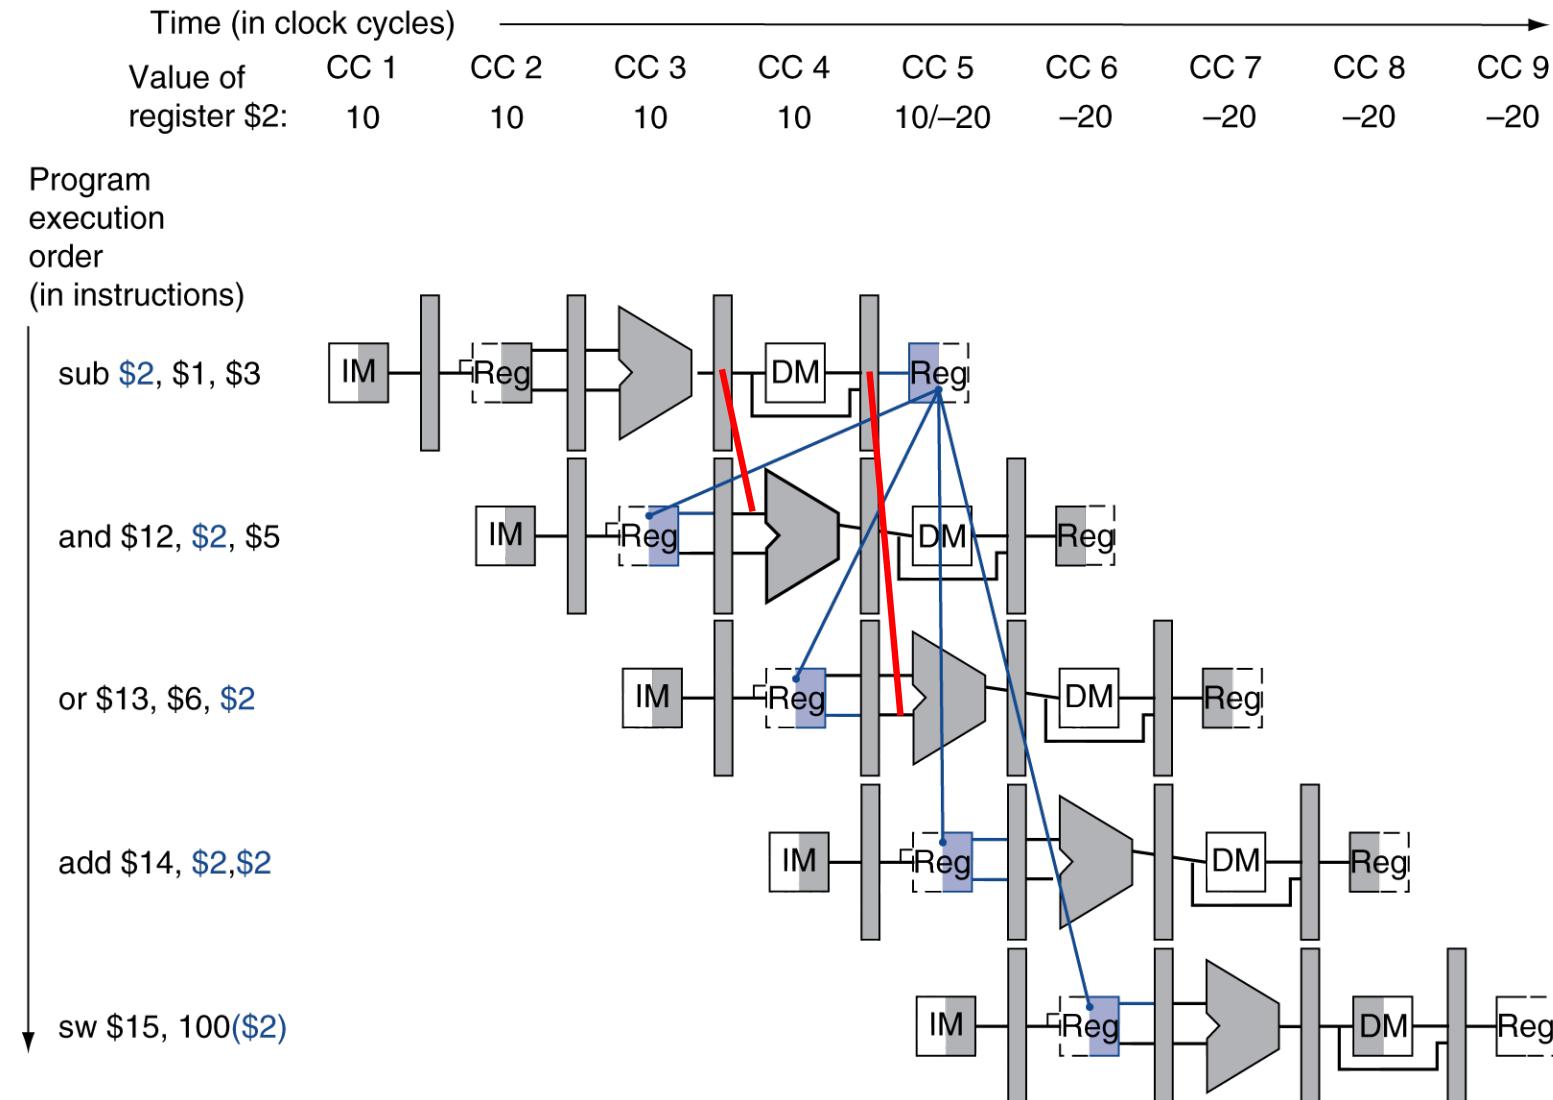
# Data Hazards in ALU Instructions

- Consider this sequence:

```
sub $2, $1,$3  
and $12, $2,$5  
or $13, $6,$2  
add $14, $2,$2  
sw $15,100($2)
```

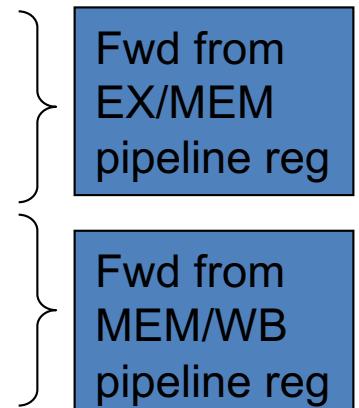
- We can resolve hazards with forwarding
  - How do we detect when to forward?

# Dependencies & Forwarding



# Detecting the Need to Forward

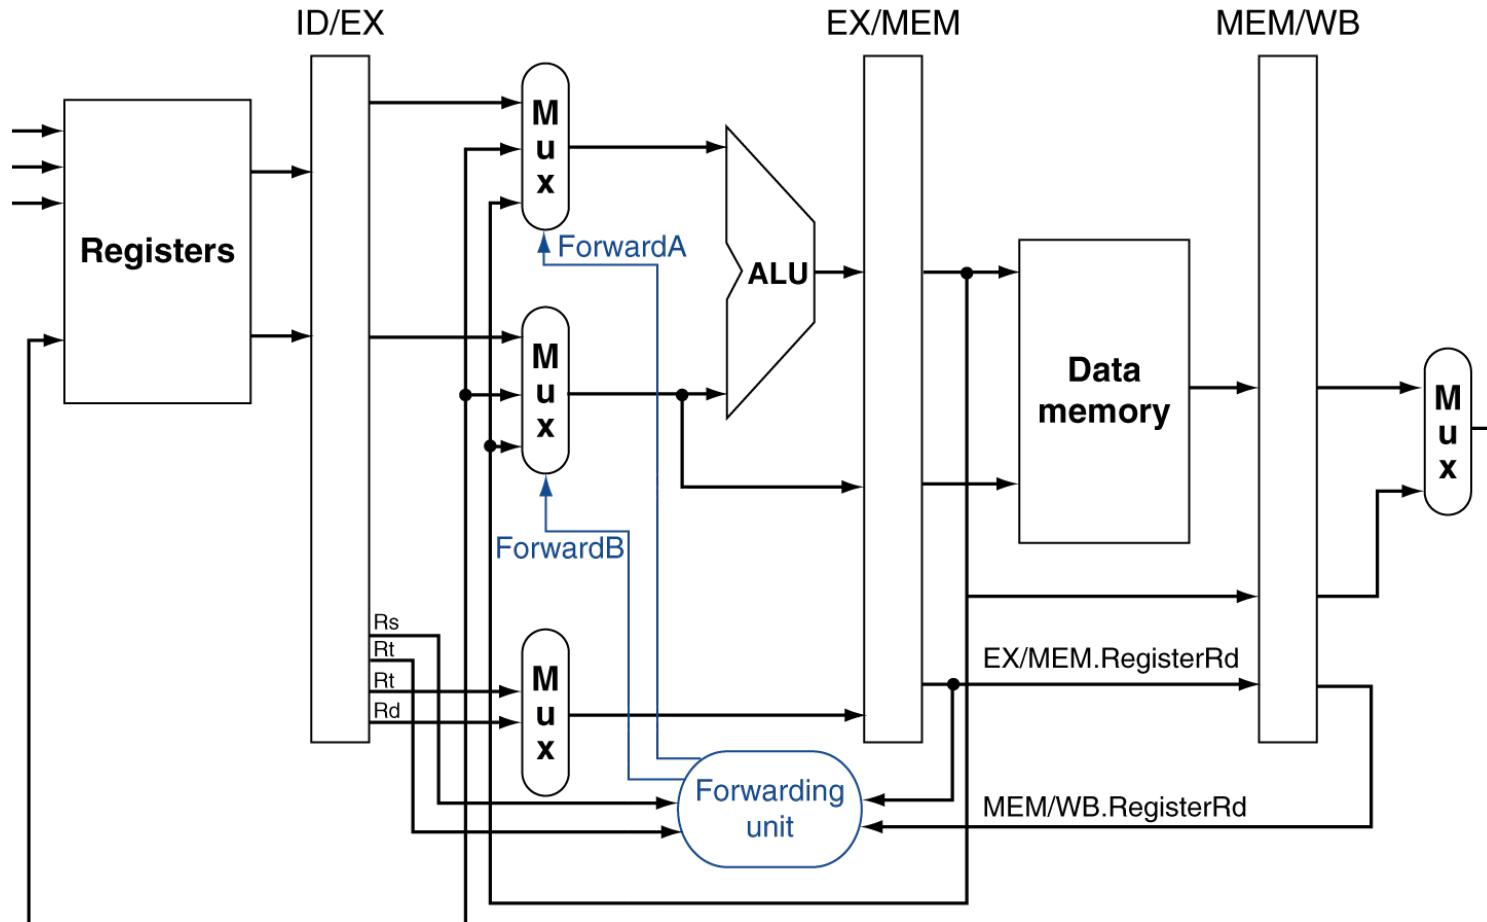
- Pass register numbers along pipeline
  - e.g., ID/EX.RegisterRs = register number for Rs sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
  - ID/EX.RegisterRs, ID/EX.RegisterRt
- Data hazards when
  - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
  - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
  - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
  - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt



# Detecting the Need to Forward

- But only if forwarding instruction will write to a register!
  - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not \$zero
  - EX/MEM.RegisterRd  $\neq 0$ ,
  - MEM/WB.RegisterRd  $\neq 0$

# Forwarding Paths



b. With forwarding

# Forwarding Conditions

- EX hazard
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))  
**ForwardA = 10**
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))  
**ForwardB = 10**
- MEM hazard
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
**ForwardA = 01**
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
**ForwardB = 01**

# Double Data Hazard

- Consider the sequence:

add \$1, \$1, \$2

add \$1, \$1, \$3

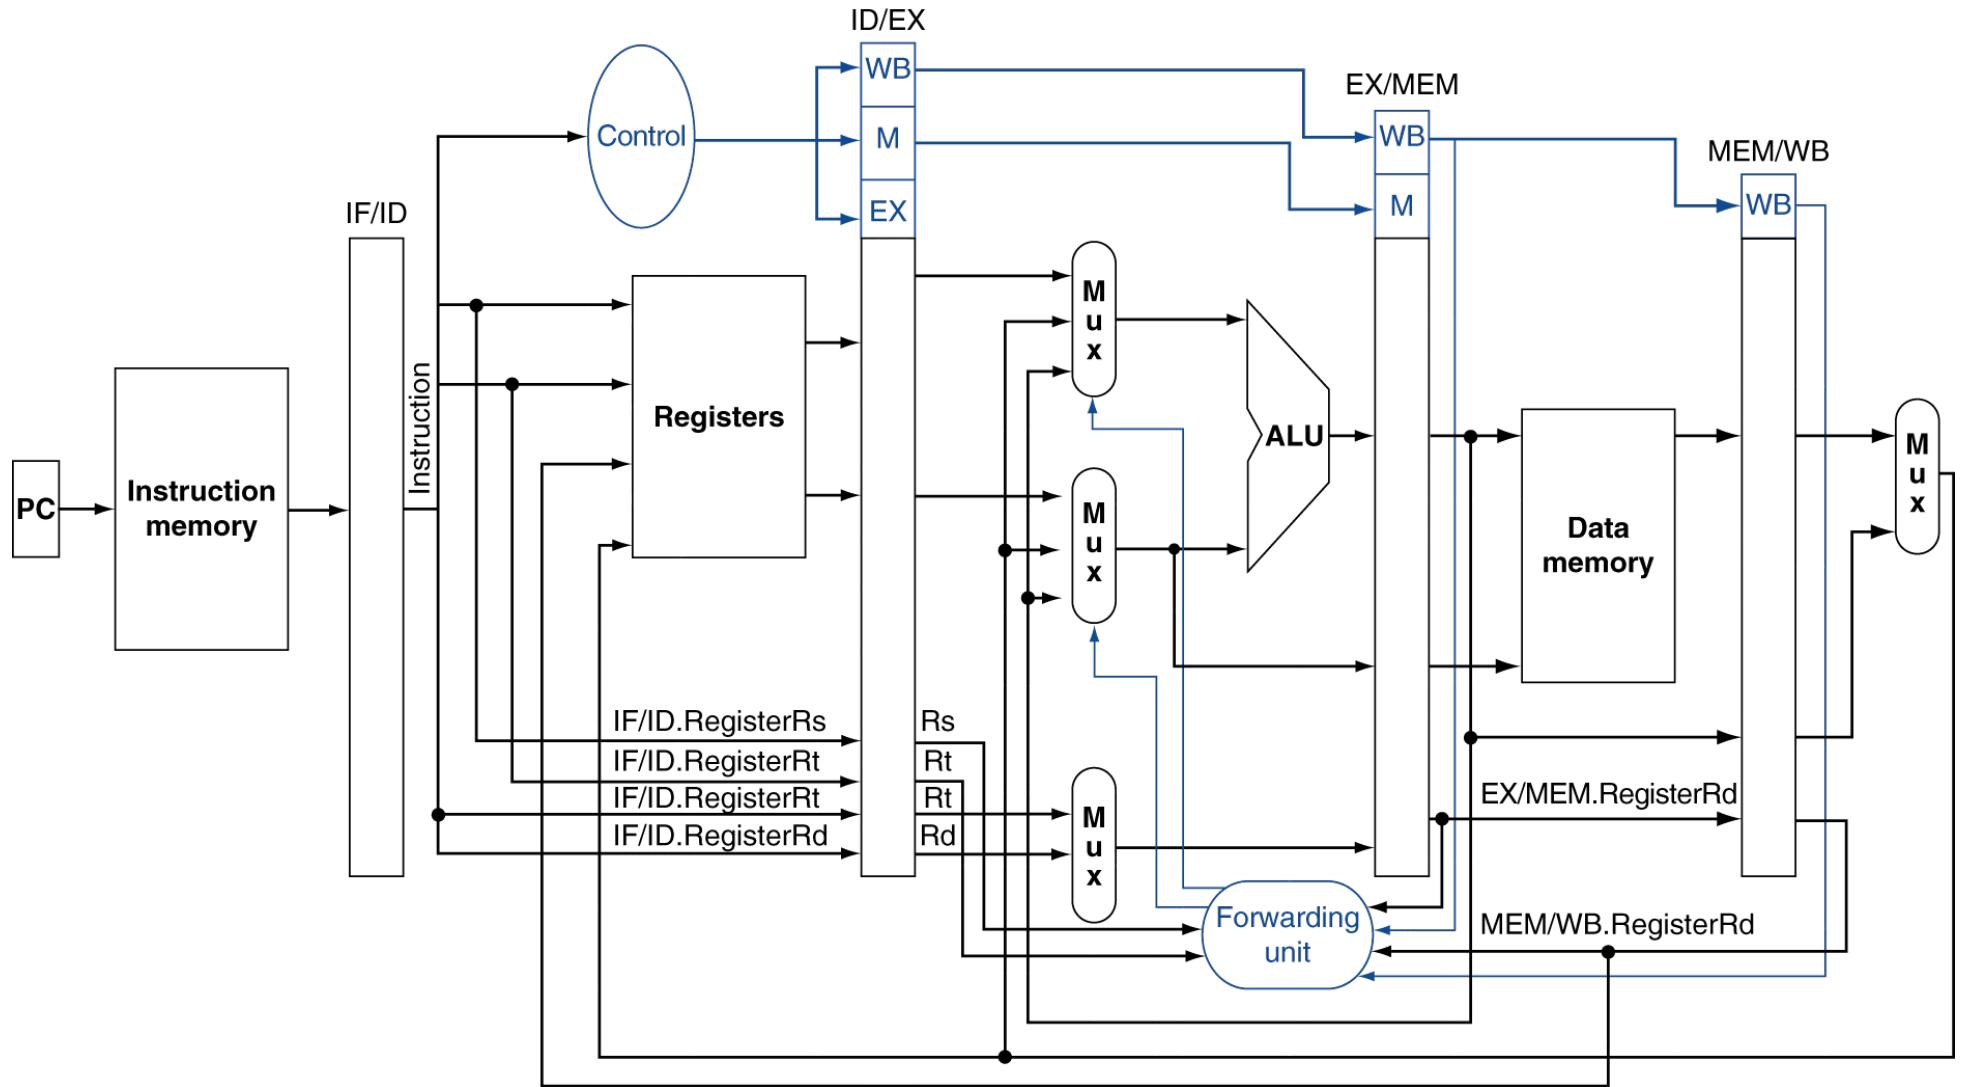
add \$1, \$1, \$4

- Both hazards occur
  - Want to use the most recent
- Revise MEM hazard condition
  - Only fwd if EX hazard condition isn't true

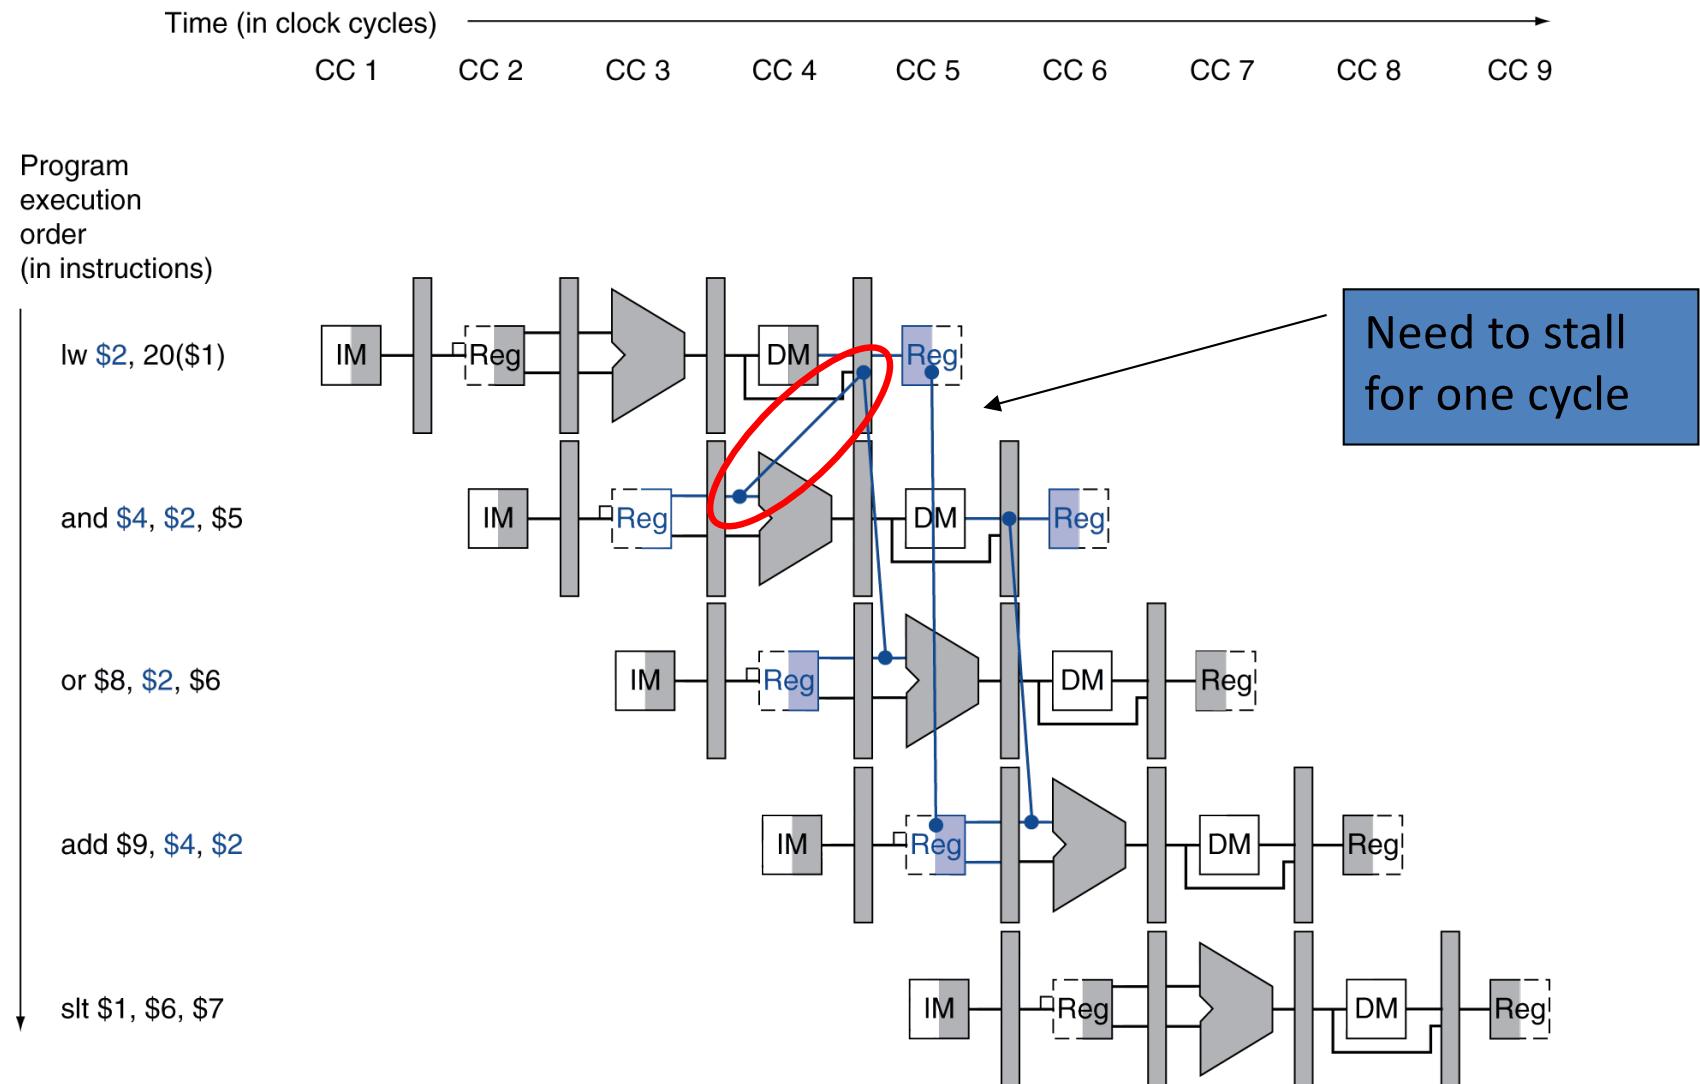
# Revised Forwarding Condition

- MEM hazard
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 01
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 01

# Datapath with Forwarding



# Load-Use Data Hazard



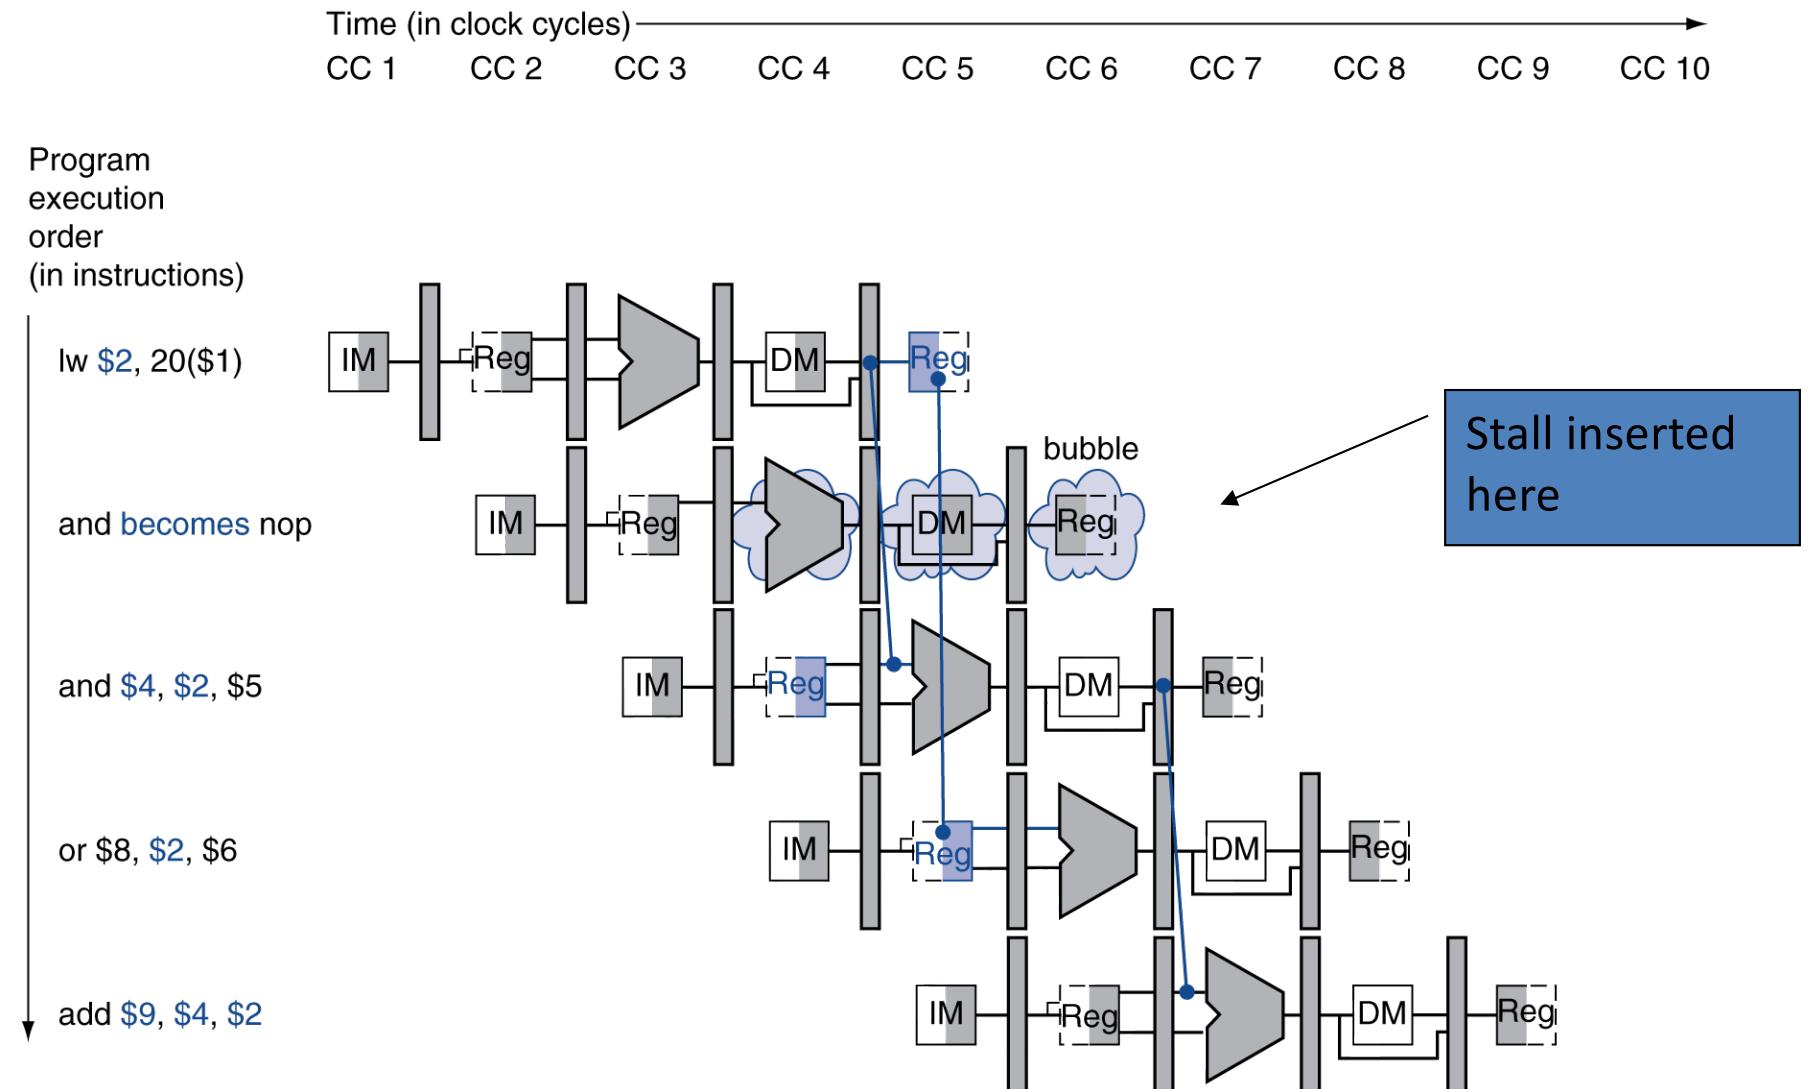
# Load-Use Hazard Detection

- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
  - IF/ID.RegisterRs, IF/ID.RegisterRt
- Load-use hazard when
  - ID/EX.MemRead and
$$((ID/EX.RegisterRt = IF/ID.RegisterRs) \text{ or } (ID/EX.RegisterRt = IF/ID.RegisterRt))$$
- If detected, stall and insert bubble

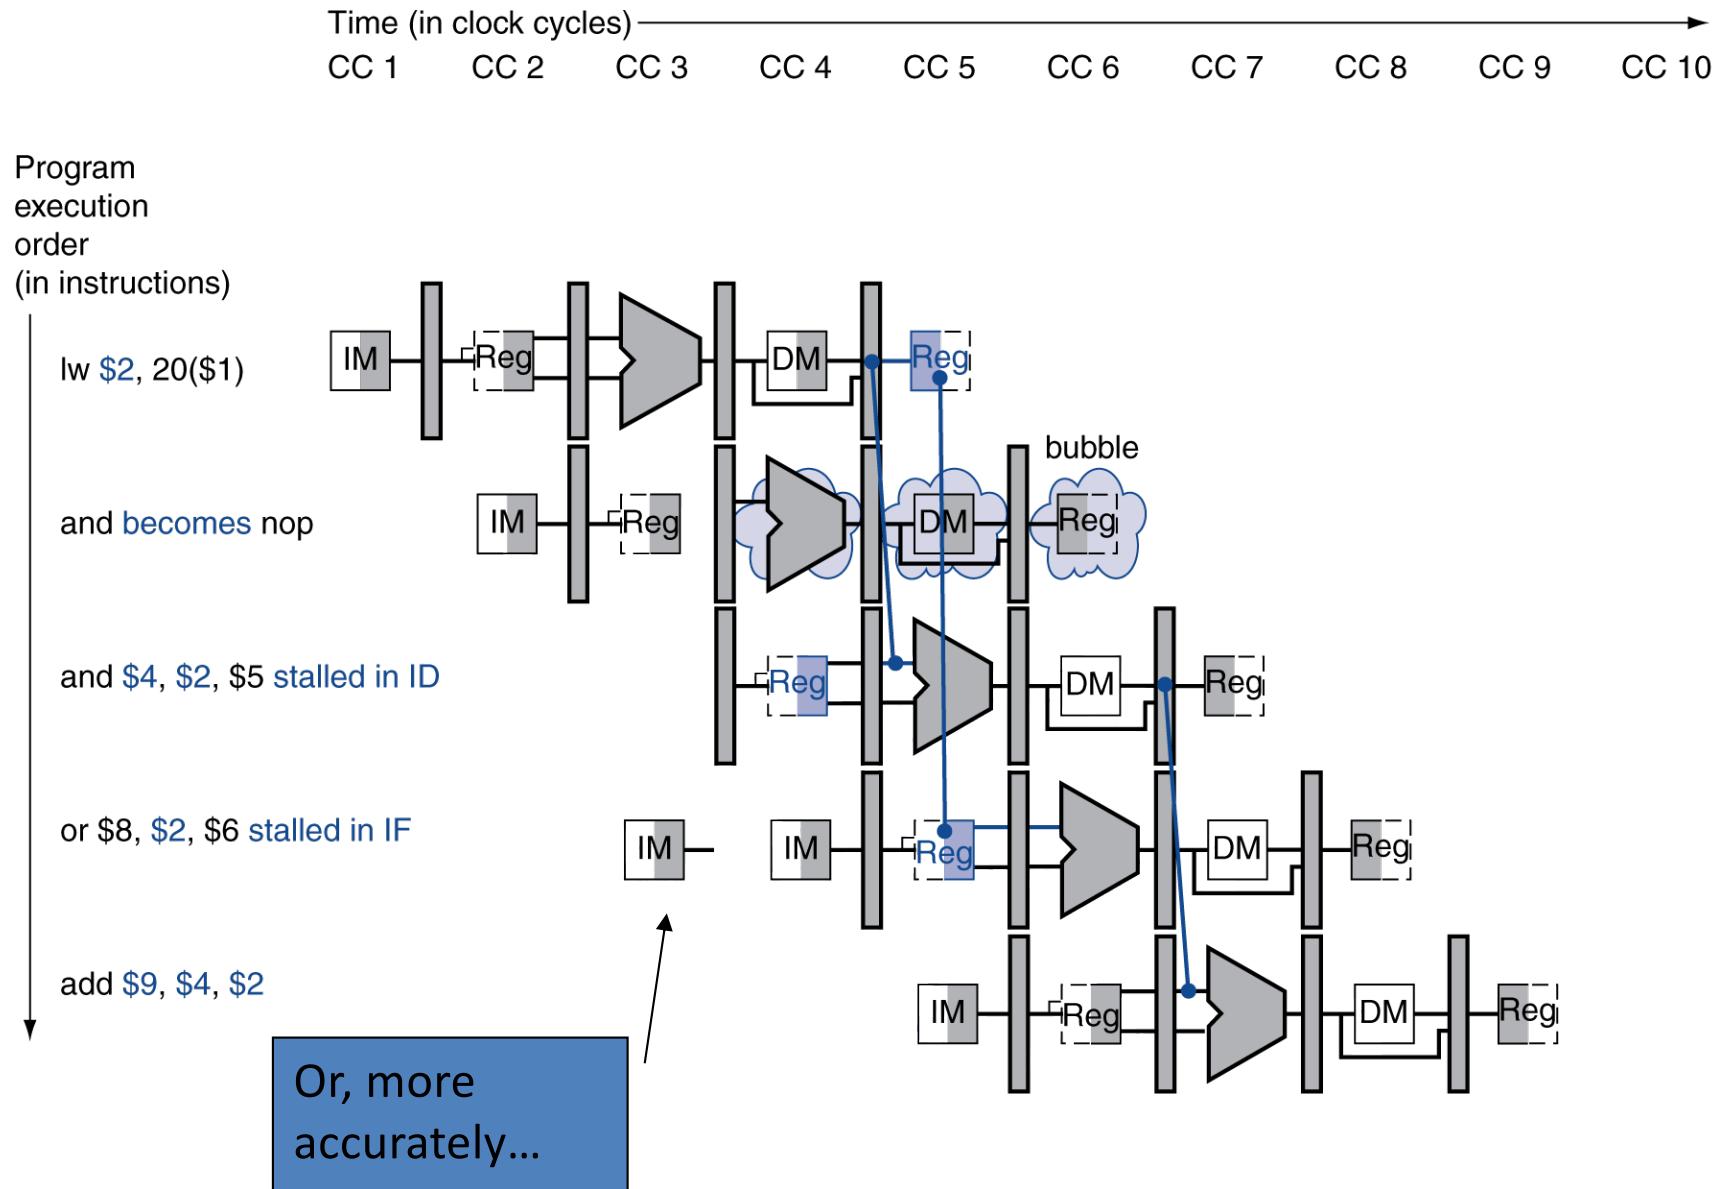
# How to Stall the Pipeline

- Force control values in ID/EX register to 0
  - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
  - Using instruction is decoded again
  - Following instruction is fetched again
  - 1-cycle stall allows MEM to read data for 1w
    - Can subsequently forward to EX stage

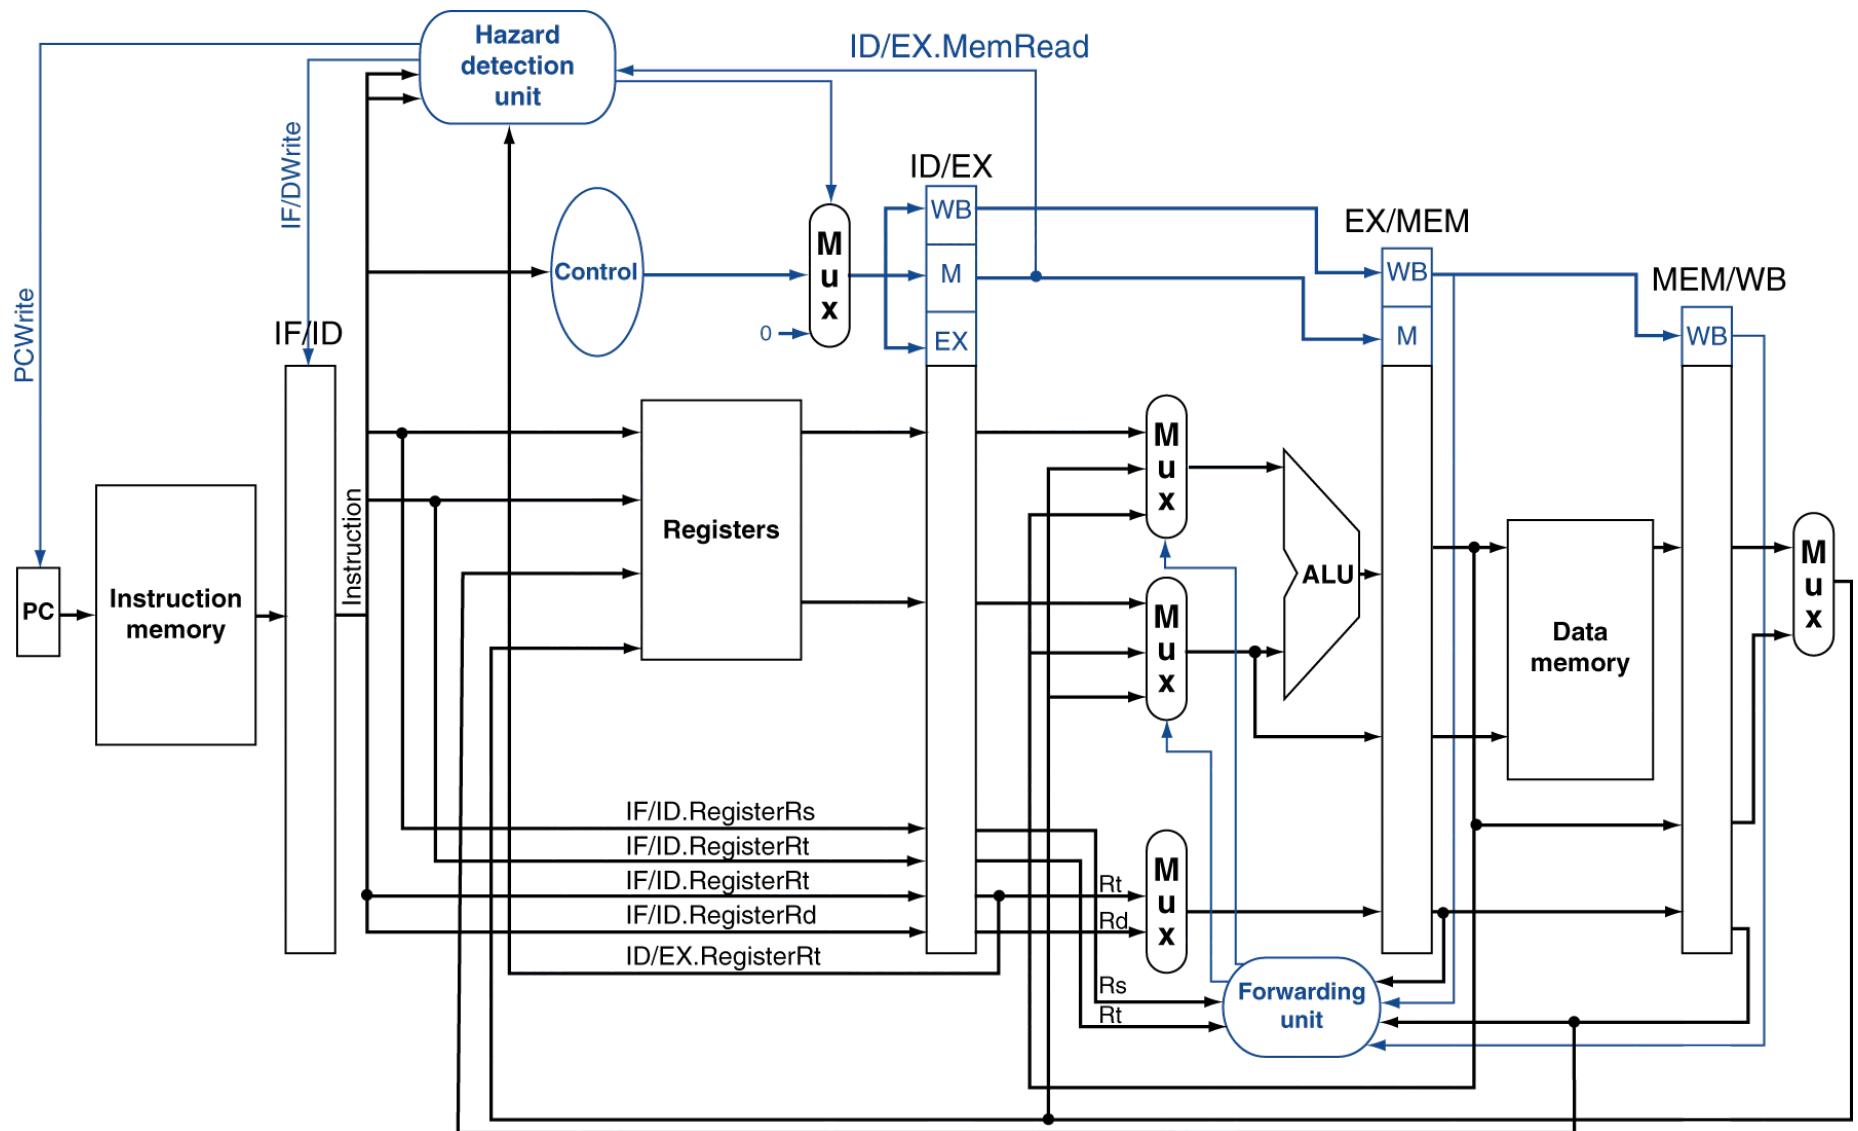
# Stall/Bubble in the Pipeline



# Stall/Bubble in the Pipeline



# Datapath with Hazard Detection



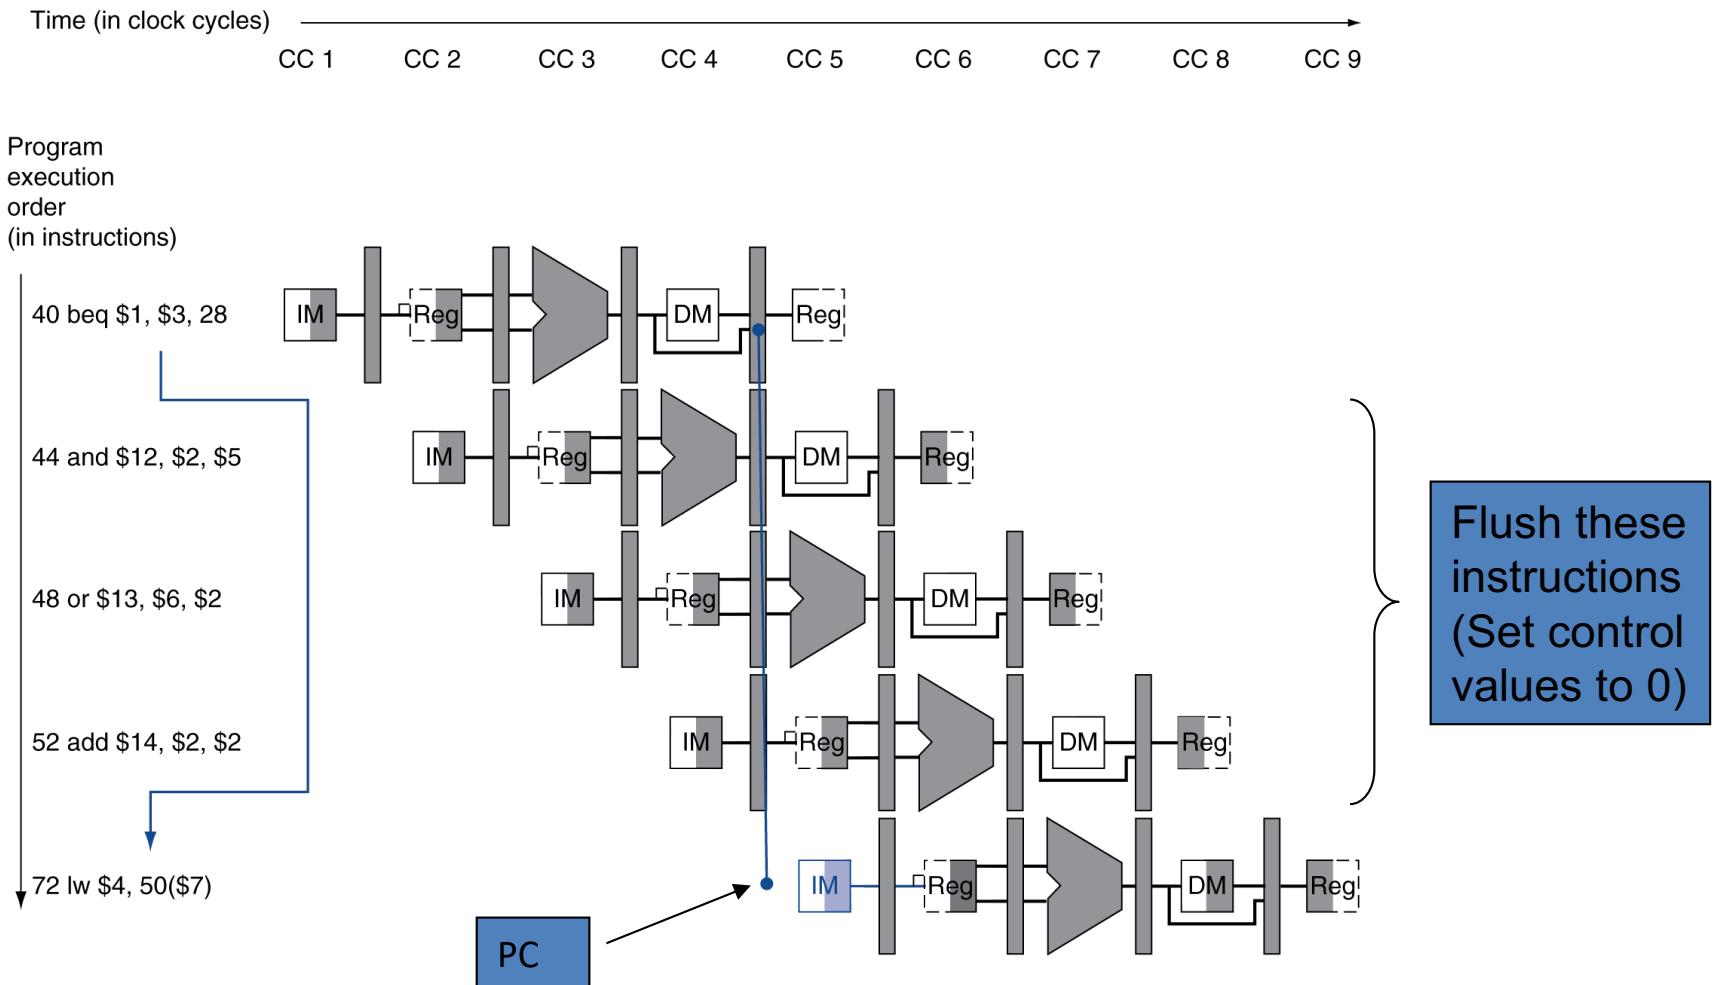
# Stalls and Performance

## The BIG Picture

- Stalls reduce performance
  - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
  - Requires knowledge of the pipeline structure

# Branch Hazards

- If branch outcome determined in MEM

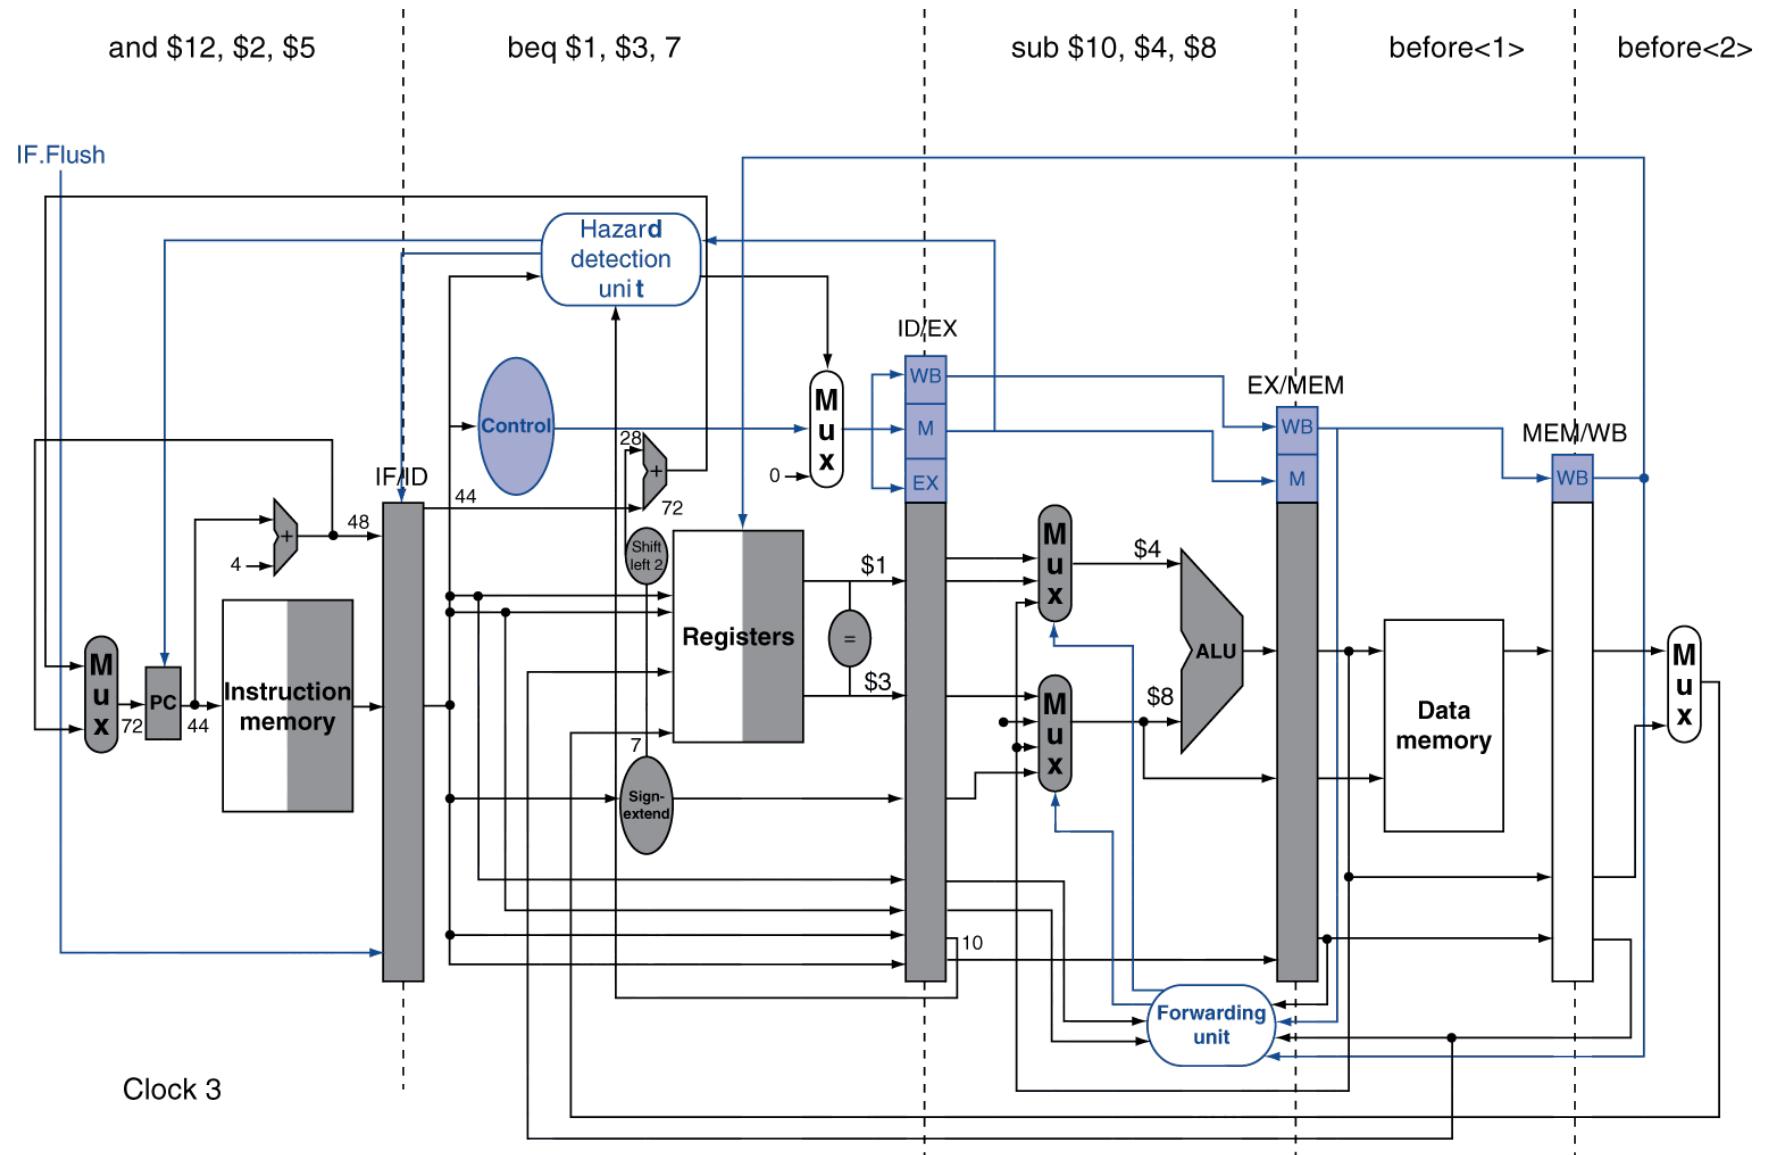


# Reducing Branch Delay

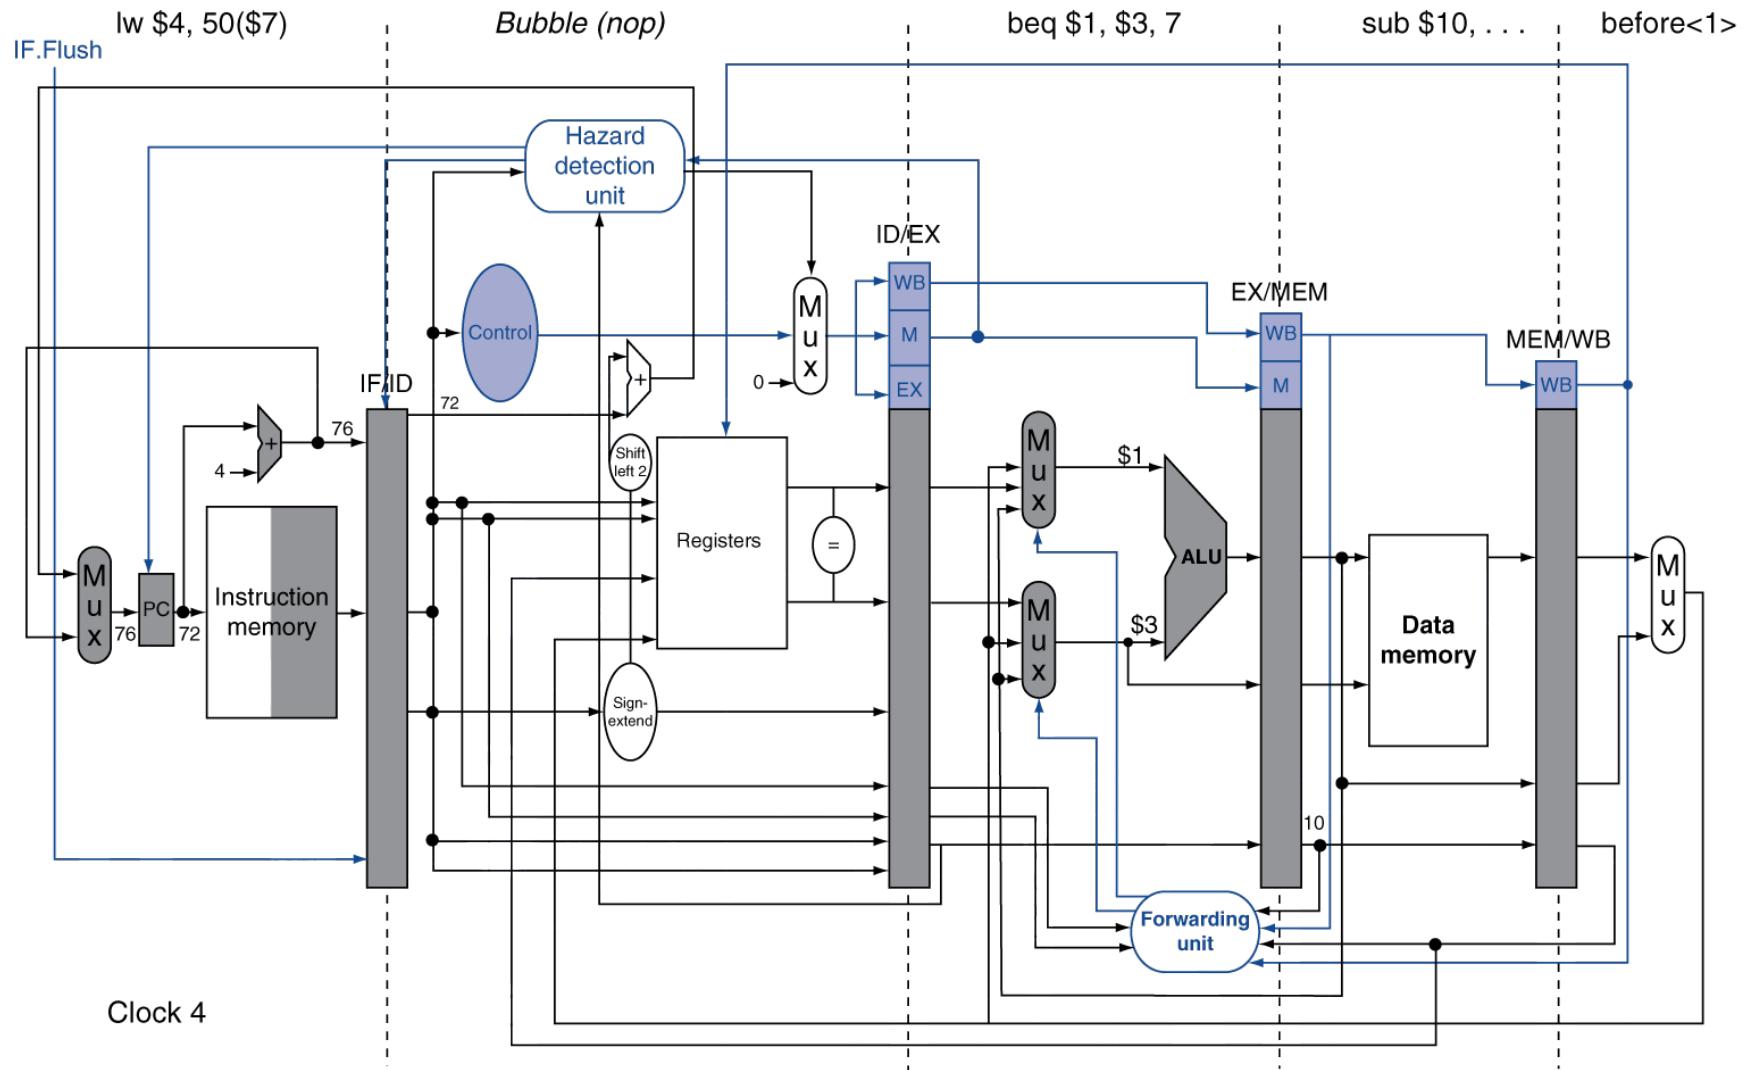
- Move hardware to determine outcome to ID stage
  - Target address adder
  - Register comparator
- Example: branch taken

```
36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
      ...
72: lw $4, 50($7)
```

# Example: Branch Taken

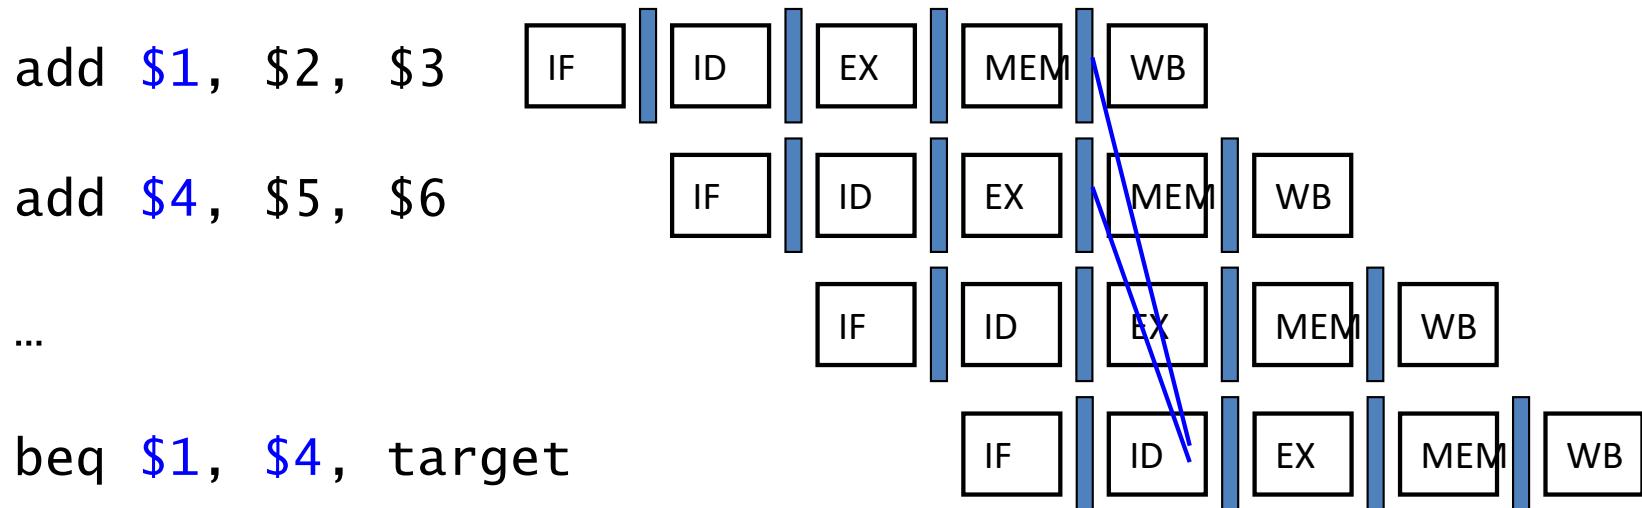


# Example: Branch Taken



# Data Hazards for Branches

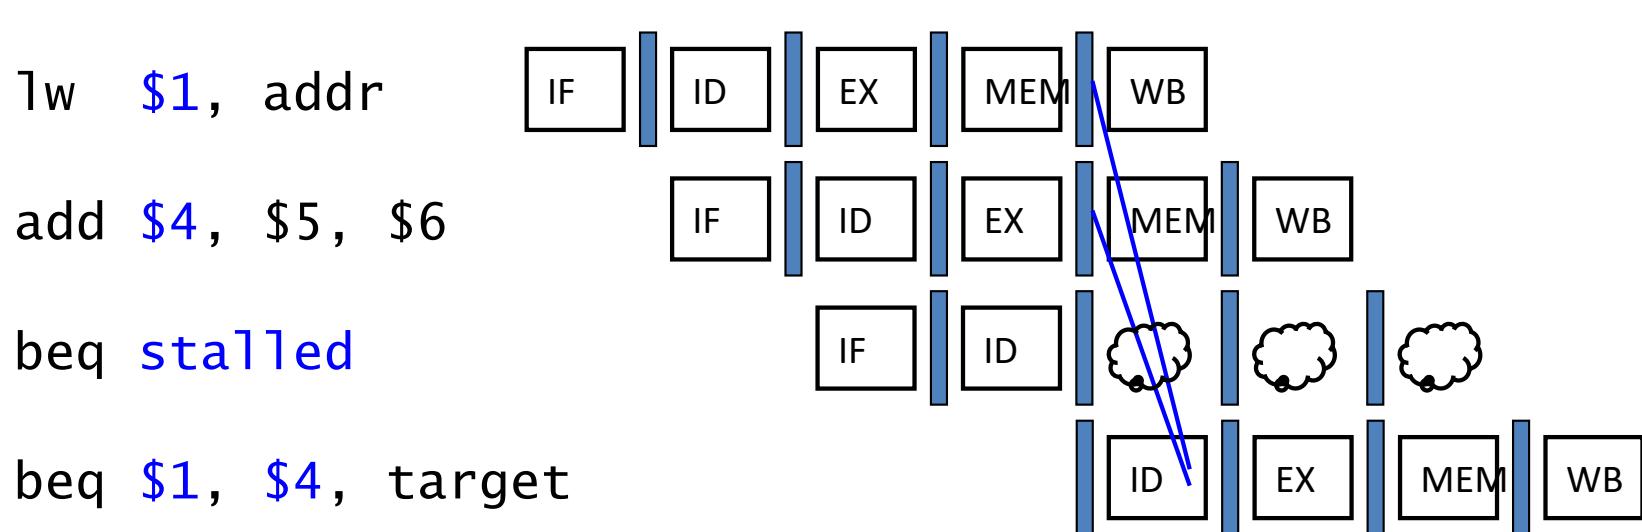
- If a comparison register is a destination of 2<sup>nd</sup> or 3<sup>rd</sup> preceding ALU instruction



- Can resolve using forwarding

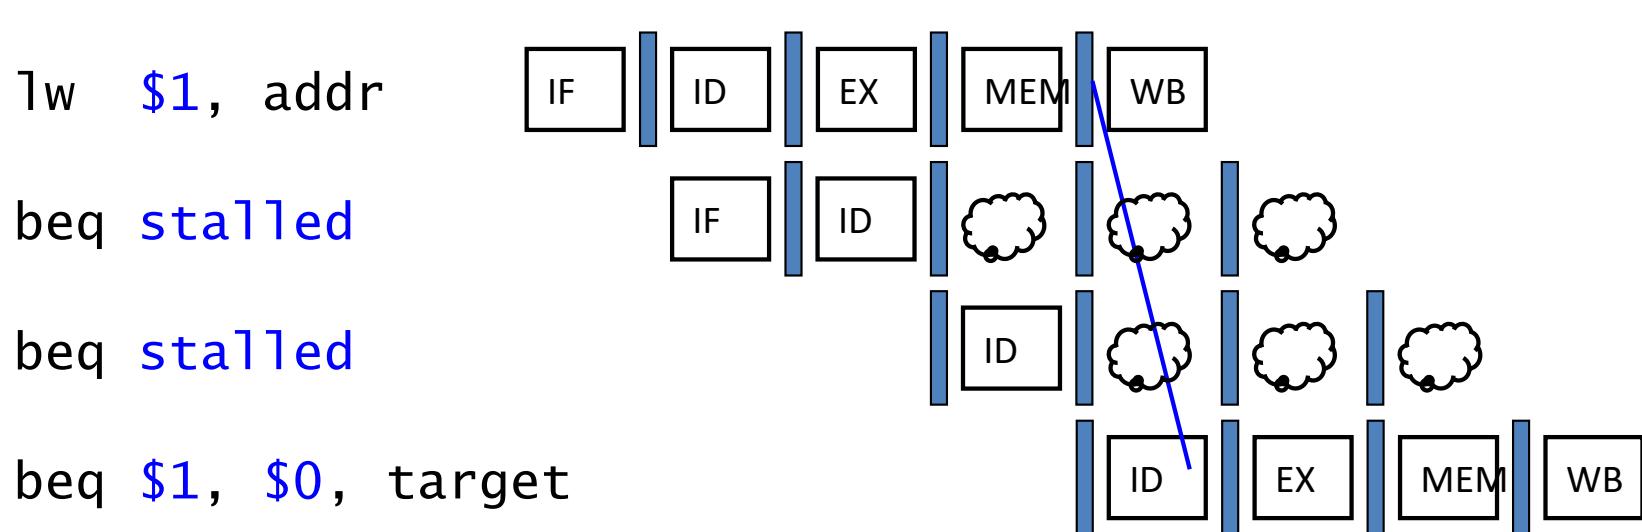
# Data Hazards for Branches

- If a comparison register is a destination of preceding ALU instruction or 2<sup>nd</sup> preceding load instruction
  - Need 1 stall cycle



# Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
  - Need 2 stall cycles



# Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
  - Different ISAs use the terms differently
- Exception
  - Arises within the CPU
    - e.g., undefined opcode, overflow, syscall, ...
- Interrupt
  - From an external I/O controller
- Dealing with them without sacrificing performance is hard

# Handling Exceptions

- In MIPS, exceptions managed by a System Control Coprocessor (CP0)
- Save PC of offending (or interrupted) instruction
  - In MIPS: Exception Program Counter (EPC)
- Save indication of the problem
  - In MIPS: Cause register
  - We'll assume 1-bit
    - 0 for undefined opcode, 1 for overflow
- Jump to handler at 8000 00180

# An Alternate Mechanism

- Vectored Interrupts
  - Handler address determined by the cause
- Example:
  - Undefined opcode: C000 0000
  - Overflow: C000 0020
  - ...: C000 0040
- Instructions either
  - Deal with the interrupt, or
  - Jump to real handler

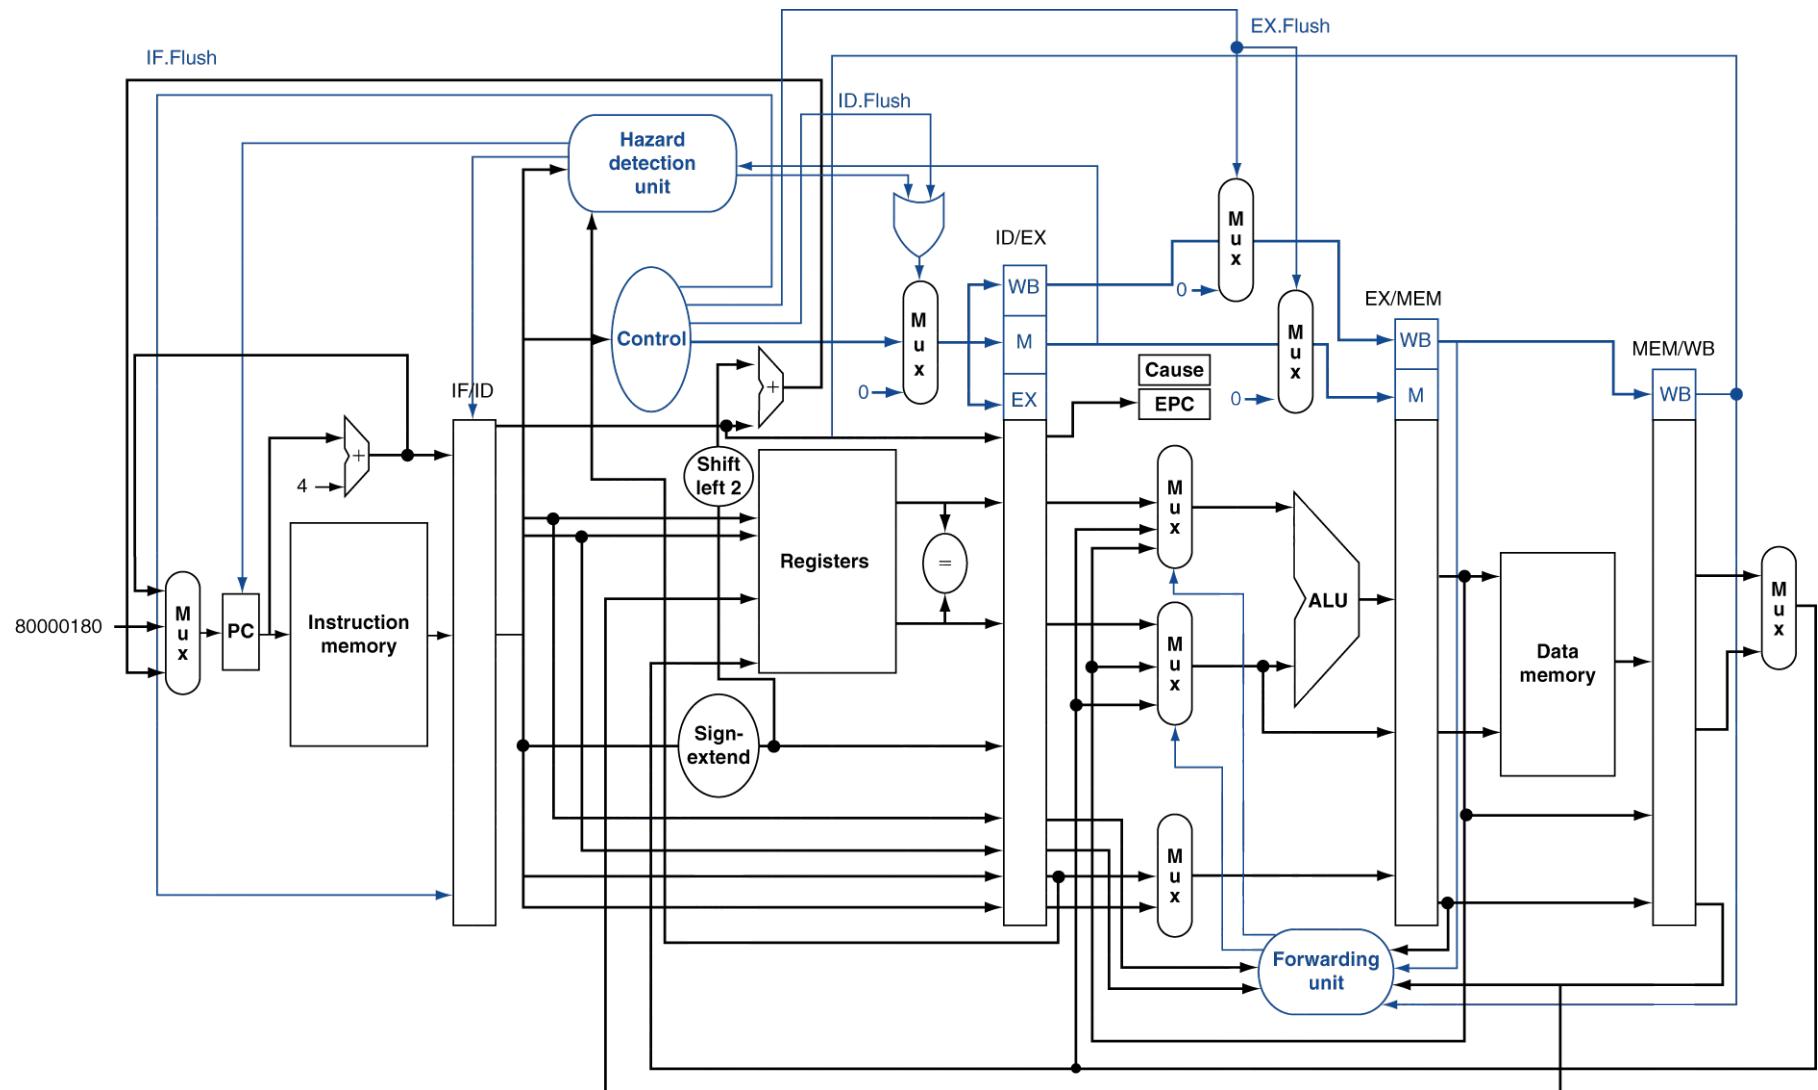
# Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
  - Take corrective action
  - use EPC to return to program
- Otherwise
  - Terminate program
  - Report error using EPC, cause, ...

# Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage
  - add \$1, \$2, \$1
    - Prevent \$1 from being clobbered
    - Complete previous instructions
    - Flush add and subsequent instructions
    - Set Cause and EPC register values
    - Transfer control to handler
- Similar to mispredicted branch
  - Use much of the same hardware

# Pipeline with Exceptions



# Exception Properties

- Restartable exceptions
  - Pipeline can flush the instruction
  - Handler executes, then returns to the instruction
    - Refetched and executed from scratch
- PC saved in EPC register
  - Identifies causing instruction
  - Actually PC + 4 is saved
    - Handler must adjust

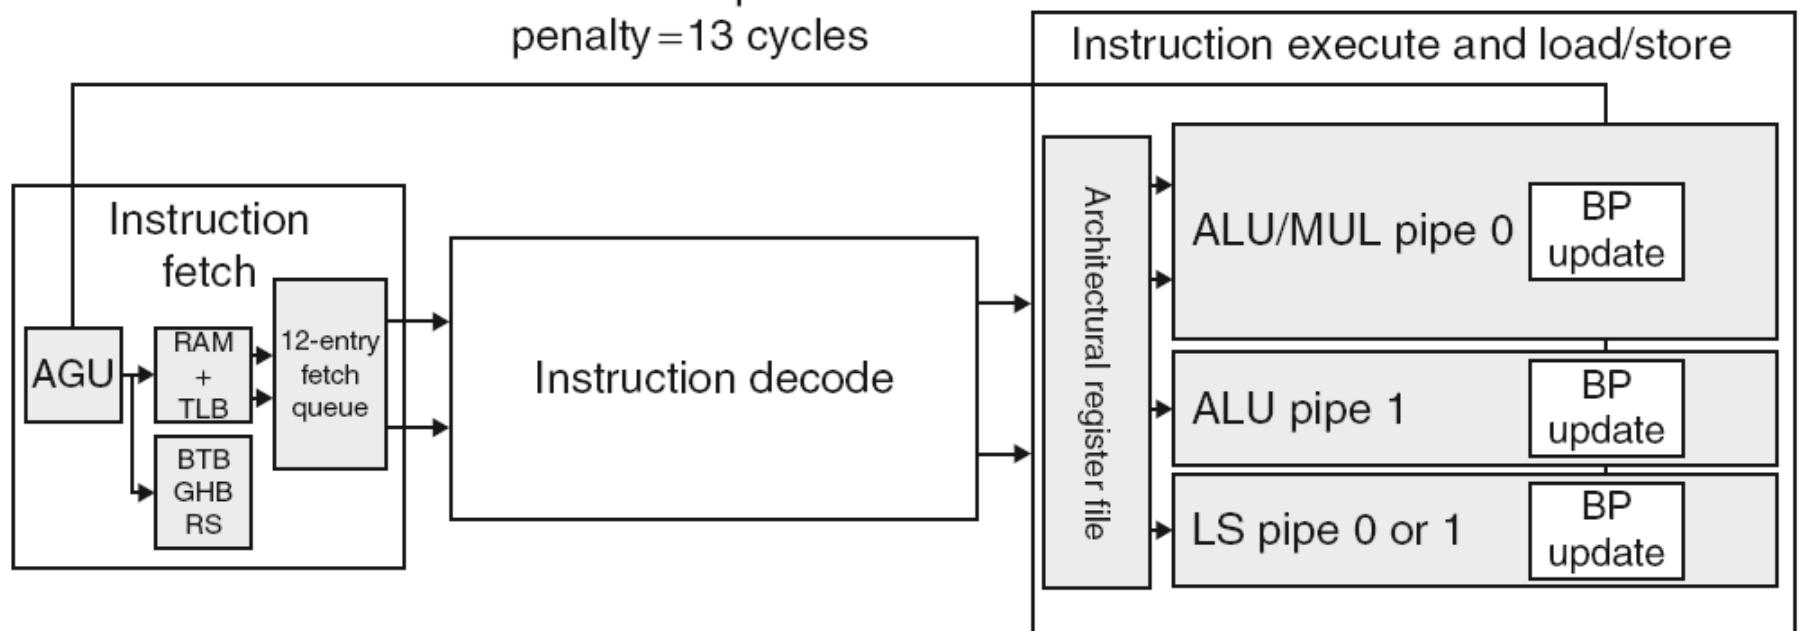
# Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 <sup>st</sup> level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 <sup>nd</sup> level caches/core	128-1024 KiB	256 KiB
3 <sup>rd</sup> level caches (shared)	-	2- 8 MB

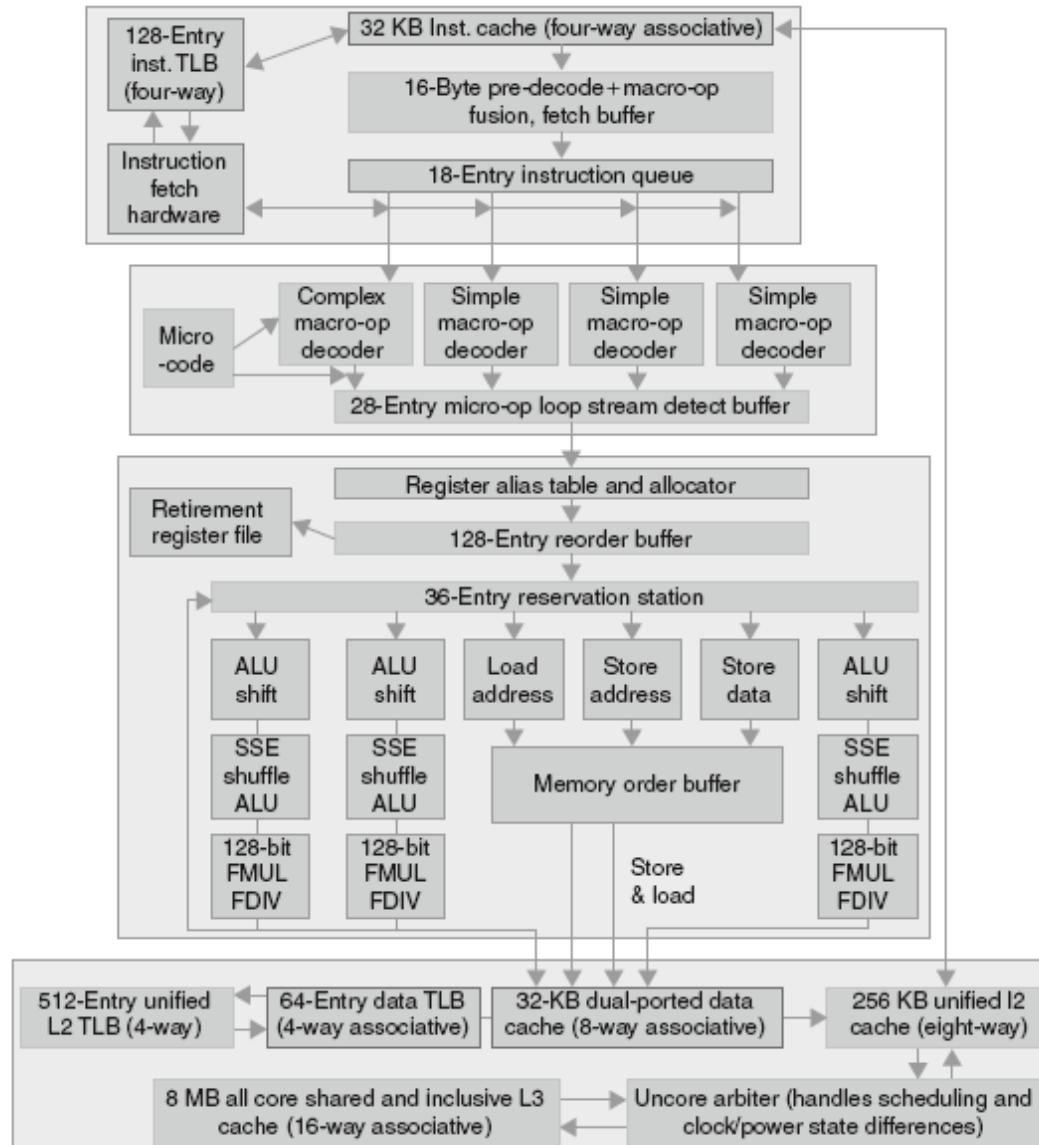
# ARM Cortex-A8 Pipeline

F0 F1 F2 D0 D1 D2 D3 D4 E0 E1 E2 E3 E4 E5

Branch mispredict  
penalty = 13 cycles



# Core i7 Pipeline



**Question:** For each code sequences below, choose one of the statements below:

1:

```
lw $t0,0($t0)  
add $t1,$t0,$t0
```

2:

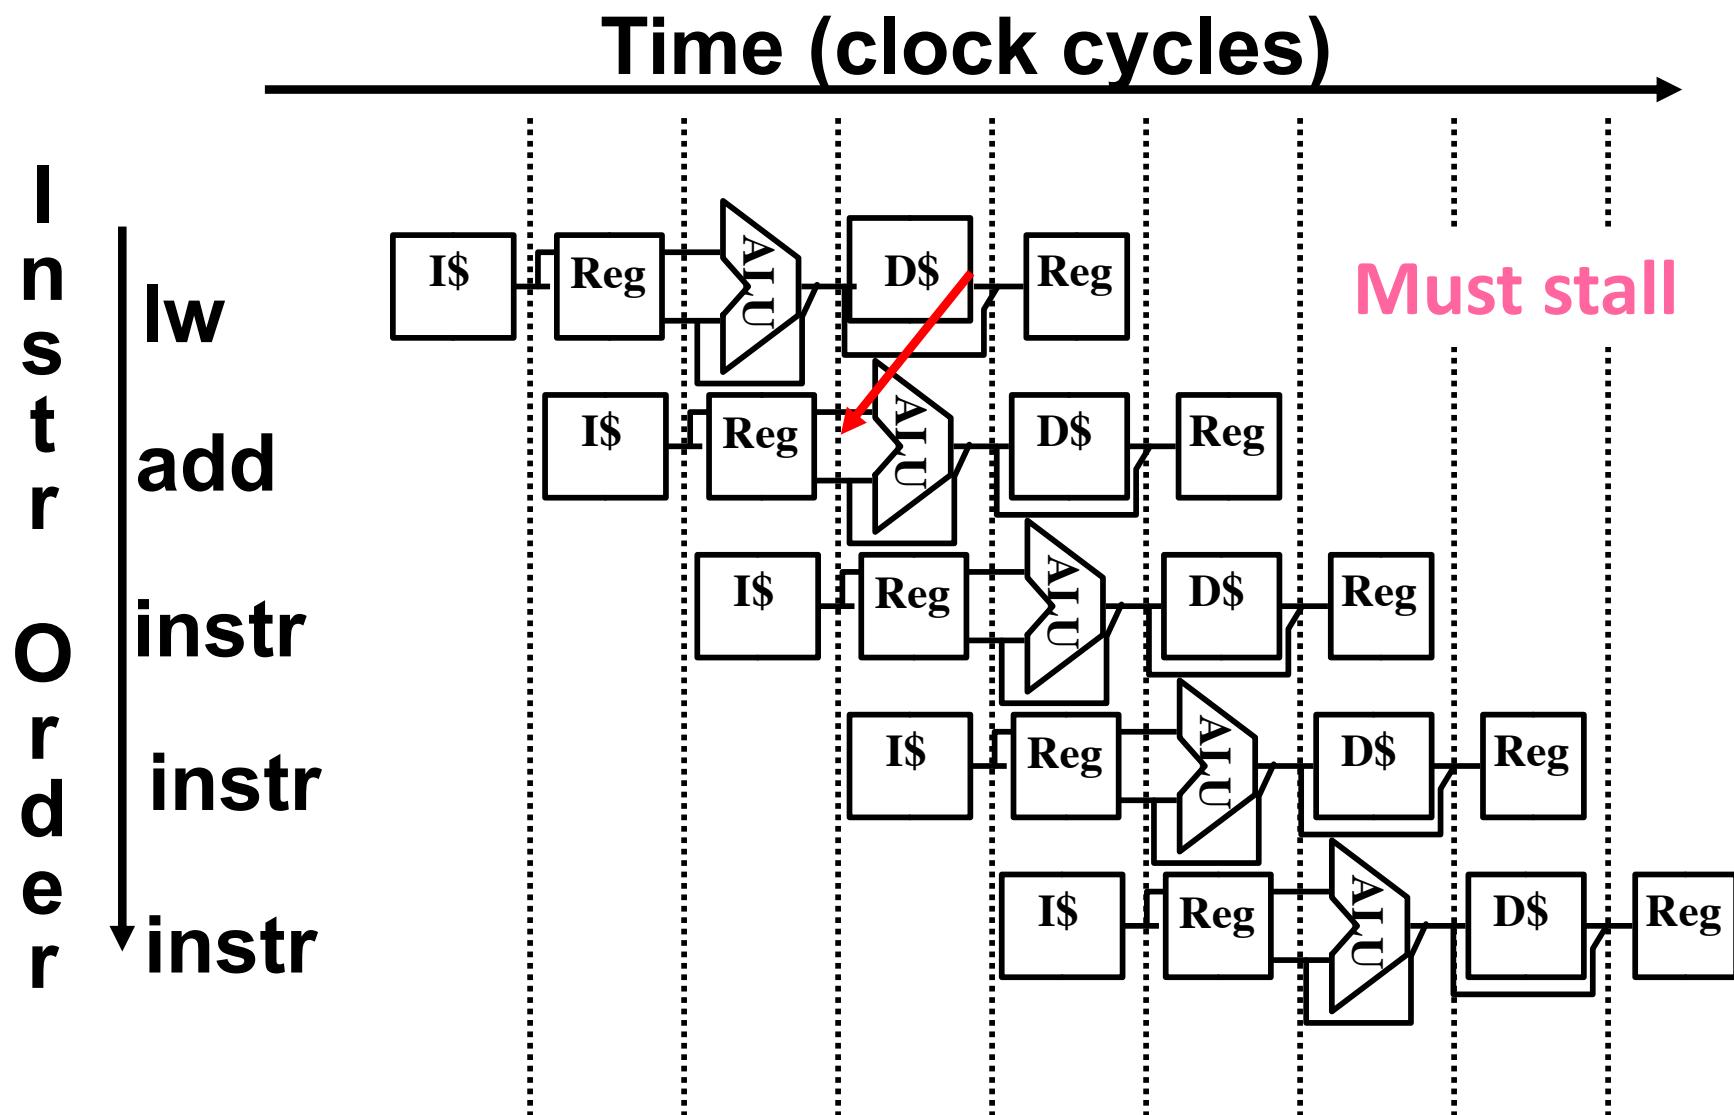
```
add $t1,$t0,$t0  
addi $t2,$t0,5  
addi $t4,$t1,5
```

3:

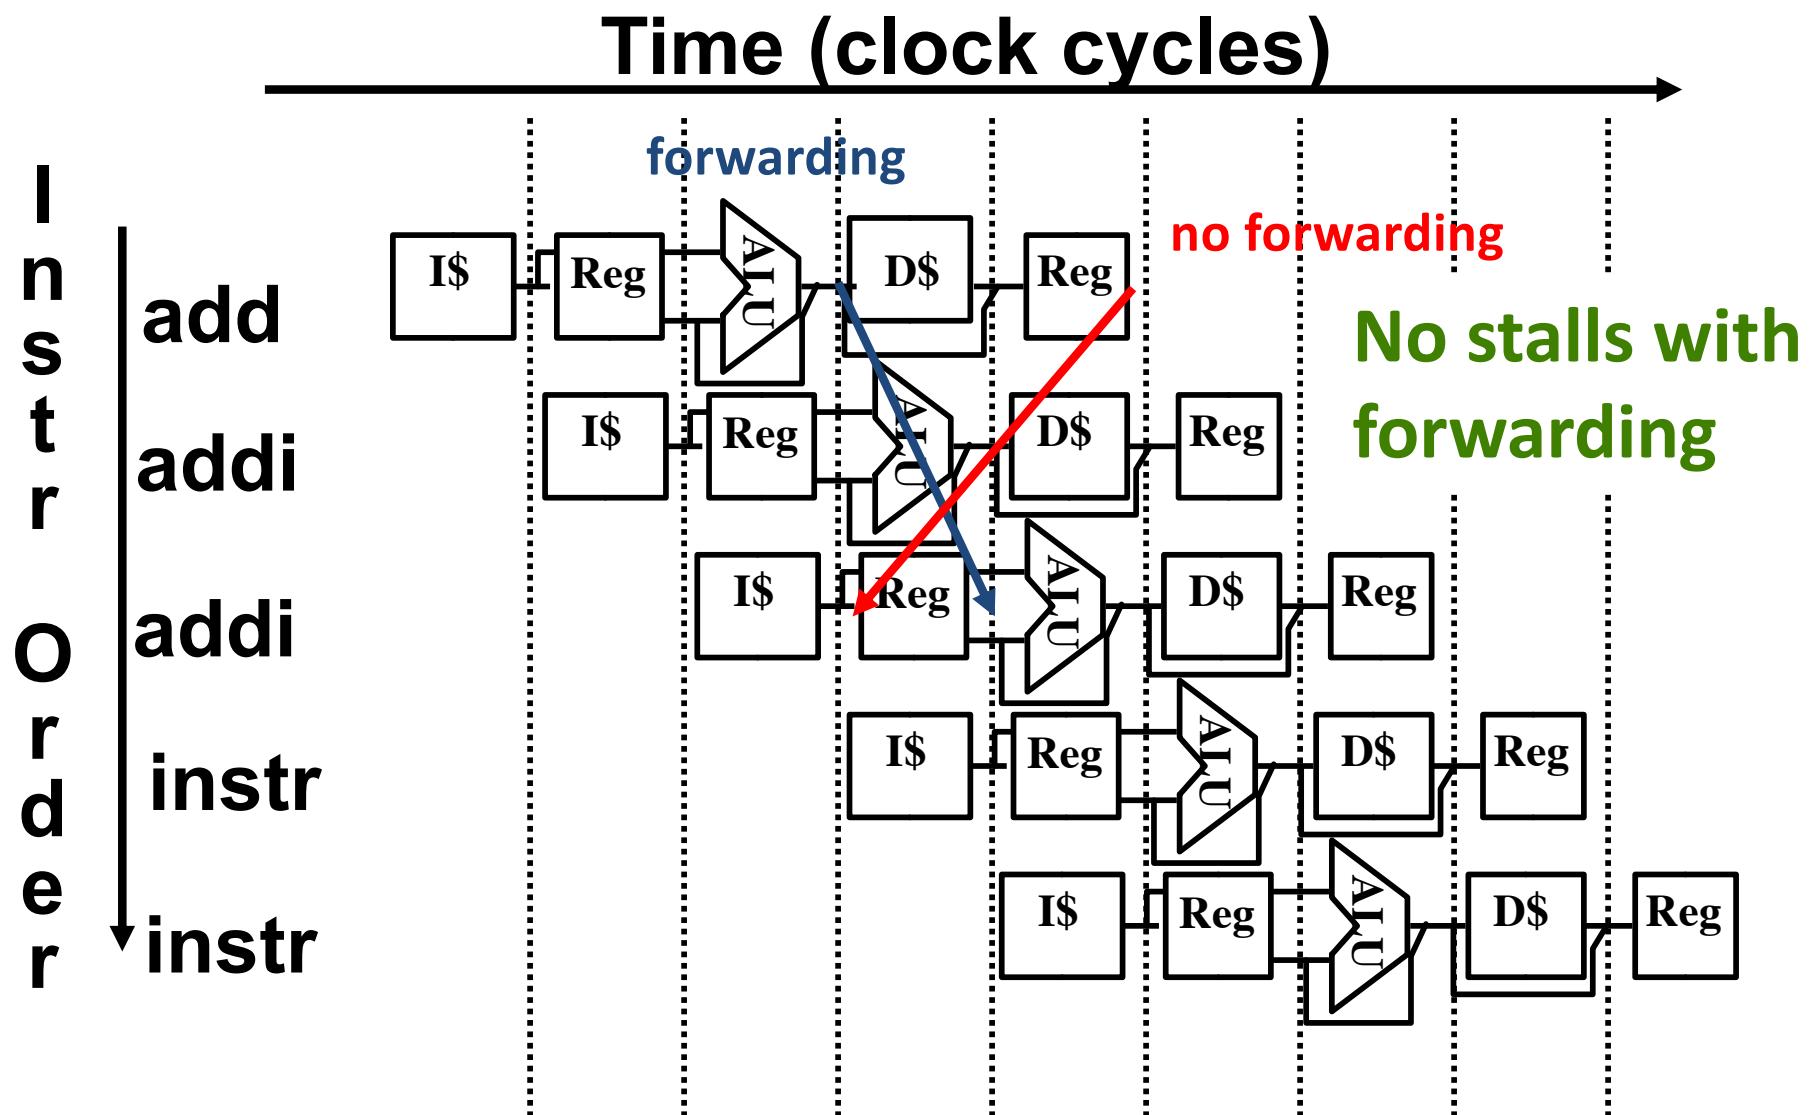
```
addi $t1,$t0,1  
addi $t2,$t0,2  
addi $t3,$t0,2  
addi $t3,$t0,4  
addi $t5,$t1,5
```

- No stalls as is**
- No stalls with forwarding**
- Must stall**

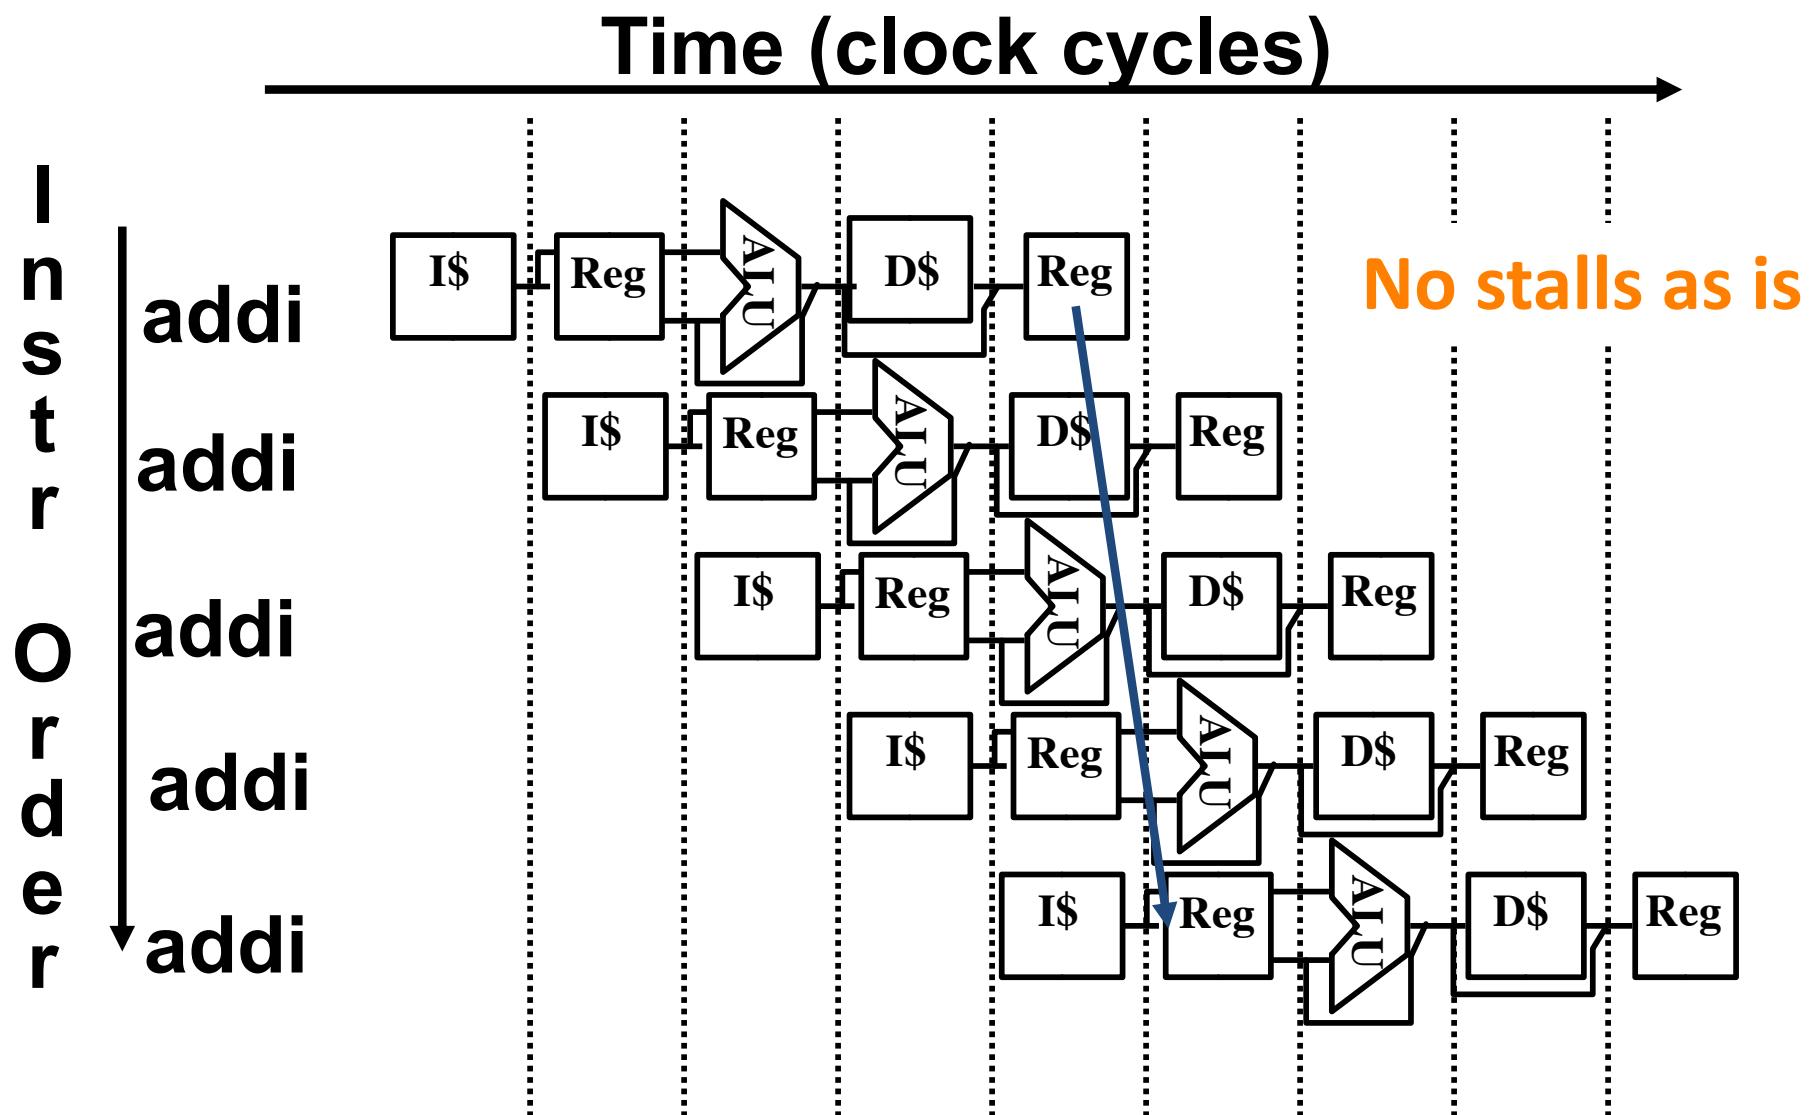
# Code Sequence 1



# Code Sequence 2



# Code Sequence 3



# Summary

- Hazards reduce effectiveness of pipelining
  - Cause stalls/bubbles
- Structural Hazards
  - Conflict in use of datapath component
- Data Hazards
  - Need to wait for result of a previous instruction
- Control Hazards
  - Address of next instruction uncertain/unknown
  - Branch and jump delay slots